

quantum algorithms for sieving short(est) vectors in lattices

Elena Kirshanova, Erik Mårtensson, Eamonn Postlethwaite, Subhayan Roy Moulik

London-ish Lattice Coding & Crypto Meeting
Imperial College
18th September, 2019



NIST PQC Standardisation Process - Round 2

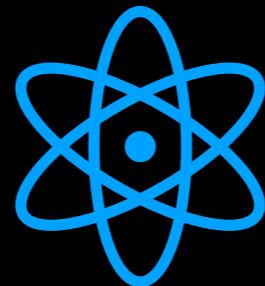
Bike	Classic McEliece	CRYSTALS-DILITHIUM	CRYSTALS-KYBER
FALCON	FrodoKEM	GeMSS	HQC
LAC	LEDAcrypt	LUOV	MQDSS
NewHope	NTRU	NTRU Prime	NTS-KEM
Picnic	qTESLA	Rainbow	ROLLO
Round5	RQC	SABER	SIKE
	SPHINCS+	Three Bears	

NIST PQC Standardisation Process - Round 2

Bike	Classic McEliece	CRYSTALS-DILITHIUM	CRYSTALS-KYBER
FALCON	FrodoKEM	GeMSS	HQC
LAC	LEDAcrypt	LUOV	MQDSS
NewHope	NTRU	NTRU Prime	NTS-KEM
Picnic	qTESLA	Rainbow	ROLLO
Round5	RQC	SABER	SIKE
	SPHINCS+	Three Bears	



001001010
111000010110



0010100100101

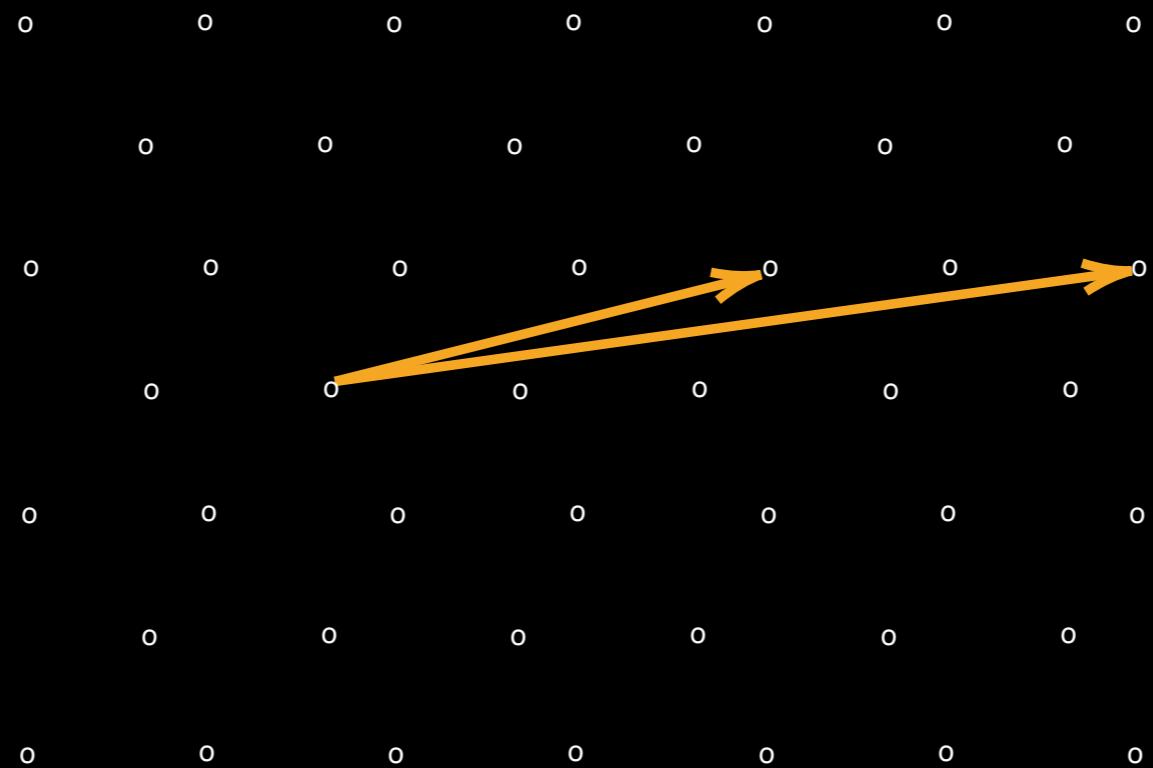
Lattices and SVP

Given m linear independent vectors,

$$B = \{b_1, \dots, b_m : b_i \in \mathbb{R}^n\}$$

the lattice generated by them is defined as

$$\mathcal{L}(B) = \left\{ \sum_{i=1}^m z_i b_i \mid z_i \in \mathbb{Z} \right\}$$



Lattices and SVP

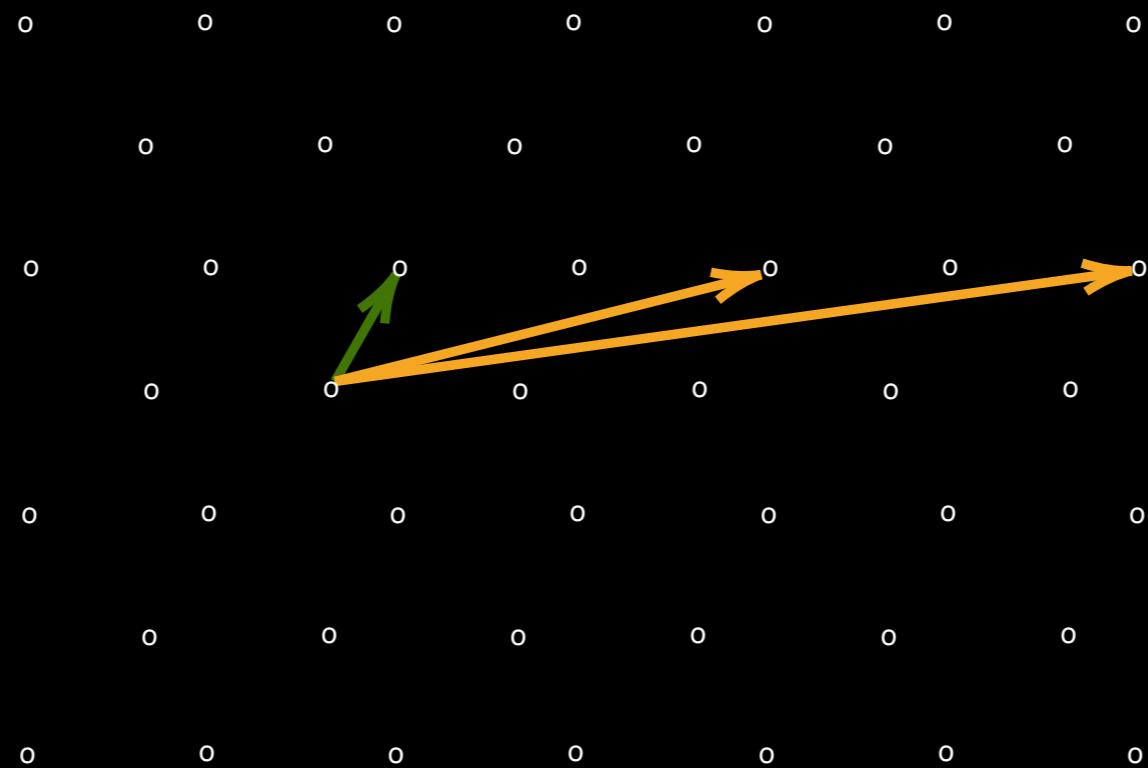
Given m linear independent vectors,

$$B = \{b_1, \dots, b_m : b_i \in \mathbb{R}^n\}$$

the lattice generated by them is defined as

$$\mathcal{L}(B) = \left\{ \sum_{i=1}^m z_i b_i \mid z_i \in \mathbb{Z} \right\}$$

Shortest Vector Problem (SVP): Given $\mathcal{L}(B)$, find the shortest non-zero vector



Lattices and SVP

Given m linear independent vectors,

$$B = \{b_1, \dots, b_m : b_i \in \mathbb{R}^n\}$$

the lattice generated by them is defined as

$$\mathcal{L}(B) = \left\{ \sum_{i=1}^m z_i b_i \mid z_i \in \mathbb{Z} \right\}$$

Shortest Vector Problem (SVP): Given $\mathcal{L}(B)$, find the shortest non-zero vector

Best known algorithms are atleast exponential time
provable vs heuristic

SVP hardness

Theory

	Algorithm	$\log_2(\text{Time})$	$\log_2(\text{Space})$
Proven SVP	Enumeration [Poh81, Kan83, ..., MW15, AN17]	$O(n \log n)$	$O(\log n)$
	AKS-sieve [AKS01, NV08, MV10, HPS11]	$3.398n$	$1.985n$
	ListSieve [MV10, MDB14]	$3.199n$	$1.327n$
	Birthday sieves [PS09, HPS11]	$2.465n$	$1.233n$
	Enumeration/DGS hybrid [CCL17]	$2.048n$	$0.500n$
	Voronoi cell algorithm [AEVZ02, MV10b]	$2.000n$	$1.000n$
	Quantum sieve [LMP13, LMP15]	$1.799n$	$1.286n$
	Quantum enum/DGS [CCL17]	$1.256n$	0.500n
Heuristic SVP	Discrete Gaussian sampling [ADRS15, ADS15, AS18]	1.000n	$1.000n$
	The Nguyen–Vidick sieve [NV08]	$0.415n$	$0.208n$
	The GaussSieve [MV10, ..., IKMT14, BNvdP16, YKYC17]	$0.415n$	$0.208n$
	Triple sieve [BLS16, HK17]	$0.396n$	$0.189n$
	Two-level sieve [WLTB11]	$0.384n$	$0.256n$
	Three-level sieve [ZPH13]	$0.3778n$	$0.283n$
	Overlattice sieve [BGJ14]	$0.3774n$	$0.293n$
	Triple sieve with NNS [HK17, HKL18]	$0.359n$	0.189n
	Hyperplane LSH [Cha02, Laa15, ..., LM18, Duc18]	$0.337n$	$0.337n$
	Graph-based NNS [EPY99, DCL11, MPLK14, Laa18]	$0.327n$	$0.282n$
	Hypercube LSH [TT07, Laa17]	$0.322n$	$0.322n$
	Quantum sieve [LMP13, LMP15]	$0.312n$	$0.208n$
	May–Ozerov NNS [MO15, BGJ15]	$0.311n$	$0.311n$
	Spherical LSH [AINR14, LdW15]	$0.298n$	$0.298n$
	Cross-polytope LSH [TT07, AILRS15, BL16, KW17]	$0.298n$	$0.298n$
	Spherical LSF [BDGL16, MLB17, ALRW17, Chr17]	0.292n	$0.292n$
	Quantum NNS sieve [LMP15, Laa16]	0.265n	$0.265n$

n = dimension

slide stolen from Thijs Laarhoven
<http://www.thijs.com/docs/pqc18-slides.pdf>

BUT EXPONENTIAL TIME ALGORITHM !!

BUT EXPONENTIAL TIME ALGORITHM !!

Efficiency vs Security: Battle for dimensions

HALL OF FAME

Position	Dimension	Euclidean norm	Seed	Contestant	Solution
1	157	3320	0	L. Ducas, M. Stevens, W. van Woerden M. Albrecht, L. Ducas, G. Herold, E. Kirshanova, E. Postlethwaite, M. Stevens, P. Karpman	vec
2	155	3165	0	Martin Albrecht, Leo Ducas, Gottfried Herold, Elena Kirshanova, Eamonn Postlethwaite, Marc Stevens	vec
3	153	3192	0	Kenji KASHIWABARA and Tadanori TERUYA	vec
4	152	3217	0	Martin Albrecht, Leo Ducas, Gottfried Herold, Elena Kirshanova, Eamonn Postlethwaite, Marc Stevens	vec
5	151	3233	0	Complete list	vec

*Darmstadt lattice challenge : SVP

SVP hardness

Theory

	Algorithm	$\log_2(\text{Time})$	$\log_2(\text{Space})$
Proven SVP	Enumeration [Poh81, Kan83, ..., MW15, AN17]	$O(n \log n)$	$O(\log n)$
	AKS-sieve [AKS01, NV08, MV10, HPS11]	$3.398n$	$1.985n$
	ListSieve [MV10, MDB14]	$3.199n$	$1.327n$
	Birthday sieves [PS09, HPS11]	$2.465n$	$1.233n$
	Enumeration/DGS hybrid [CCL17]	$2.048n$	$0.500n$
	Voronoi cell algorithm [AEVZ02, MV10b]	$2.000n$	$1.000n$
	Quantum sieve [LMP13, LMP15]	$1.799n$	$1.286n$
	Quantum enum/DGS [CCL17]	$1.256n$	0.500n
Heuristic SVP	Discrete Gaussian sampling [ADRS15, ADS15, AS18]	1.000n	$1.000n$
	The Nguyen–Vidick sieve [NV08]	$0.415n$	$0.208n$
	The GaussSieve [MV10, ..., IKMT14, BNvdP16, YKYC17]	$0.415n$	$0.208n$
	Triple sieve [BLS16, HK17]	$0.396n$	$0.189n$
	Two-level sieve [WLTB11]	$0.384n$	$0.256n$
	Three-level sieve [ZPH13]	$0.3778n$	$0.283n$
	Overlattice sieve [BGJ14]	$0.3774n$	$0.293n$
	Triple sieve with NNS [HK17, HKL18]	0.359n	0.189n
	Hyperplane LSH [Cha02, Laa15, ..., LM18, Duc18]	$0.337n$	$0.337n$
	Graph-based NNS [EPY99, DCL11, MPLK14, Laa18]	0.327n	0.282n
	Hypercube LSH [TT07, Laa17]	$0.322n$	$0.322n$
	Quantum sieve [LMP13, LMP15]	$0.312n$	$0.208n$
	May–Ozerov NNS [MO15, BGJ15]	$0.311n$	$0.311n$
	Spherical LSH [AINR14, LdW15]	$0.298n$	$0.298n$
	Cross-polytope LSH [TT07, AILRS15, BL16, KW17]	$0.298n$	$0.298n$
	Spherical LSF [BDGL16, MLB17, ALRW17, Chr17]	0.292n	$0.292n$
	Quantum NNS sieve [LMP15, Laa16]	0.265n	$0.265n$

n = dimension

slide stolen from Thijs Laarhoven
<http://www.thijs.com/docs/pqc18-slides.pdf>

Heuristic Assumption:

Normalised vectors are uniformly distributed on the unit sphere

Algorithmic idea for sieving

Input : $\{b_1, \dots, b_n\}$ s.t. $b_i \in \mathbb{R}^n$

Output : $v \in \mathbb{R}^n$

1. Sample $L = \exp(n)$ vectors, put them in a list
2. Combine $k (= \text{constant})$ of them ‘cleverly’ to obtain L' shorter vectors
3. Repeat step 2 with the new list of vectors, $\text{poly}(n)$ times
4. Output the shortest vector from the latest list.

[KMPR19] : Can solve SVP, heuristically, in time $2^{0.2989n+o(n)}$ and (classical) memory $2^{0.1385n+o(n)}$, in the QRAM model. This can be implemented on a quantum computer equipped with $\text{poly}(n)$ quantum memory.

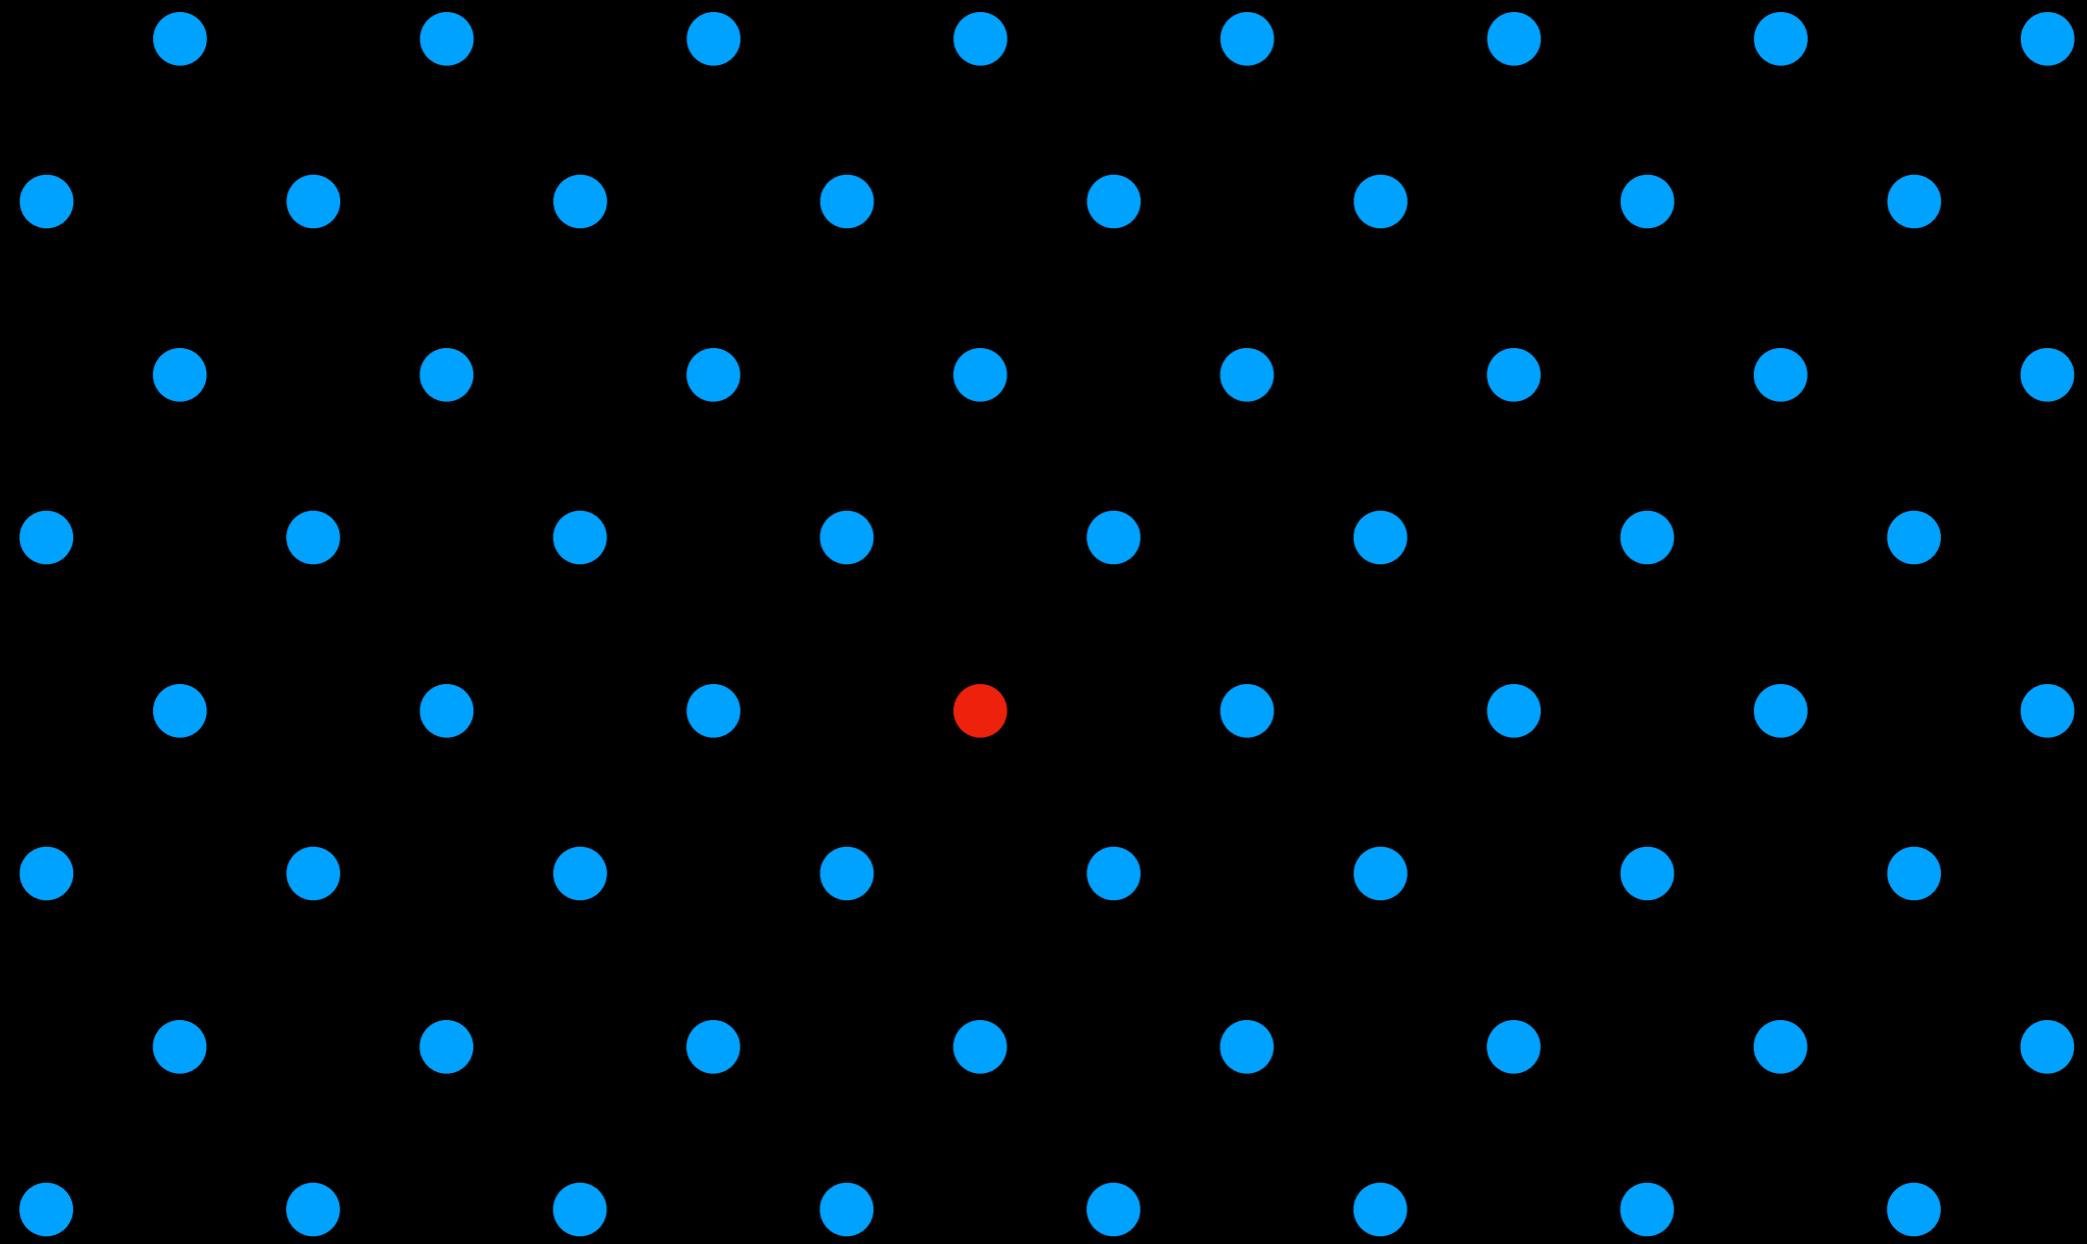
[KMPR19] : Can solve SVP, heuristically, in time $2^{0.1037n+o(n)}$ and (quantum) memory $2^{0.2075n+o(n)}$, in the quantum circuit model.

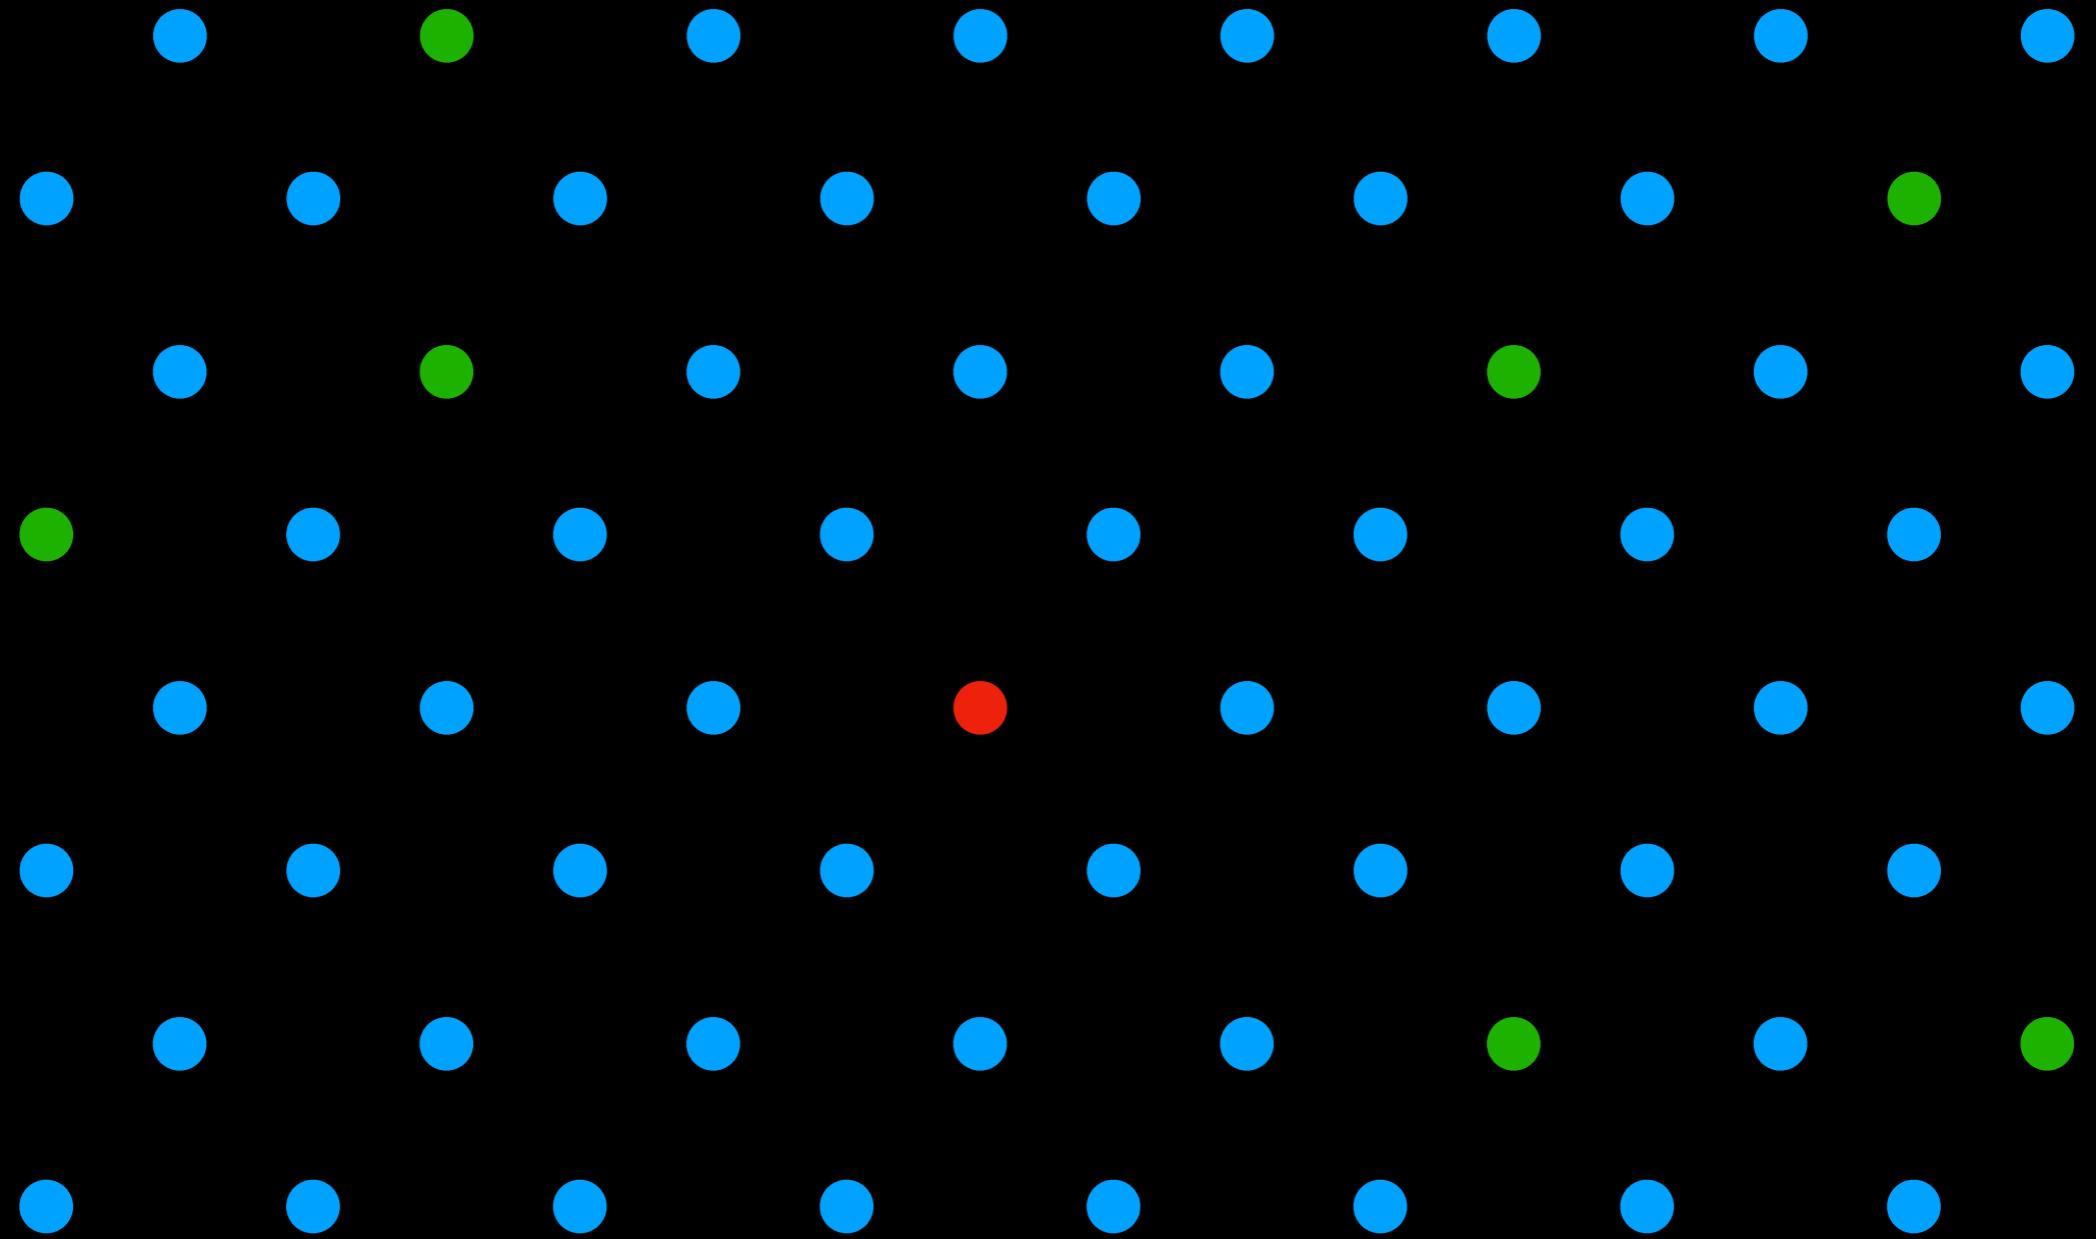
[KMPR19] : Can solve SVP, heuristically, in time $2^{0.2989n+o(n)}$ and (classical) memory $2^{0.1385n+o(n)}$, in the QRAM model. This can be implemented on a quantum computer equipped with $\text{poly}(n)$ quantum memory.

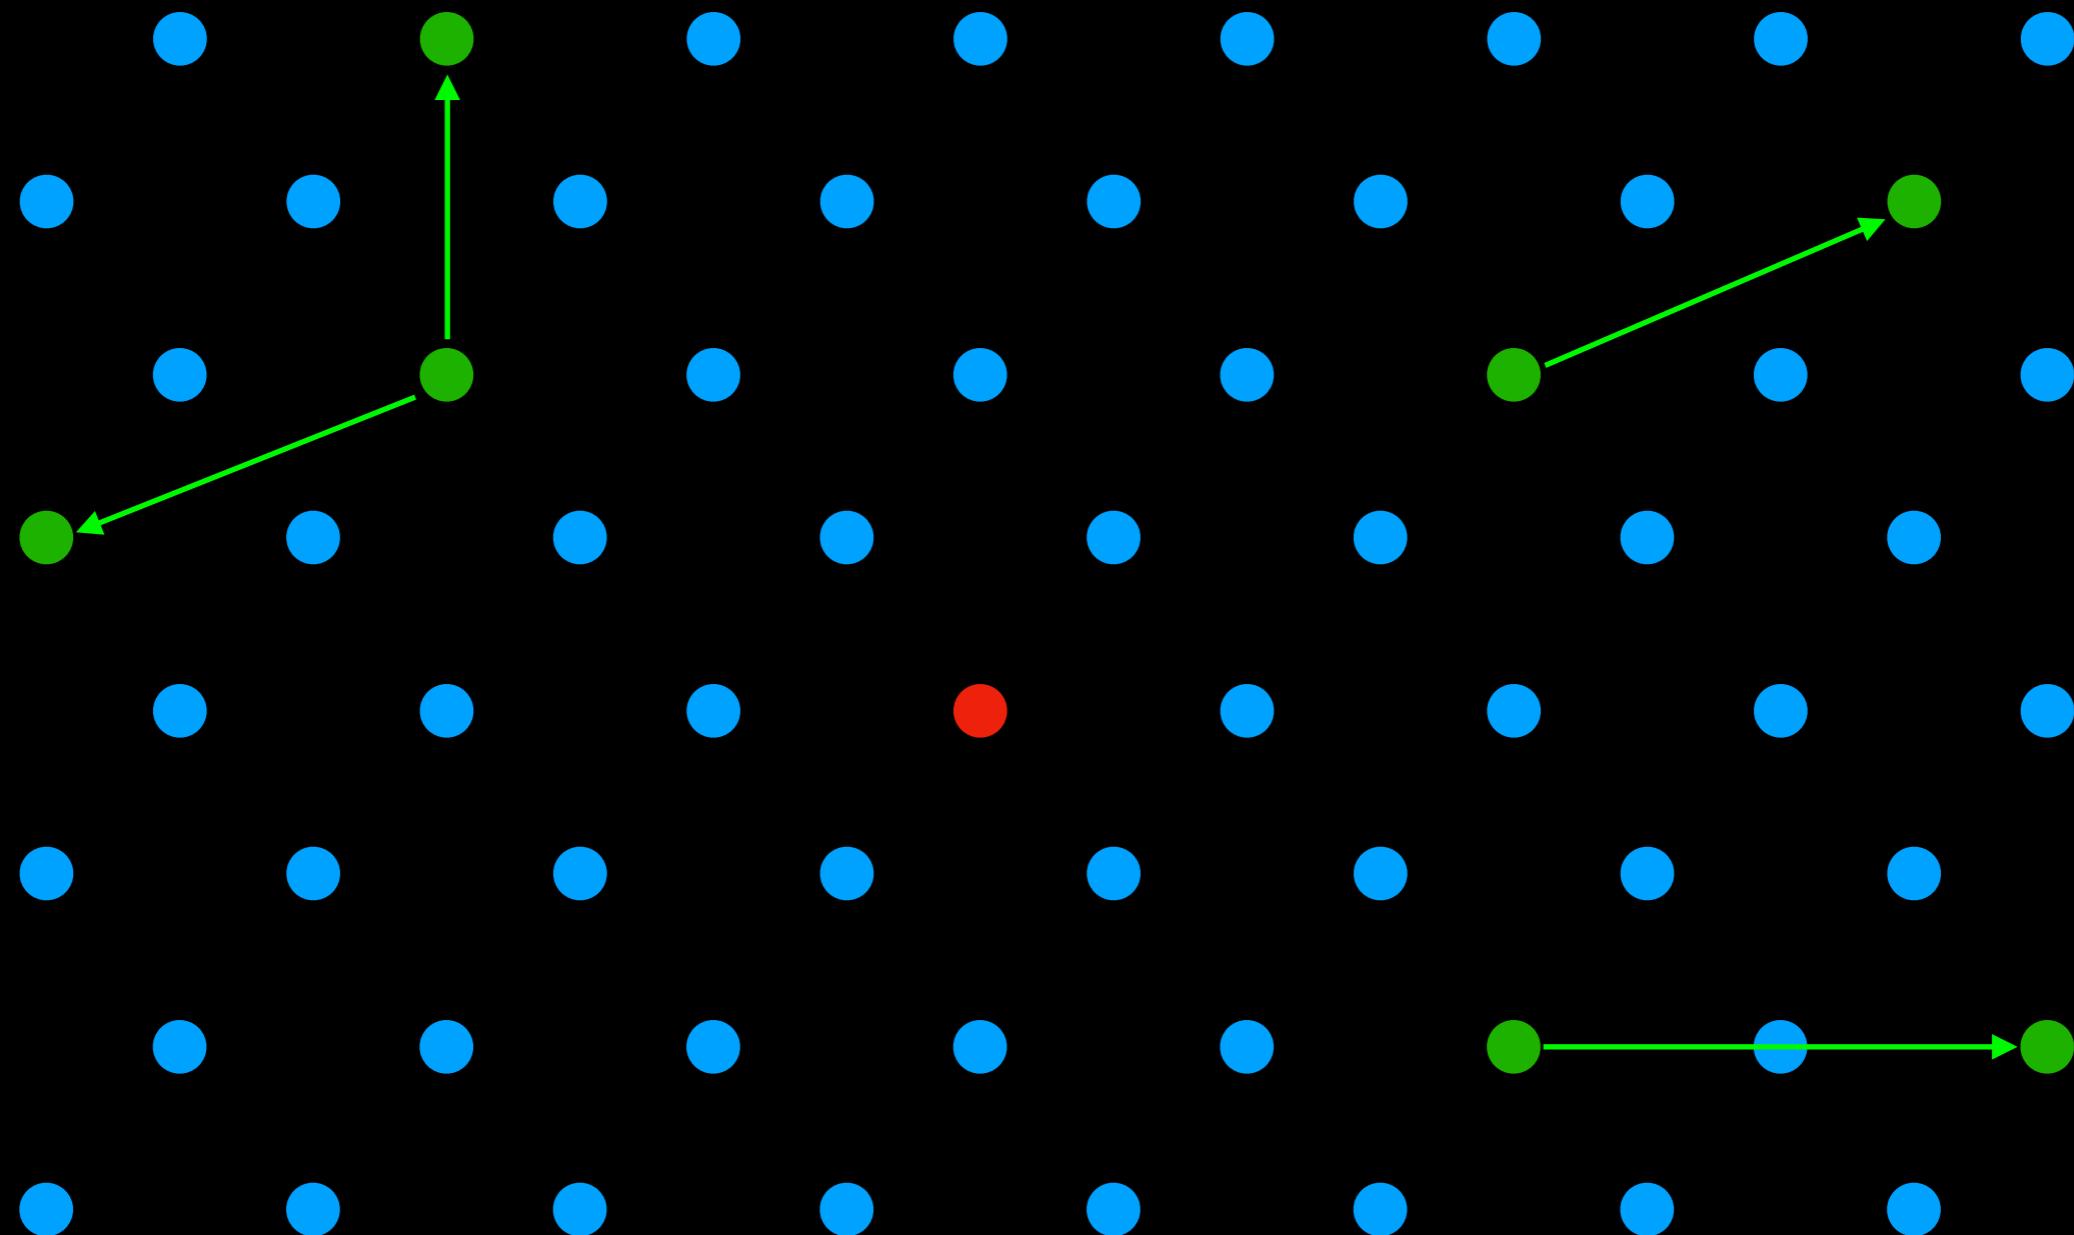
[KMPR19] : Can solve SVP, heuristically, in time $2^{0.1037n+o(n)}$ and (quantum) memory $2^{0.2075n+o(n)}$, in the quantum circuit model.

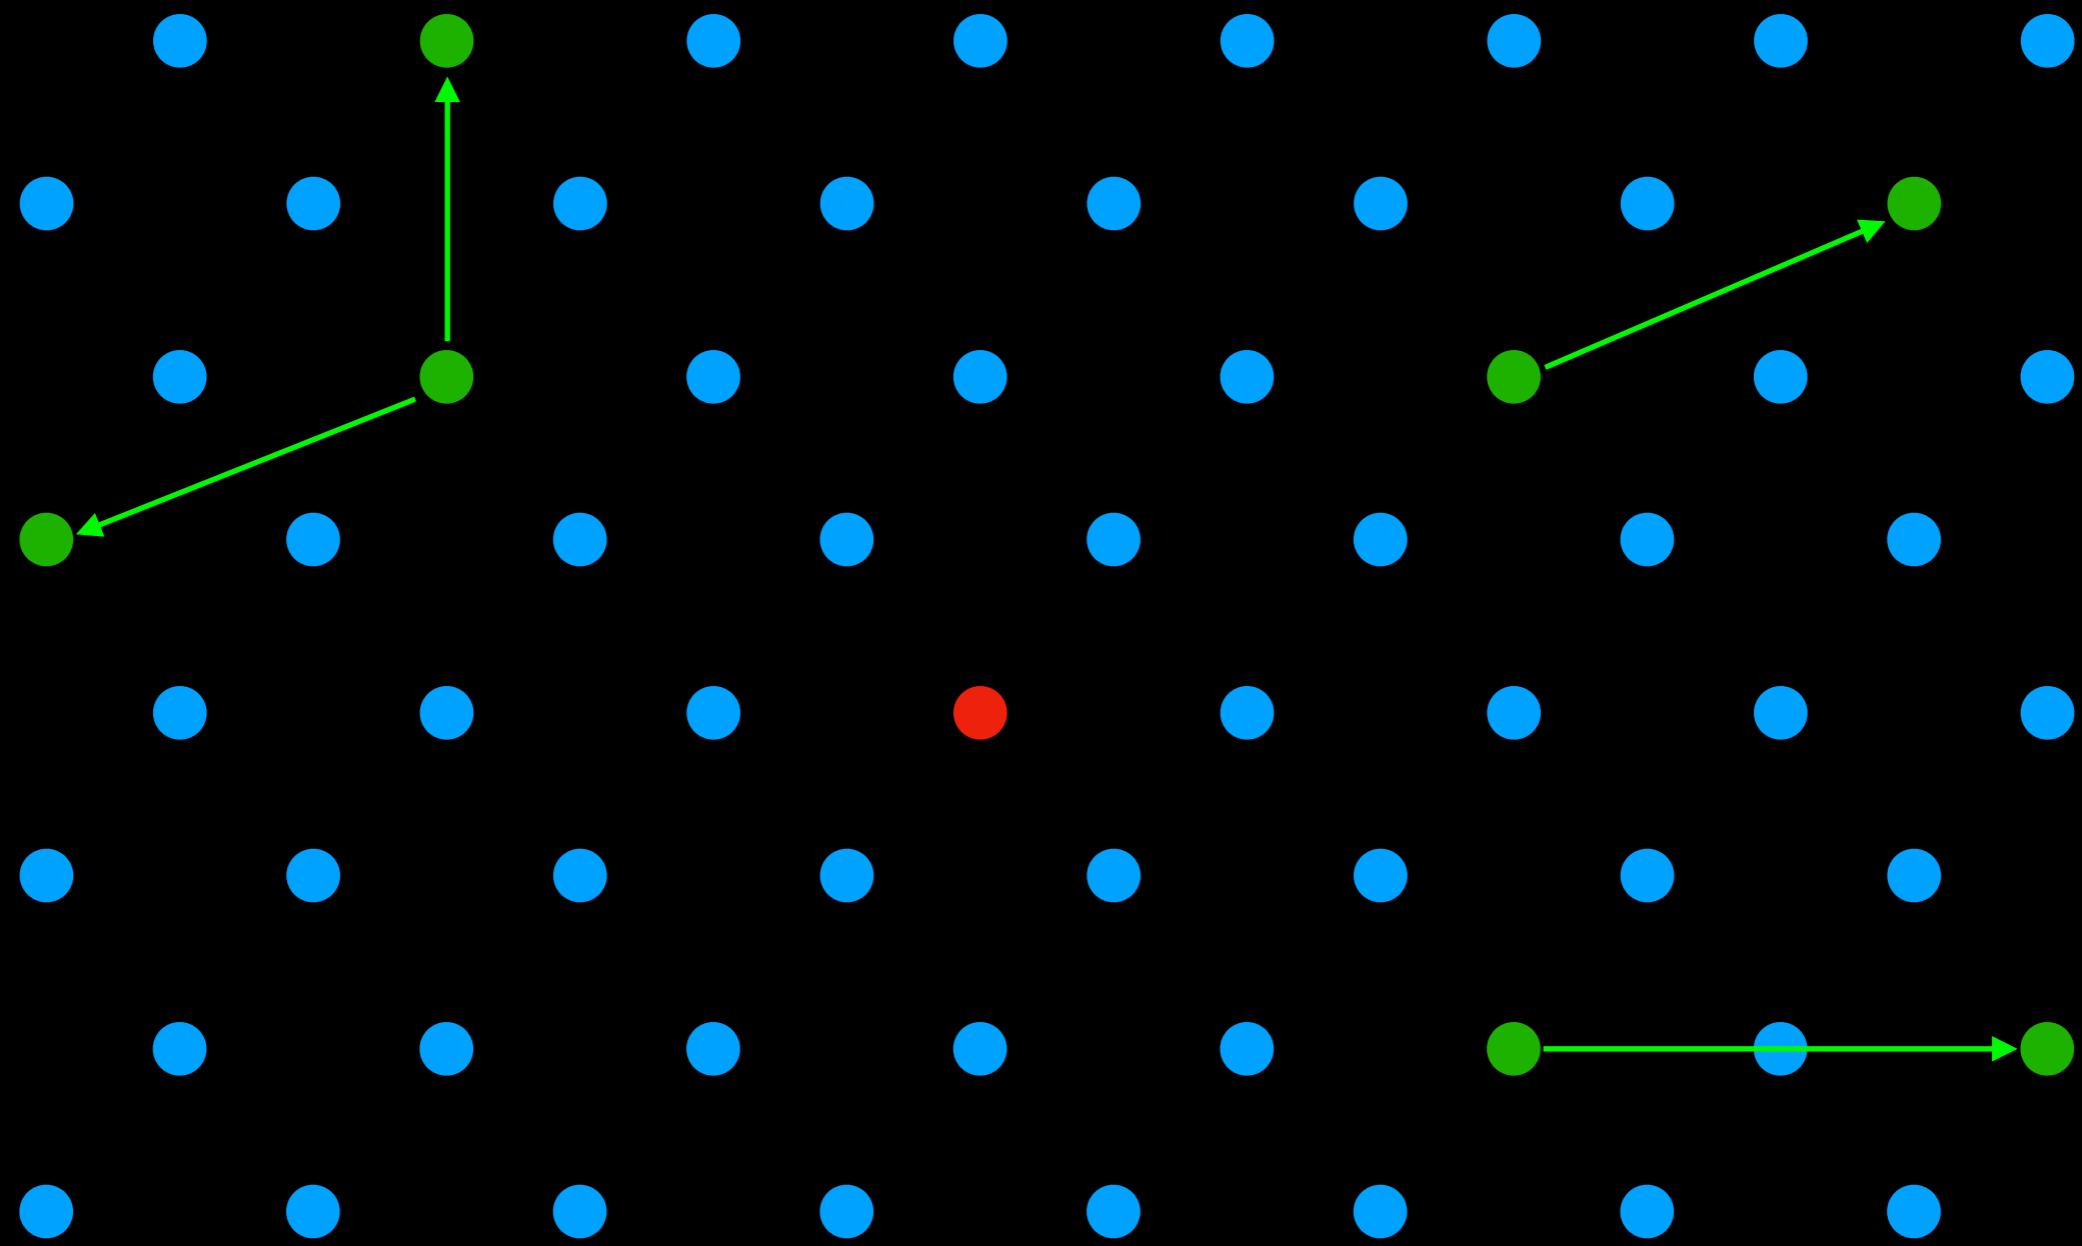
	log2(Time)	log2(Memory)
The Nguyen–Vidick sieve [NV08]	$0.415n$	$0.208n$
The GaussSieve [MV10, ..., IKMT14, BNvdP16, YKYC17]	$0.415n$	$0.208n$
Triple sieve [BLS16, HK17]	$0.396n$	$0.189n$
Two-level sieve [WLTB11]	$0.384n$	$0.256n$
Three-level sieve [ZPH13]	$0.3778n$	$0.283n$
Overlattice sieve [BGJ14]	$0.3774n$	$0.293n$
Triple sieve with NNS [HK17, HKL18]	$0.359n$	0.189n
Hyperplane LSH [Cha02, Laa15, ..., LM18, Duc18]	$0.337n$	$0.337n$
Graph-based NNS [EPY99, DCL11, MPLK14, Laa18]	$0.327n$	$0.282n$
Hypercube LSH [TT07, Laa17]	$0.322n$	$0.322n$
Quantum sieve [LMP13, LMP15]	$0.312n$	$0.208n$
May–Ozerov NNS [MO15, BGJ15]	$0.311n$	$0.311n$
Spherical LSH [AINR14, LdW15]	$0.298n$	$0.298n$
Cross-polytope LSH [TT07, AILRS15, BL16, KW17]	$0.298n$	$0.298n$
Spherical LSF [BDGL16, MLB17, ALRW17, Chr17]	0.292n	$0.292n$
Quantum NNS sieve [LMP15, Laa16]	0.265n	$0.265n$

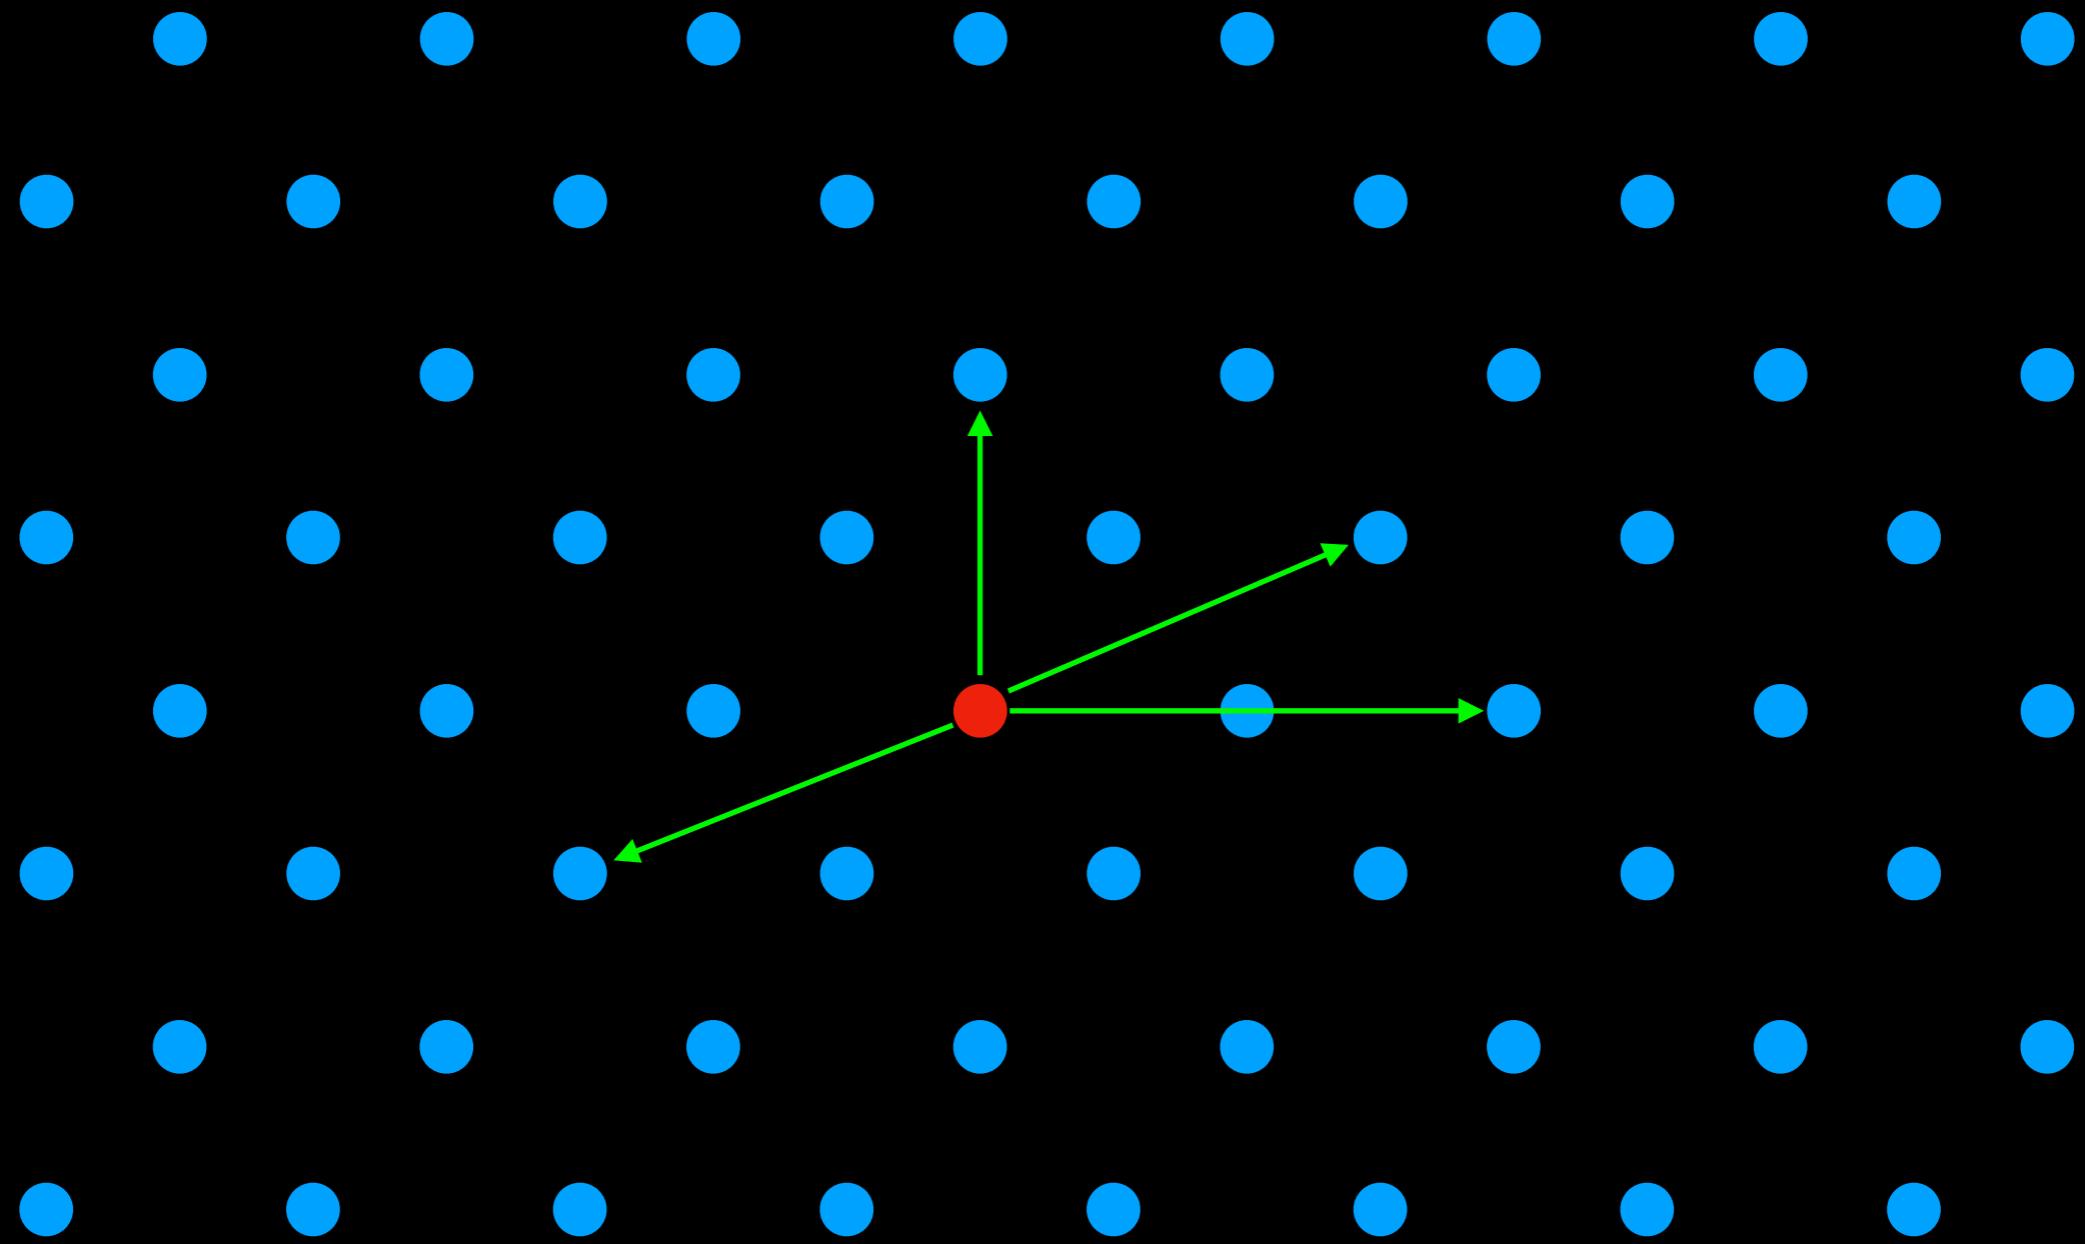
2 - Sieve

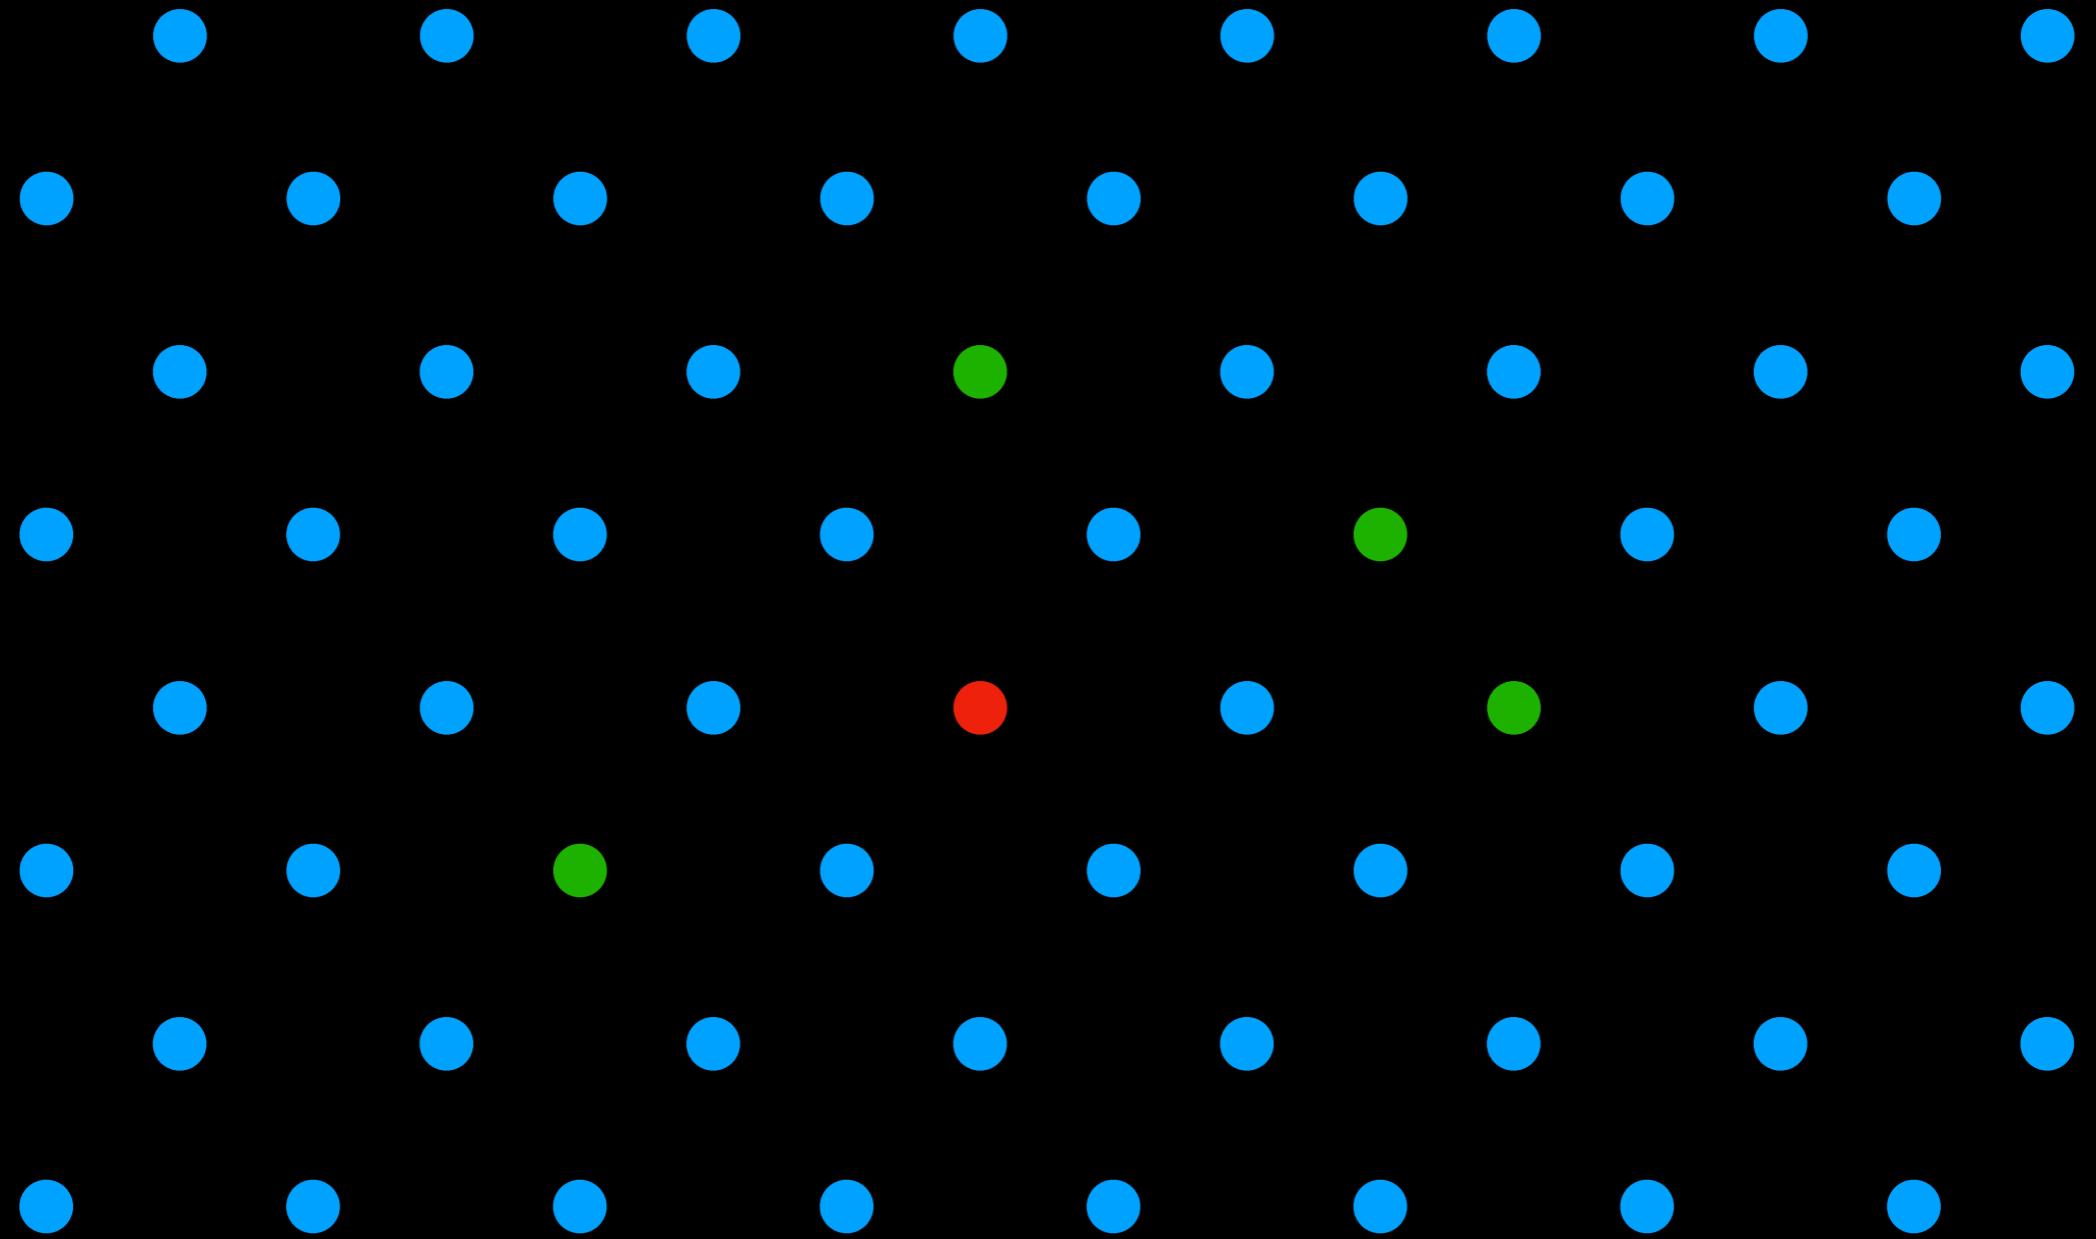


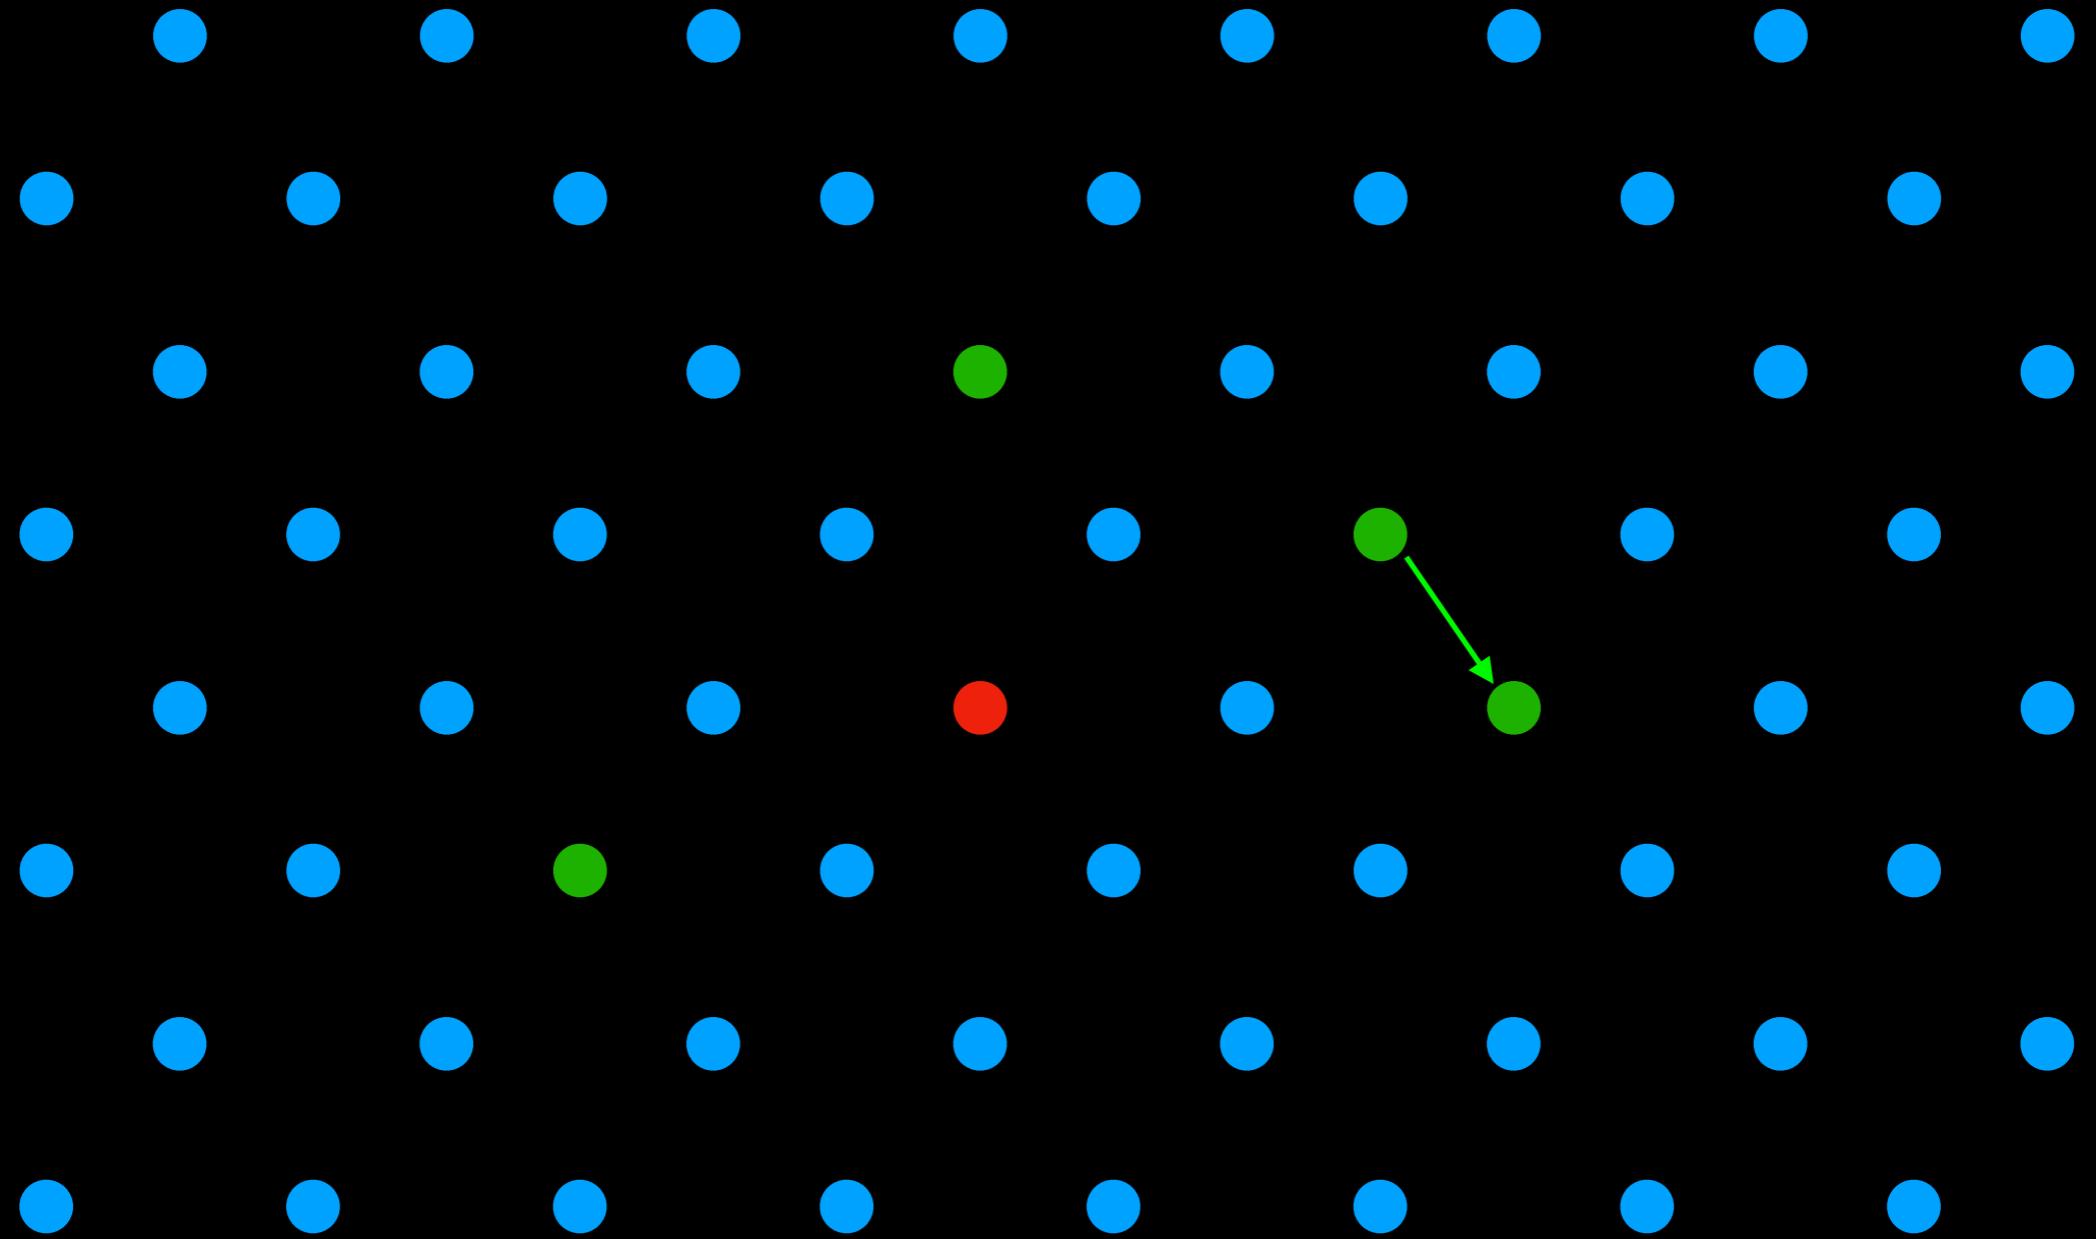


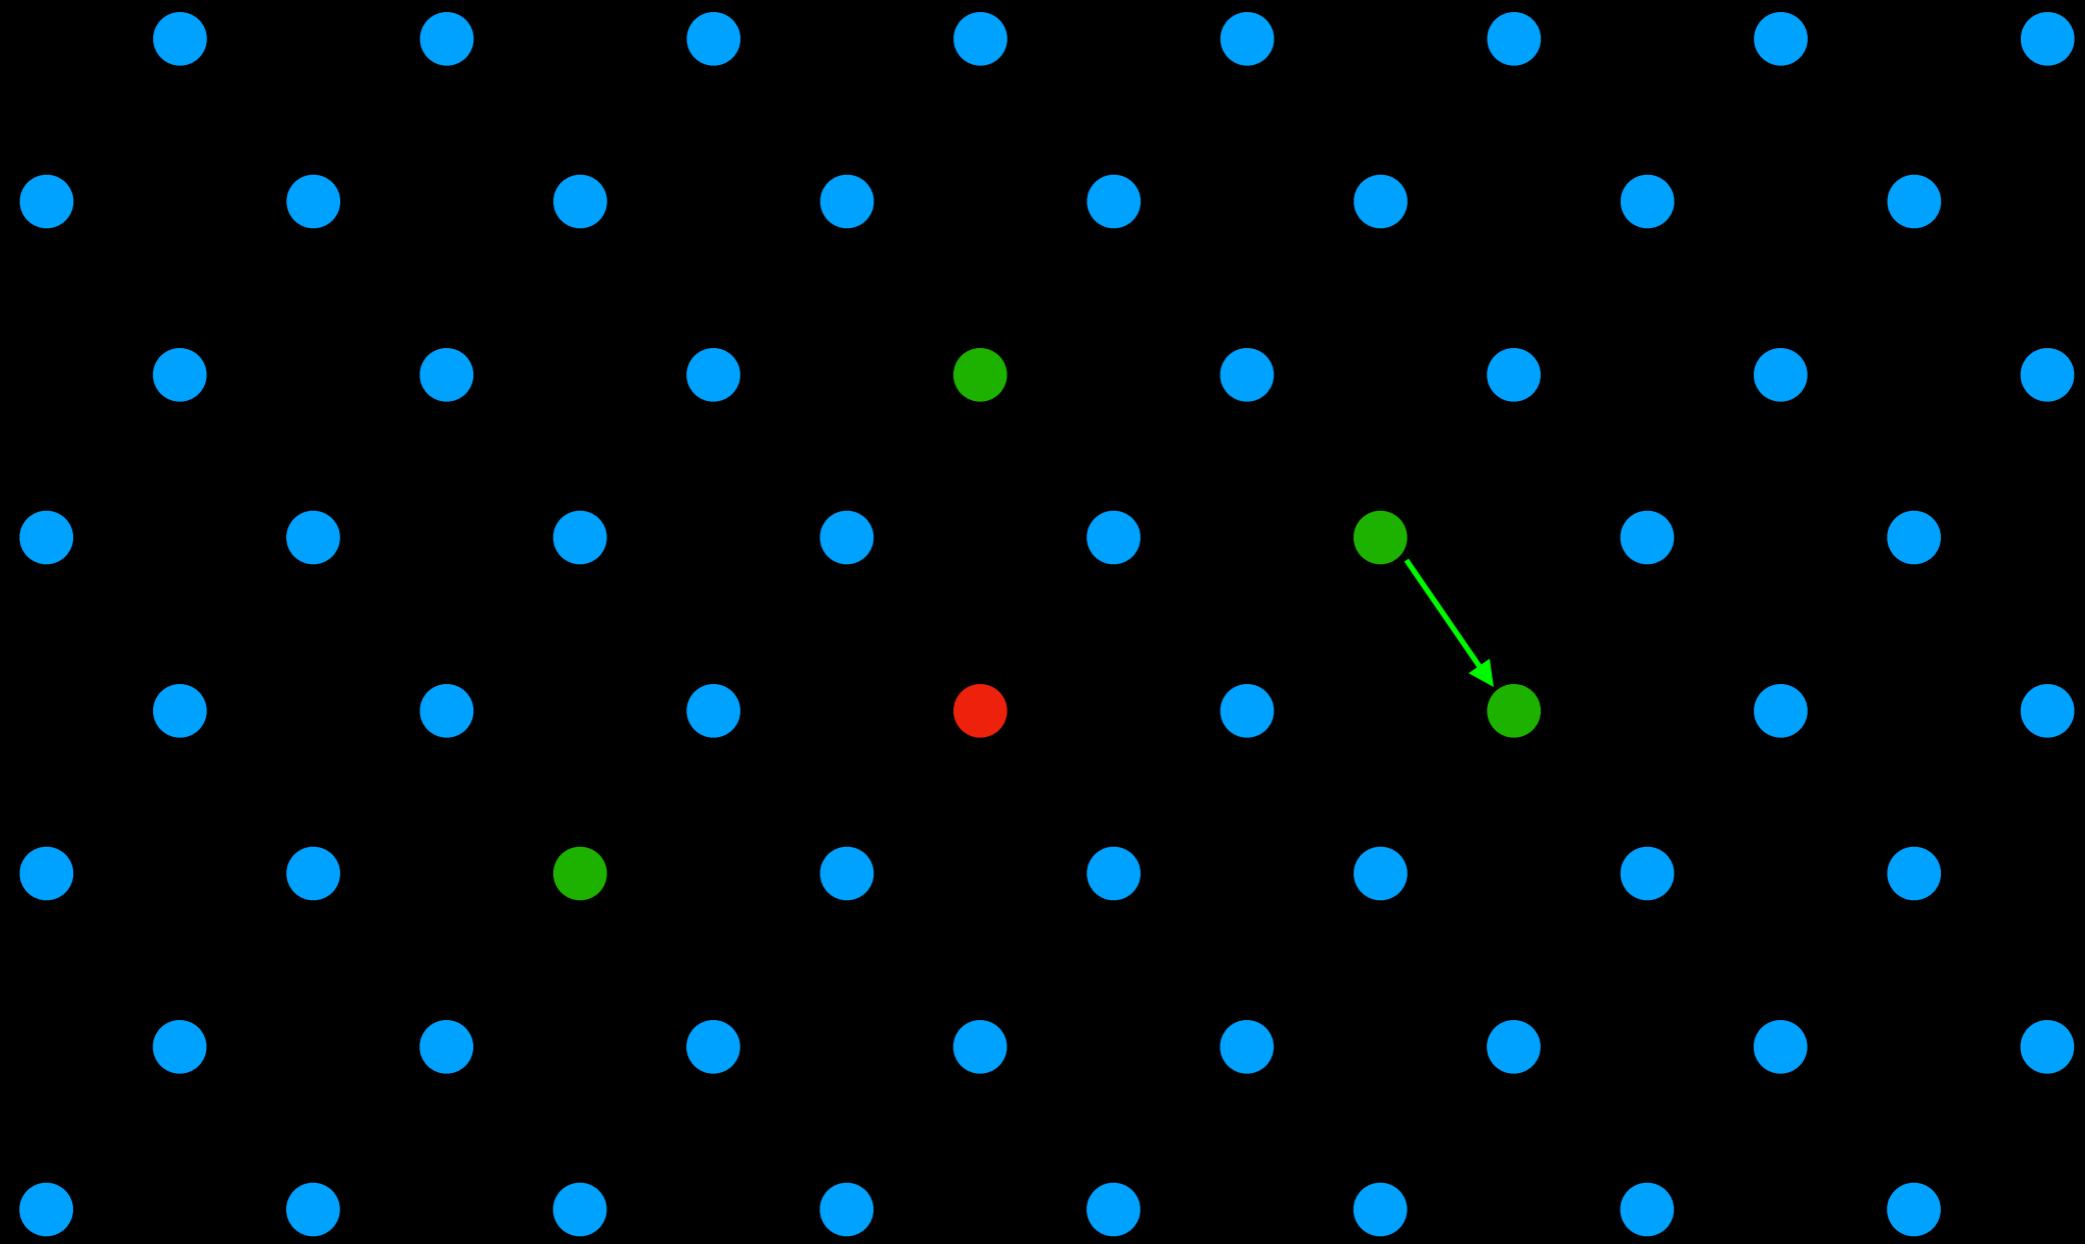


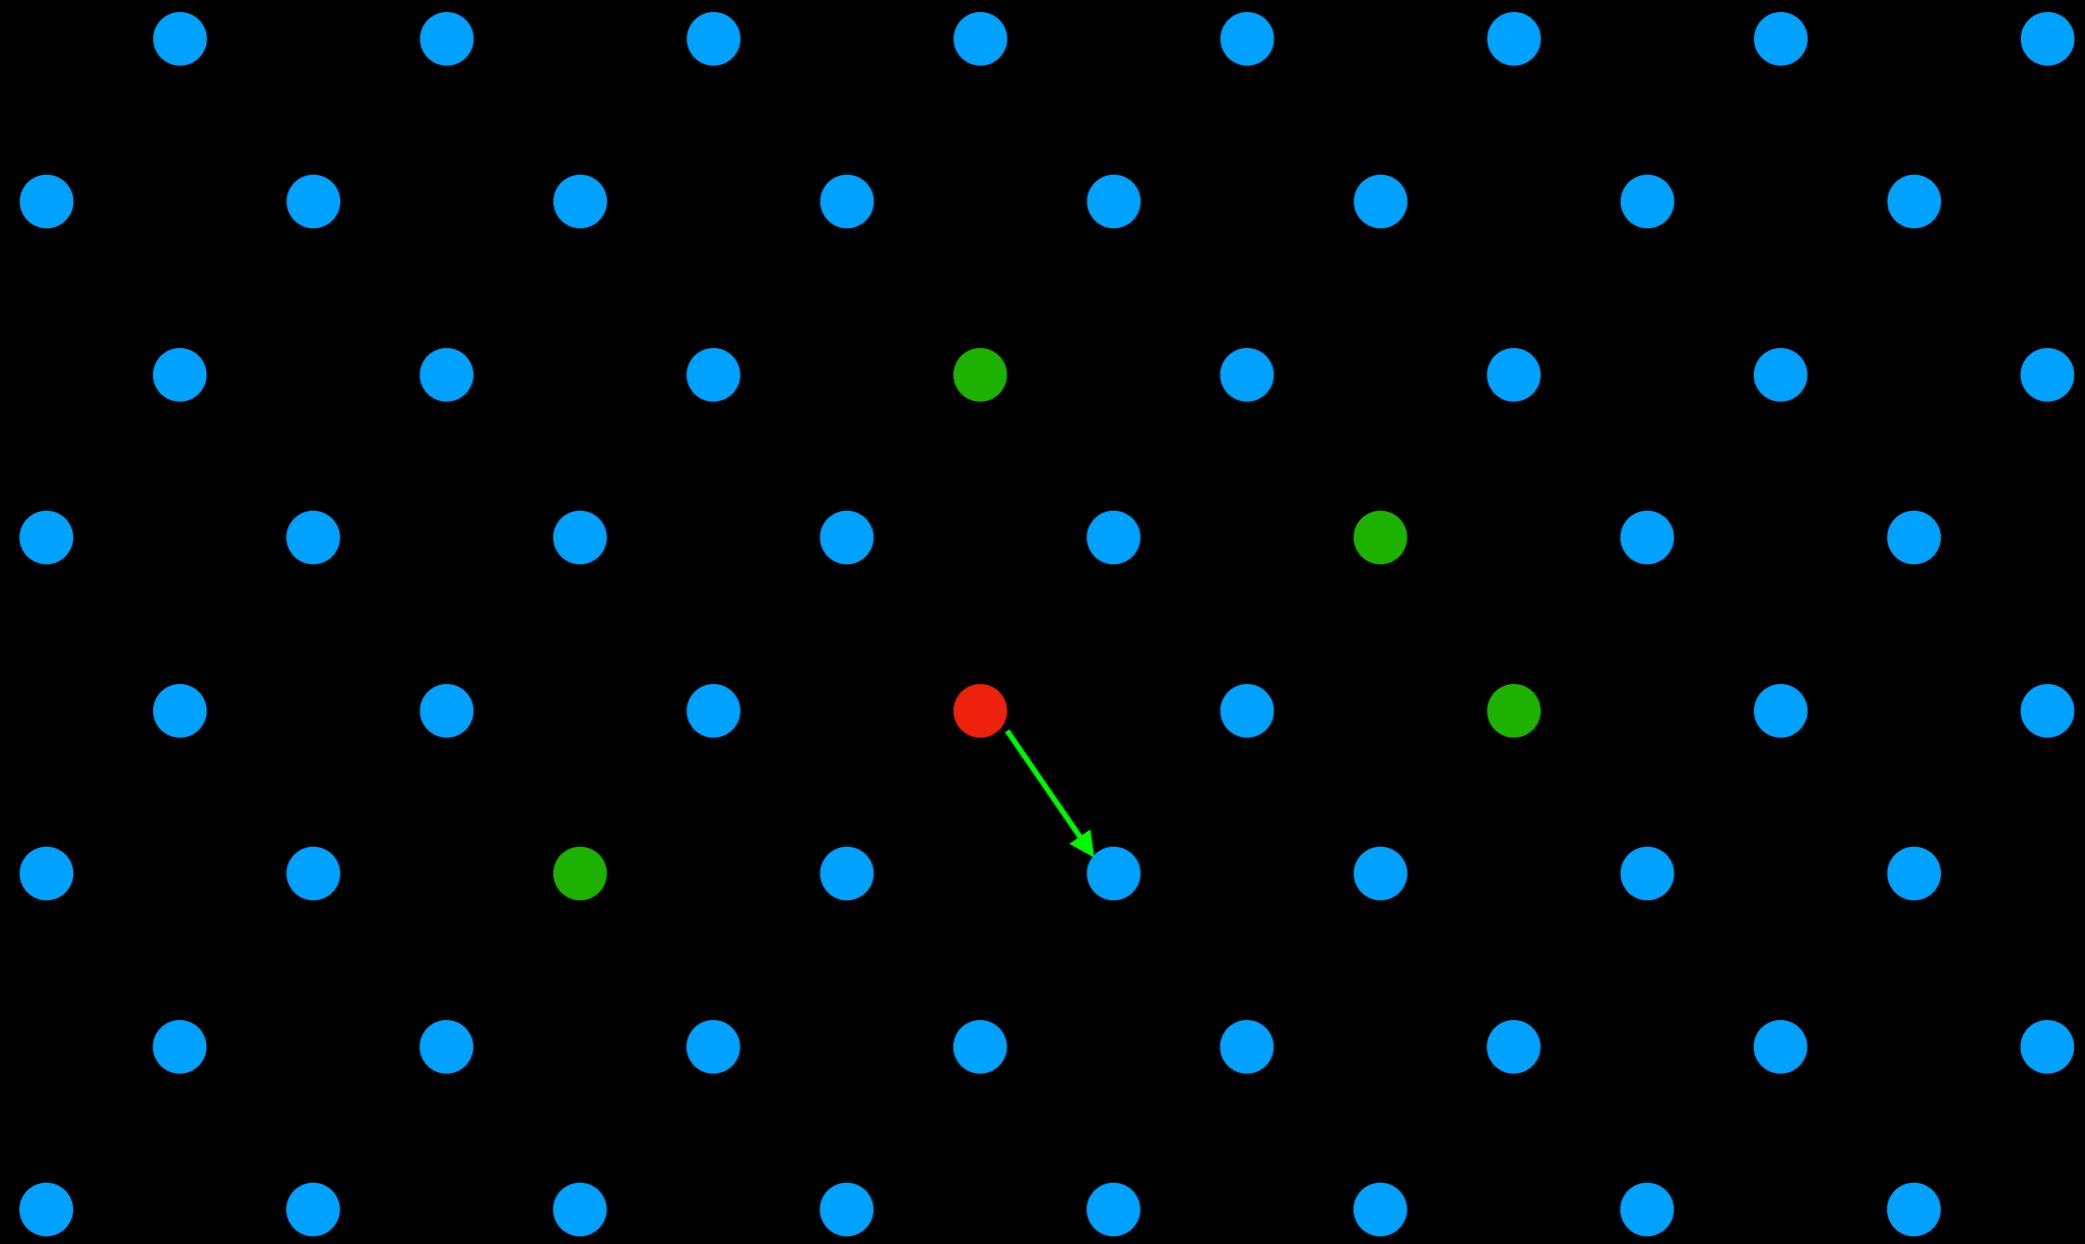


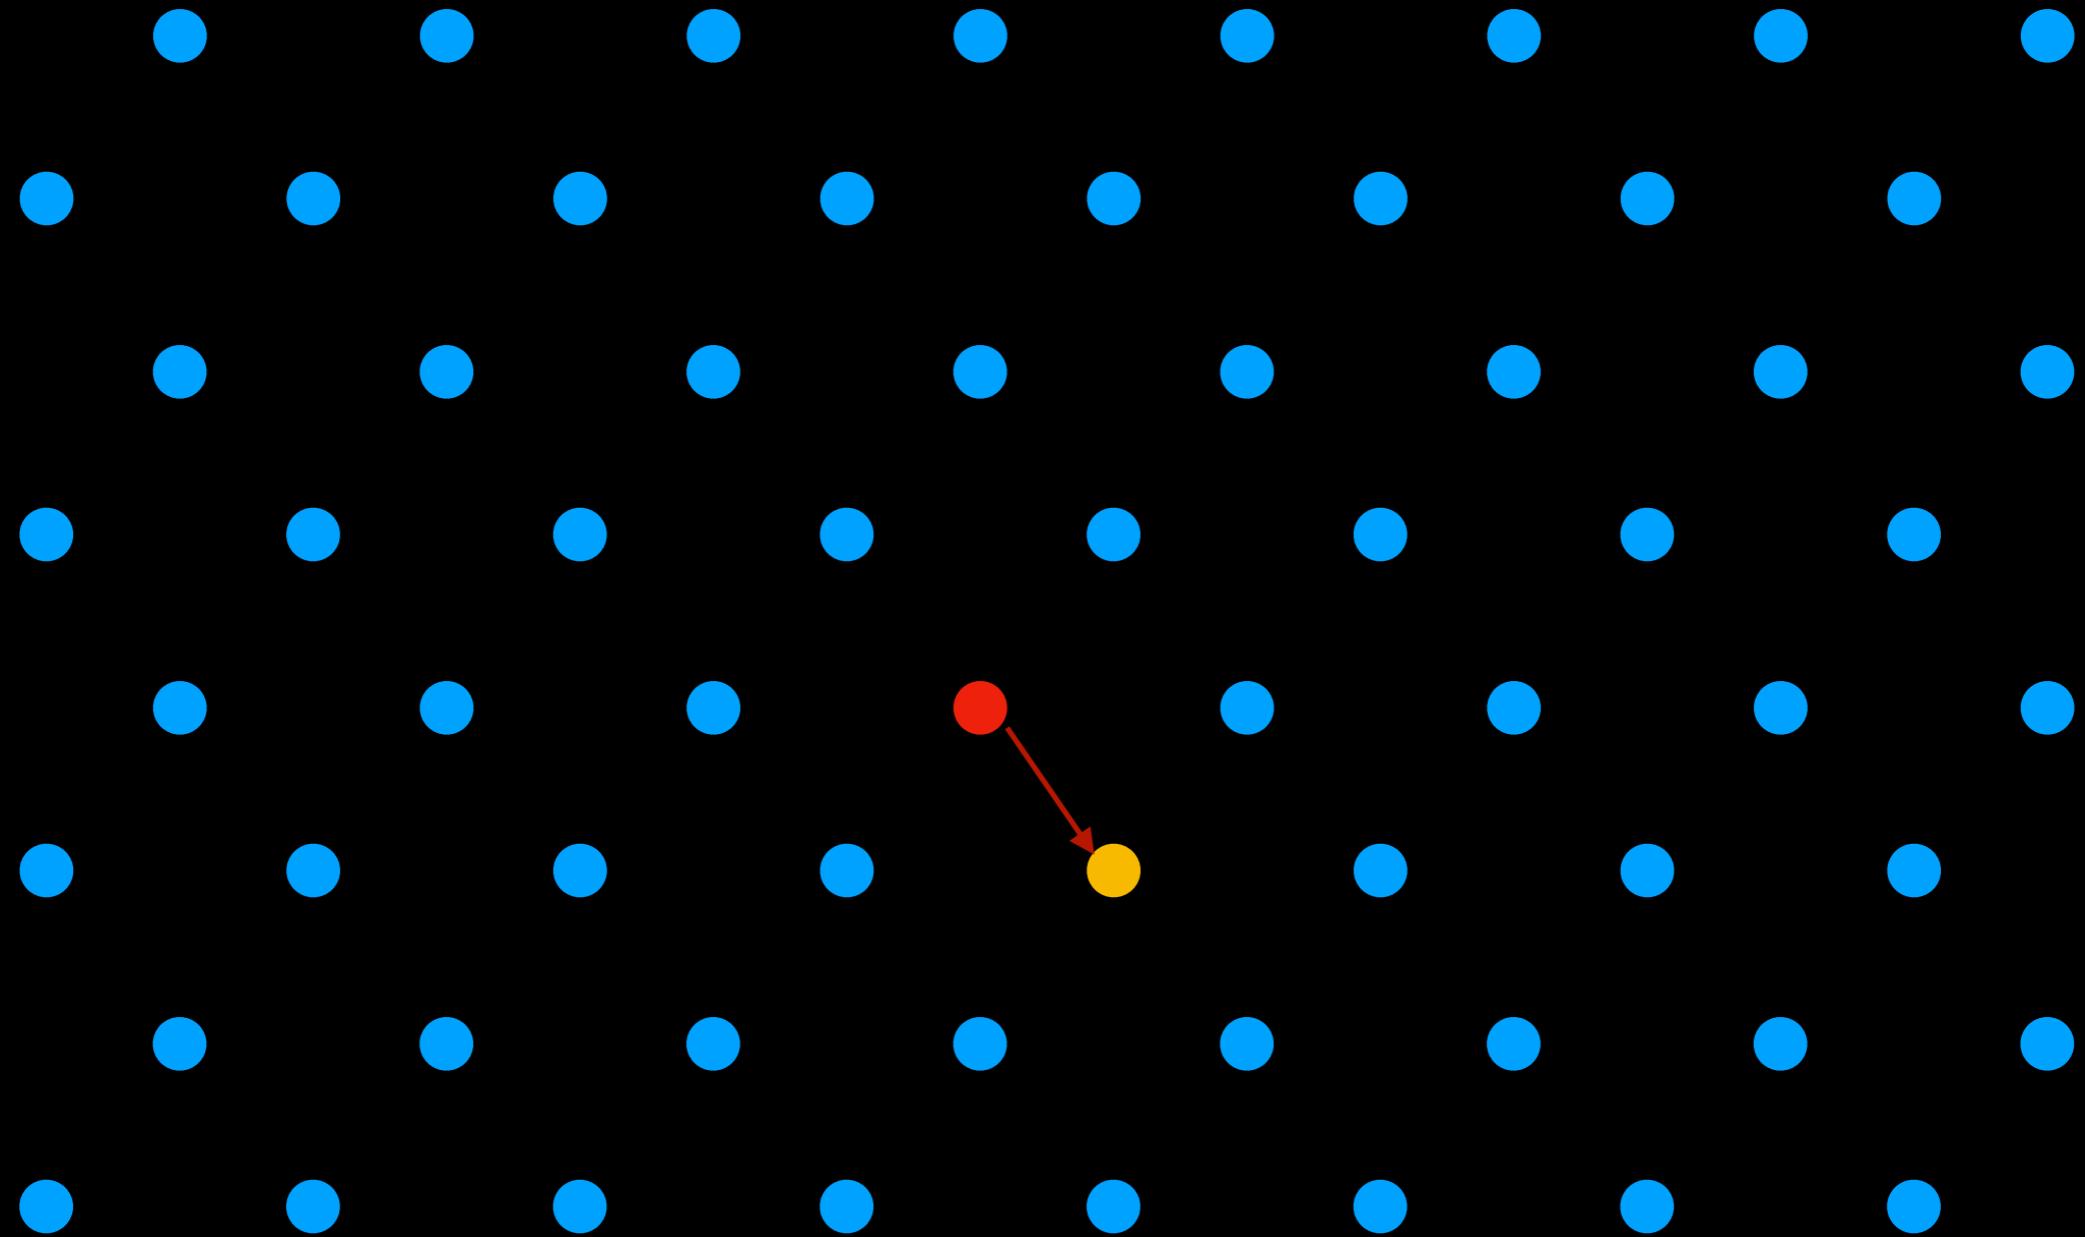












2-Sieve: Algorithm

Input: $\{b_1, \dots, b_d\} \subset R^d$ a lattice basis

Output: $v \in \mathcal{L}(B)$, a short vector from $\mathcal{L}(B)$

1. Set $N \leftarrow 2^{0.2075d+o(d)}$ and set $\lambda = \Theta(\sqrt{d} \cdot \det(B)^{1/d})$ the target length.
2. Generate a list $L_1 \leftarrow \{x_1, \dots, x_N\}$ of normalised lattice vectors using an efficient lattice sampling procedure.
3. Construct a list $L_2 \leftarrow \{x'_1, \dots, x'_N\}$ such that $| \langle x_i, x'_i \rangle | \geq 1/2$ for $x'_i \in L_1$. If no such $x'_i \in L_1$ exists, set $x'_i = \mathbf{0}$.
4. Construct a list $L_3 \leftarrow \{y_i : y_i \leftarrow \min\{||x_i \pm x'_i||\}$ for all $i \leq N\}$ and normalise its elements except for the last iteration.
5. Swap the labels L_1, L_3 . Reinitialise L_2 and L_3 to the zero state by transferring their contents to auxiliary memory.
6. Repeat Steps 3-5 $\text{poly}(d)$ times.
7. Output a vector from L_1 of Euclidean norm less than λ .

2-Sieve: Algorithm

Input: $\{b_1, \dots, b_d\} \subset R^d$ a lattice basis

Output: $v \in \mathcal{L}(B)$, a short vector from $\mathcal{L}(B)$

1. Set $N \leftarrow 2^{0.2075d+o(d)}$ and set $\lambda = \Theta(\sqrt{d} \cdot \det(B)^{1/d})$ the target length.

L_1

2. Generate a list $L_1 \leftarrow \{x_1, \dots, x_N\}$ of normalised lattice vectors using an efficient lattice sampling procedure.

x_1

x_2

x_3

\vdots

x_N

3. Construct a list $L_2 \leftarrow \{x'_1, \dots, x'_N\}$ such that $| \langle x_i, x'_i \rangle | \geq 1/2$ for $x'_i \in L_1$. If no such $x'_i \in L_1$ exists, set $x'_i = \mathbf{0}$.

4. Construct a list $L_3 \leftarrow \{y_i : y_i \leftarrow \min\{| |x_i \pm x'_i| |\}$ for all $i \leq N\}$ and normalise its elements except for the last iteration.

5. Swap the labels L_1, L_3 . Reinitialise L_2 and L_3 to the zero state by transferring their contents to auxiliary memory.

6. Repeat Steps 3-5 $poly(d)$ times.

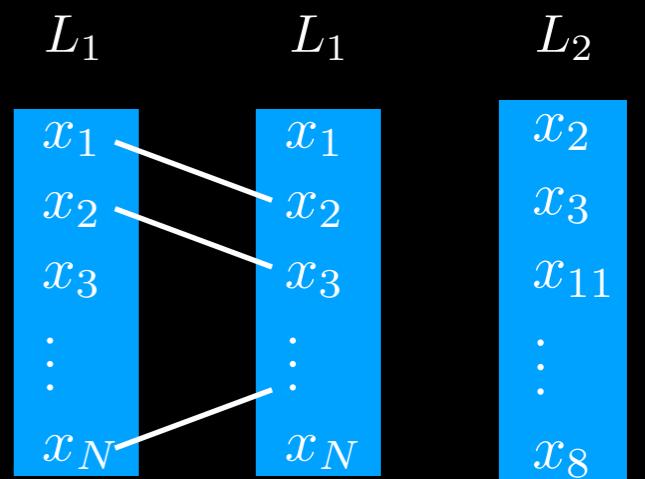
7. Output a vector from L_1 of Euclidean norm less than λ .

2-Sieve: Algorithm

Input: $\{b_1, \dots, b_d\} \subset R^d$ a lattice basis

Output: $v \in \mathcal{L}(B)$, a short vector from $\mathcal{L}(B)$

1. Set $N \leftarrow 2^{0.2075d+o(d)}$ and set $\lambda = \Theta(\sqrt{d} \cdot \det(B)^{1/d})$ the target length.
2. Generate a list $L_1 \leftarrow \{x_1, \dots, x_N\}$ of normalised lattice vectors using an efficient lattice sampling procedure.
3. Construct a list $L_2 \leftarrow \{x'_1, \dots, x'_N\}$ such that $| \langle x_i, x'_i \rangle | \geq 1/2$ for $x'_i \in L_1$. If no such $x'_i \in L_1$ exists, set $x'_i = \mathbf{0}$.
4. Construct a list $L_3 \leftarrow \{y_i : y_i \leftarrow \min\{\|x_i \pm x'_i\|\}\}$ for all $i \leq N\}$ and normalise its elements except for the last iteration.
5. Swap the labels L_1, L_3 . Reinitialise L_2 and L_3 to the zero state by transferring their contents to auxiliary memory.
6. Repeat Steps 3-5 $\text{poly}(d)$ times.
7. Output a vector from L_1 of Euclidean norm less than λ .



2-Sieve: Algorithm

Input: $\{b_1, \dots, b_d\} \subset R^d$ a lattice basis

Output: $v \in \mathcal{L}(B)$, a short vector from $\mathcal{L}(B)$

1. Set $N \leftarrow 2^{0.2075d+o(d)}$ and set $\lambda = \Theta(\sqrt{d} \cdot \det(B)^{1/d})$ the target length.

2. Generate a list $L_1 \leftarrow \{x_1, \dots, x_N\}$ of normalised lattice vectors using an efficient lattice sampling procedure.

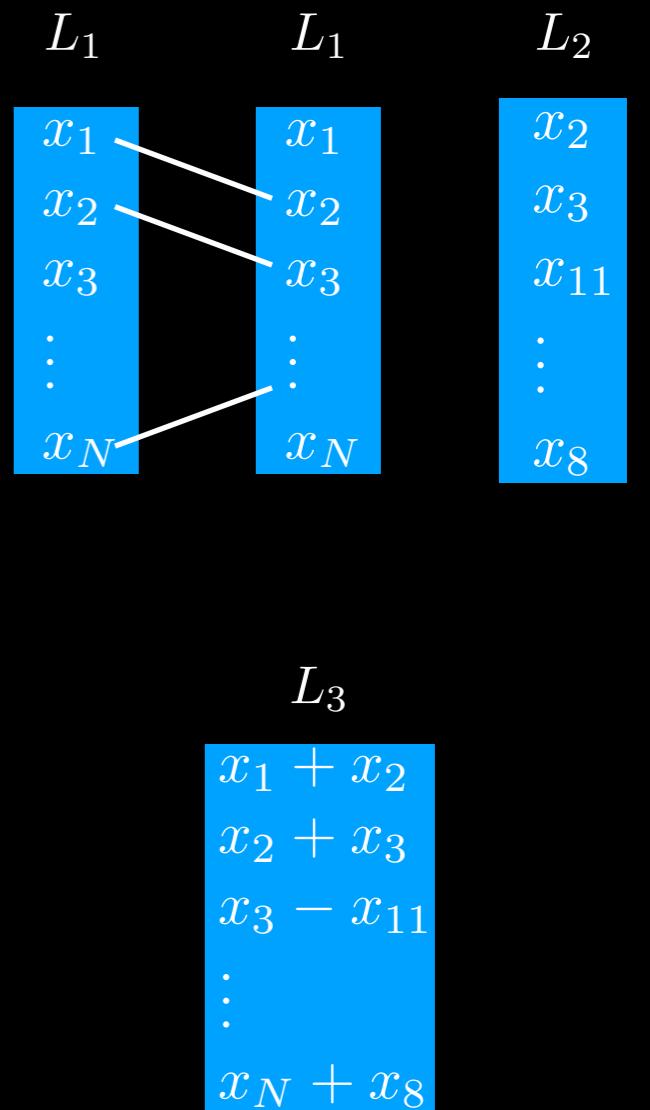
3. Construct a list $L_2 \leftarrow \{x'_1, \dots, x'_N\}$ such that $| \langle x_i, x'_i \rangle | \geq 1/2$ for $x'_i \in L_1$. If no such $x'_i \in L_1$ exists, set $x'_i = \mathbf{0}$.

4. Construct a list $L_3 \leftarrow \{y_i : y_i \leftarrow \min\{| |x_i \pm x'_i| |\}$ for all $i \leq N\}$ and normalise its elements except for the last iteration.

5. Swap the labels L_1, L_3 . Reinitialise L_2 and L_3 to the zero state by transferring their contents to auxiliary memory.

6. Repeat Steps 3-5 $\text{poly}(d)$ times.

7. Output a vector from L_1 of Euclidean norm less than λ .



2-Sieve: Algorithm

Input: $\{b_1, \dots, b_d\} \subset R^d$ a lattice basis

Output: $v \in \mathcal{L}(B)$, a short vector from $\mathcal{L}(B)$

1. Set $N \leftarrow 2^{0.2075d+o(d)}$ and set $\lambda = \Theta(\sqrt{d} \cdot \det(B)^{1/d})$ the target length.
2. Generate a list $L_1 \leftarrow \{x_1, \dots, x_N\}$ of normalised lattice vectors using an efficient lattice sampling procedure.
3. Construct a list $L_2 \leftarrow \{x'_1, \dots, x'_N\}$ such that $| \langle x_i, x'_i \rangle | \geq 1/2$ for $x'_i \in L_1$. If no such $x'_i \in L_1$ exists, set $x'_i = \mathbf{0}$.
4. Construct a list $L_3 \leftarrow \{y_i : y_i \leftarrow \min\{||x_i \pm x'_i||\} \text{ for all } i \leq N\}$ and normalise its elements except for the last iteration.
5. Swap the labels L_1, L_3 . Reinitialise L_2 and L_3 to the zero state by transferring their contents to auxiliary memory.
$$\begin{array}{l} L_1 \\ x_1 + x_2 \\ x_2 + x_3 \\ x_3 - x_{11} \\ \vdots \\ x_N + x_8 \end{array}$$
6. Repeat Steps 3-5 $\text{poly}(d)$ times.
7. Output a vector from L_1 of Euclidean norm less than λ .

L_1	L_1	L_2	L_3
x_1	x_1	x_2	$x_1 + x_2$
x_2	x_2	x_3	$x_2 + x_3$
x_3	x_3	x_{11}	$x_3 - x_{11}$
\vdots	\vdots	\vdots	\vdots
x_N	x_N	x_8	$x_N + x_8$

Idea 1: Pick a vector and search the whole list



Idea 1: Pick a vector and search the whole list

Idea 2: Search the list simultaneously

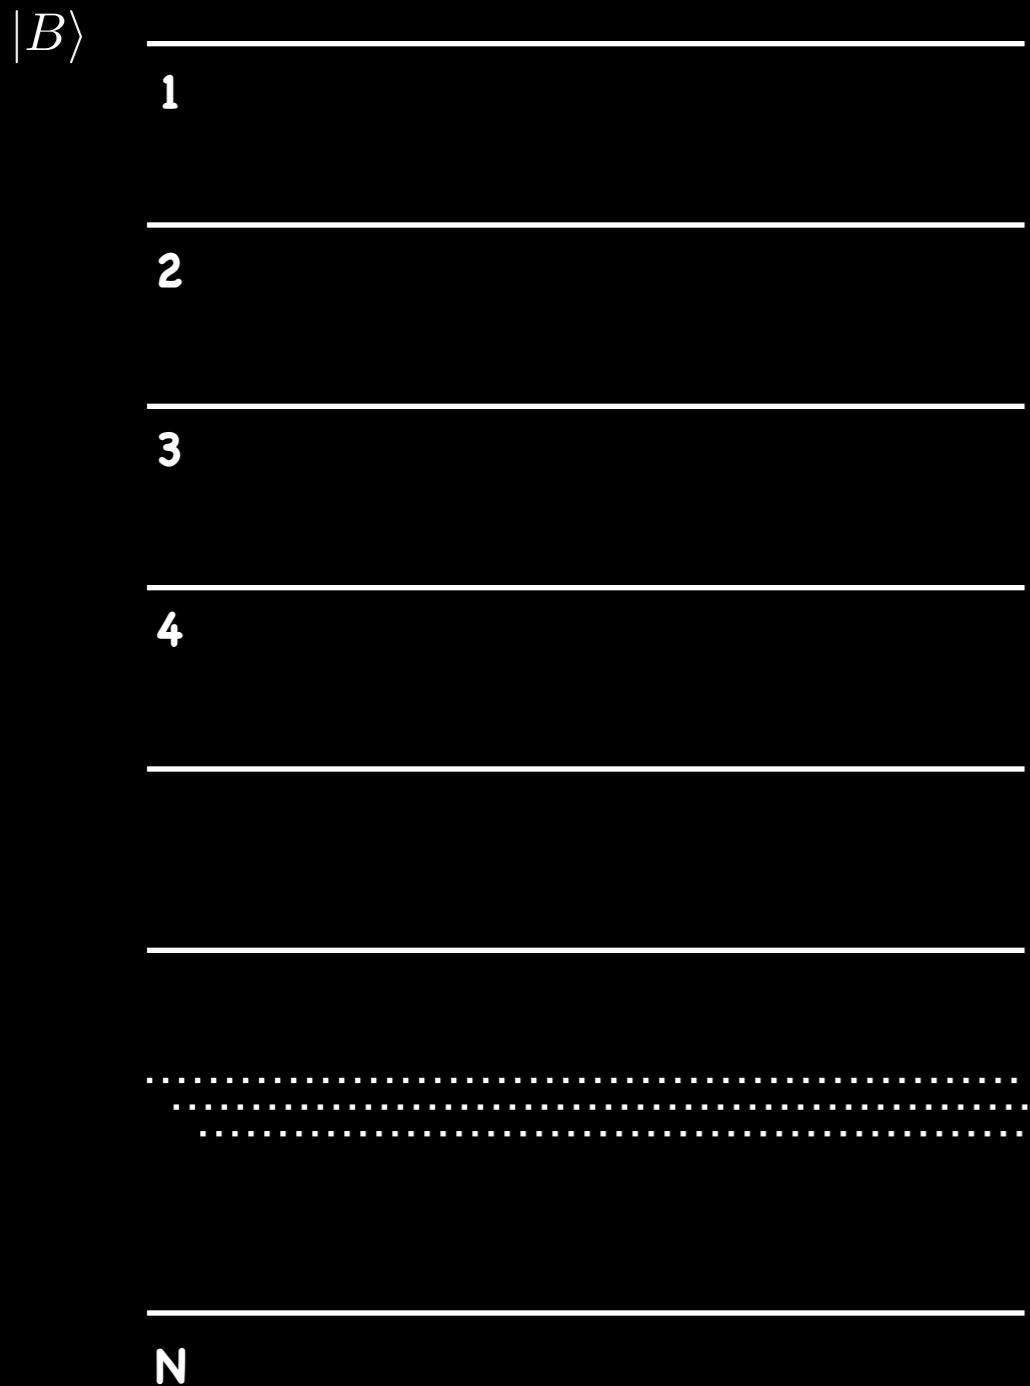
Given a list L of size N , with each element of the list being d bits, and N functions, $f_i : L \rightarrow \{0, 1\}$ for $i = [1\dots N]$, we wish to find solution vectors $s_i \in [1\dots N]$.

That is we wish to find s_i for a f_i such that $f_i(x_{s_i}) = 1$ for all $i \in [1\dots N]$.

Given unitaries $U_{f_i} : |x\rangle|b\rangle \rightarrow |x\rangle|b \oplus f_i(x)\rangle$ there exists a quantum algorithm that, for each $i \in [1\dots N]$, either returns a solution s_i or if there is no such solution, returns *no solution*.

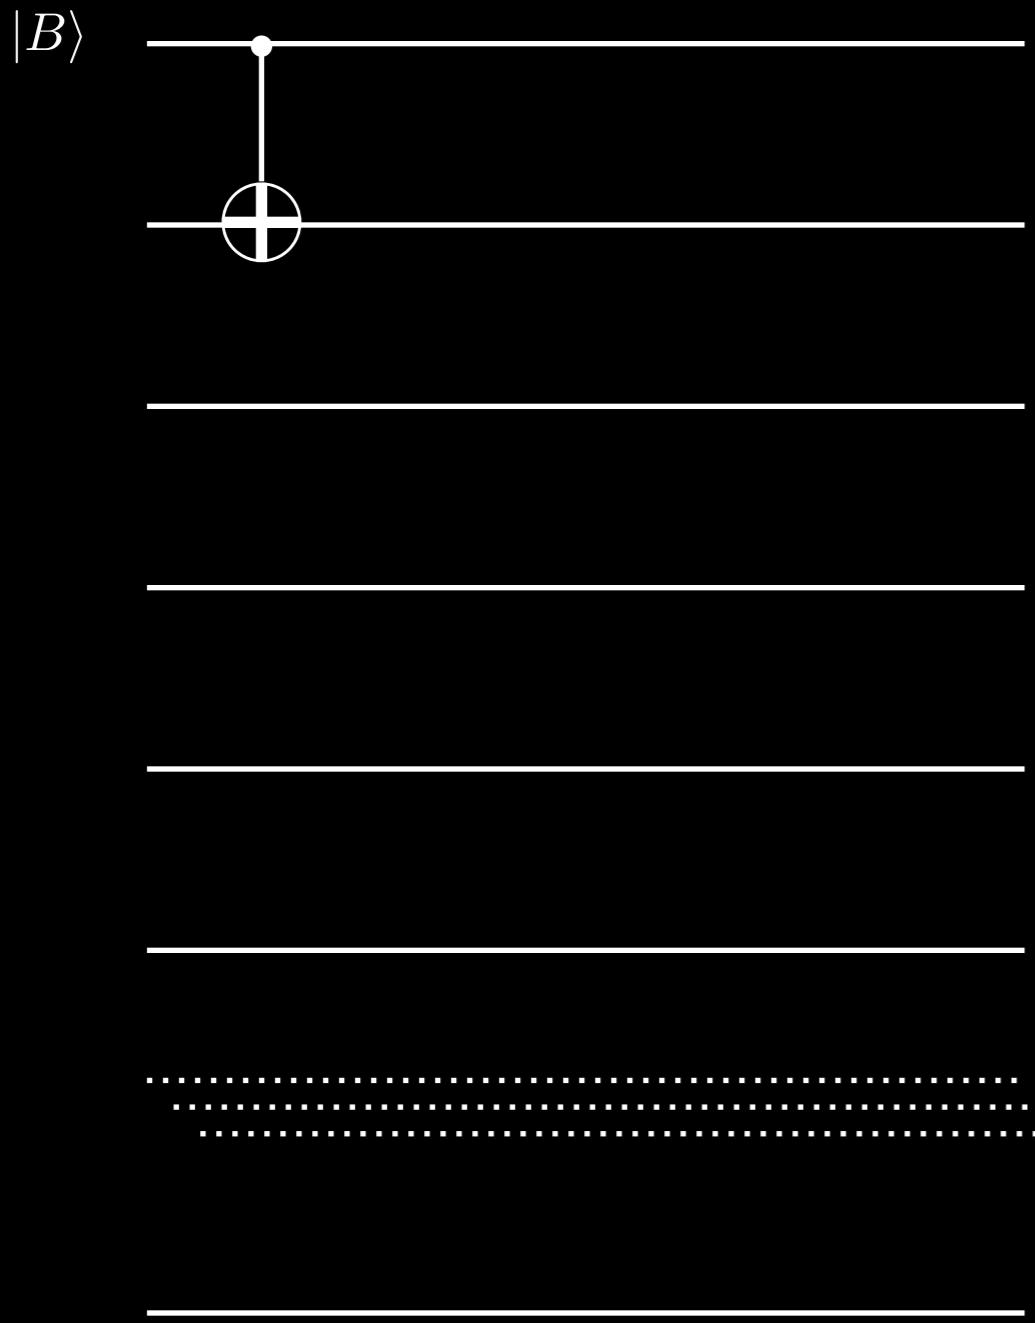
The algorithm succeeds with probability $\Theta(1)$ and, given that each U_{f_i} has depth and width $\text{polylog}(N, d)$, this can be implemented using a quantum circuit of width $\tilde{O}(N)$ and depth $\tilde{O}(\sqrt{N})$.

A quantum algorithm for parallel sieving



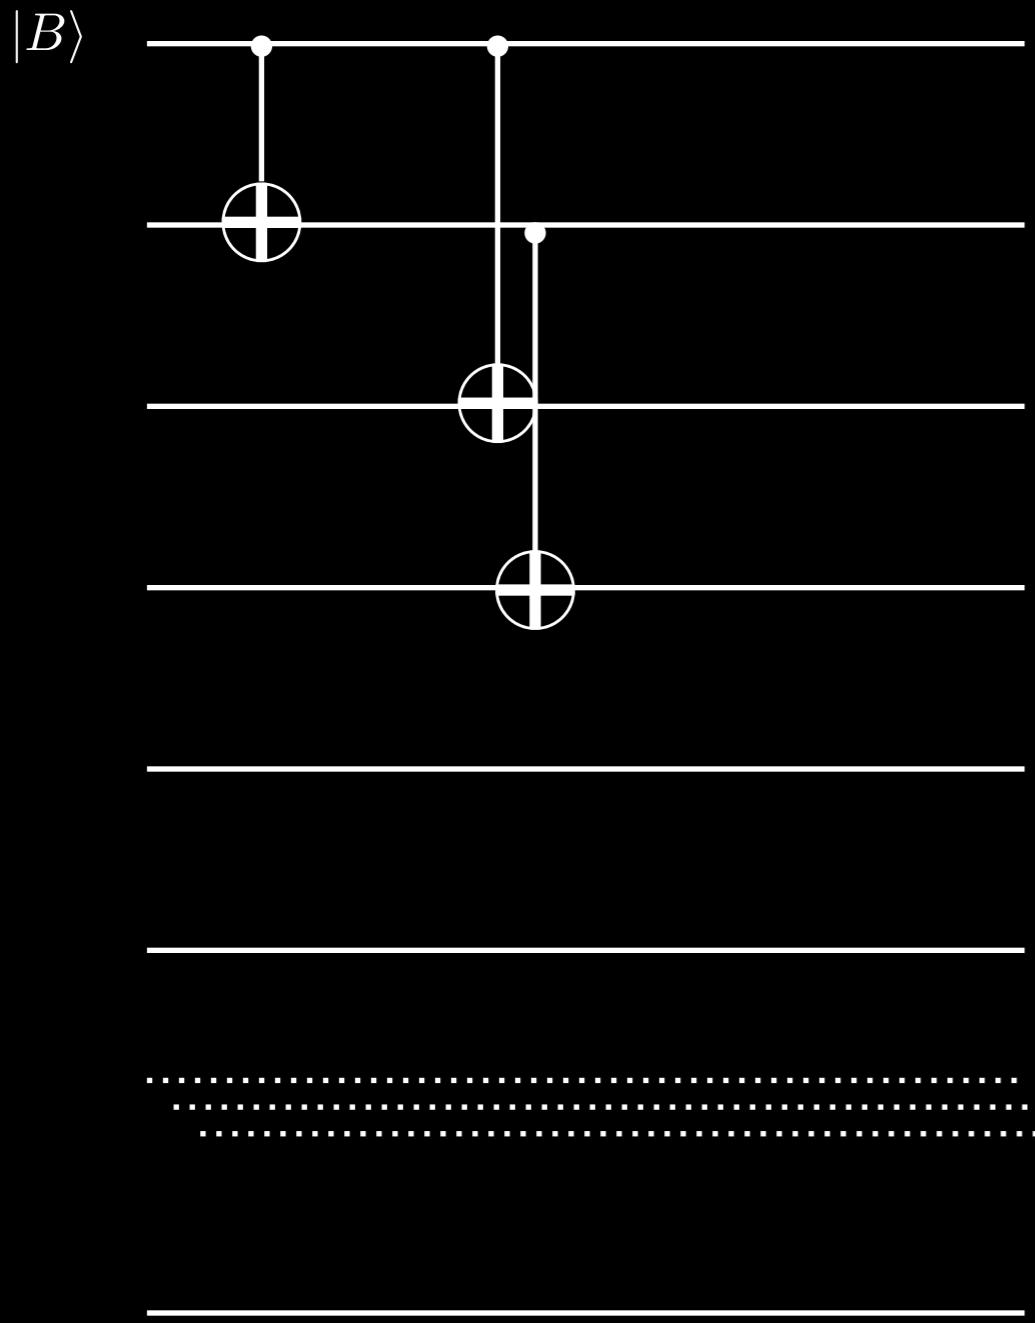
1. Start with $N^{*\text{poly}(d)}$ qubits

A quantum algorithm for parallel sieving



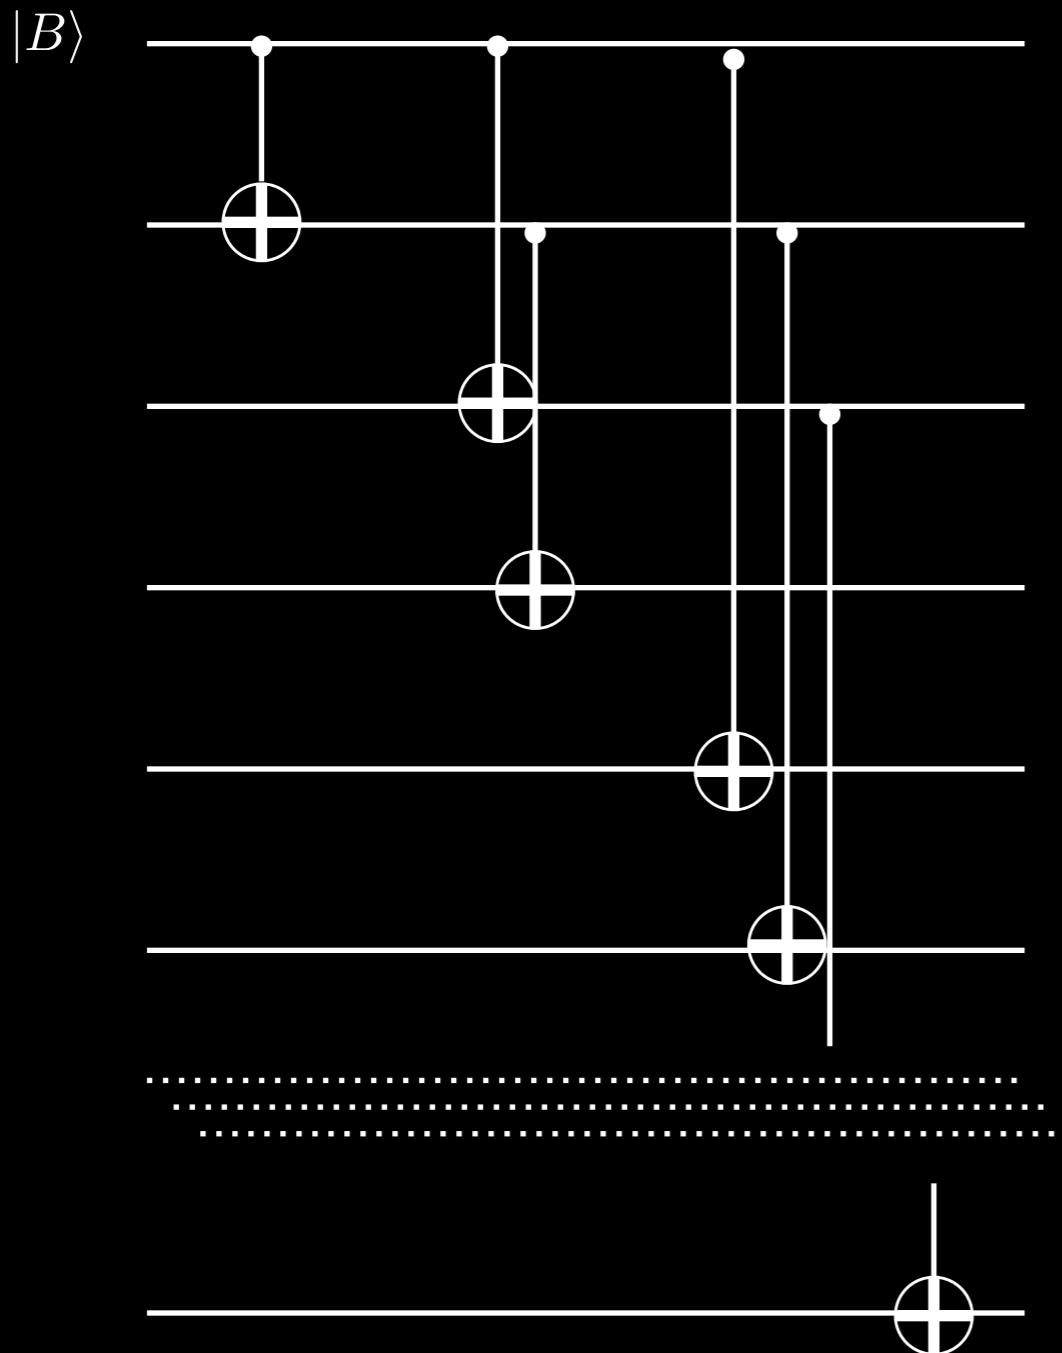
1. Start with $N^{*\text{poly}(d)}$ qubits
2. Copy the basis information

A quantum algorithm for parallel sieving



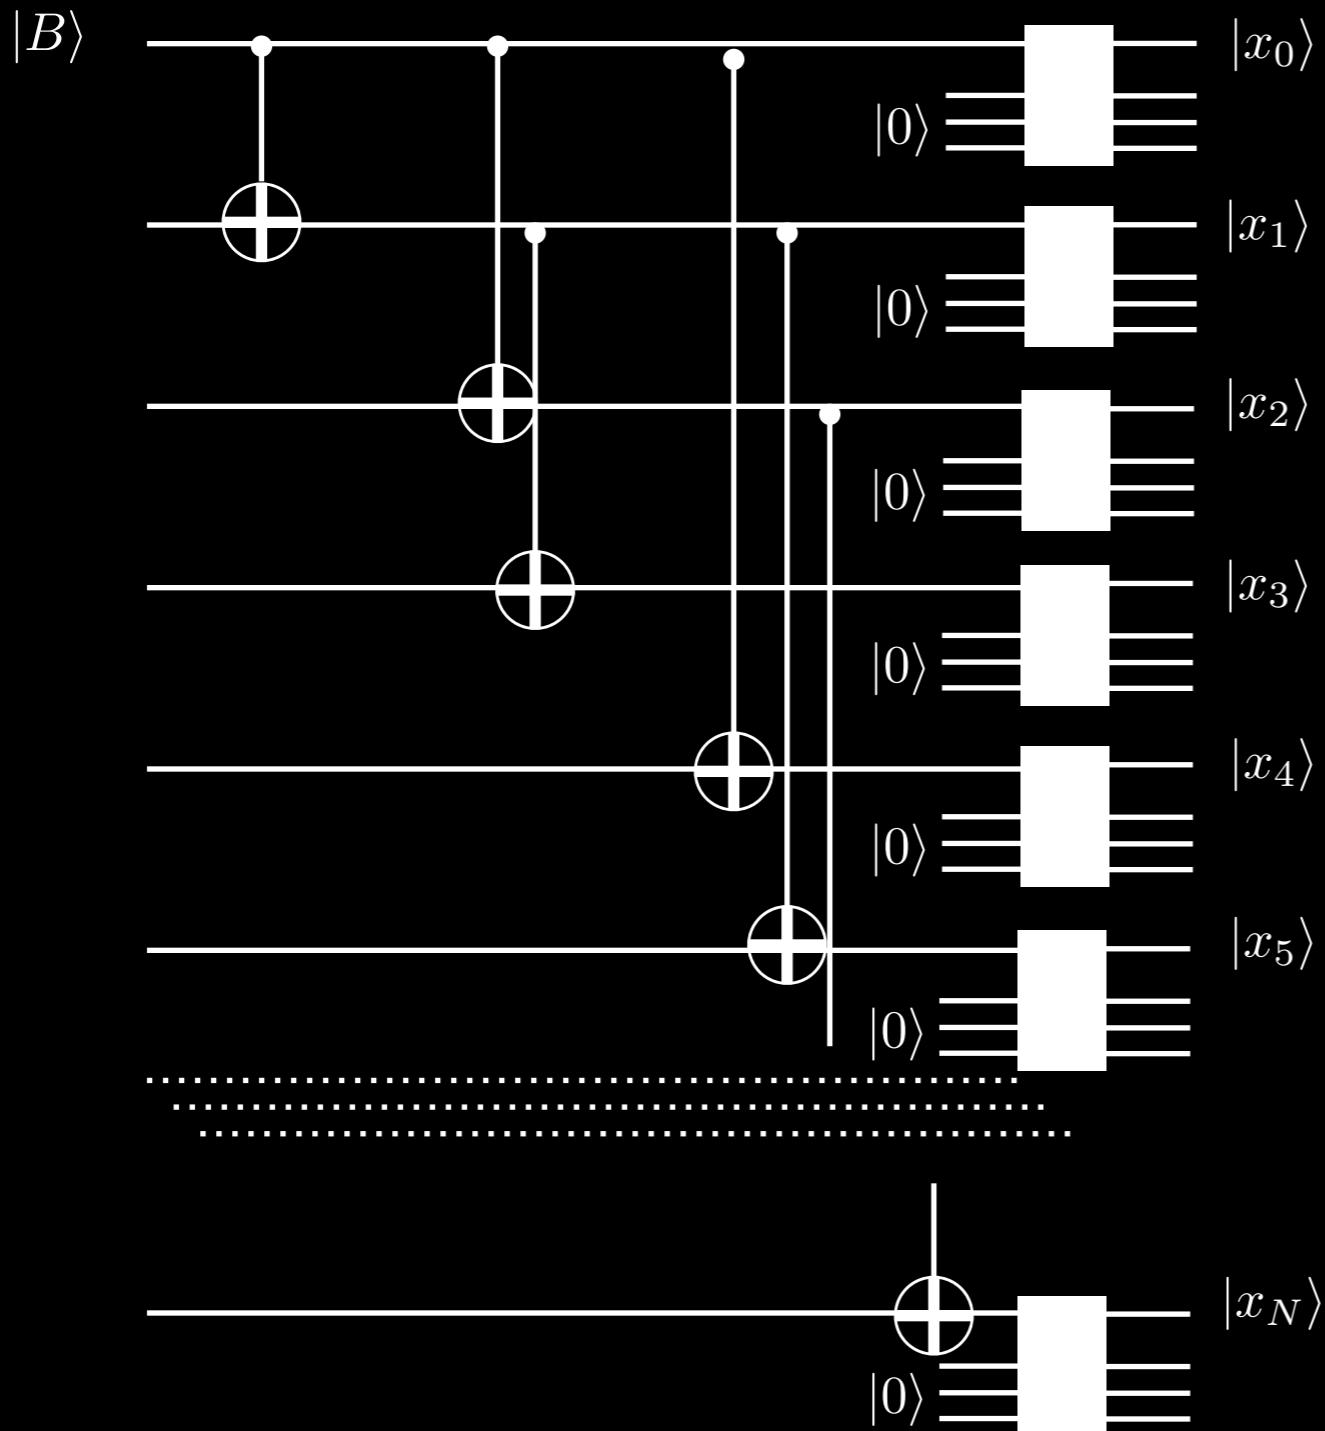
1. Start with $N^{*\text{poly}(d)}$ qubits
2. Copy the basis information

A quantum algorithm for parallel sieving



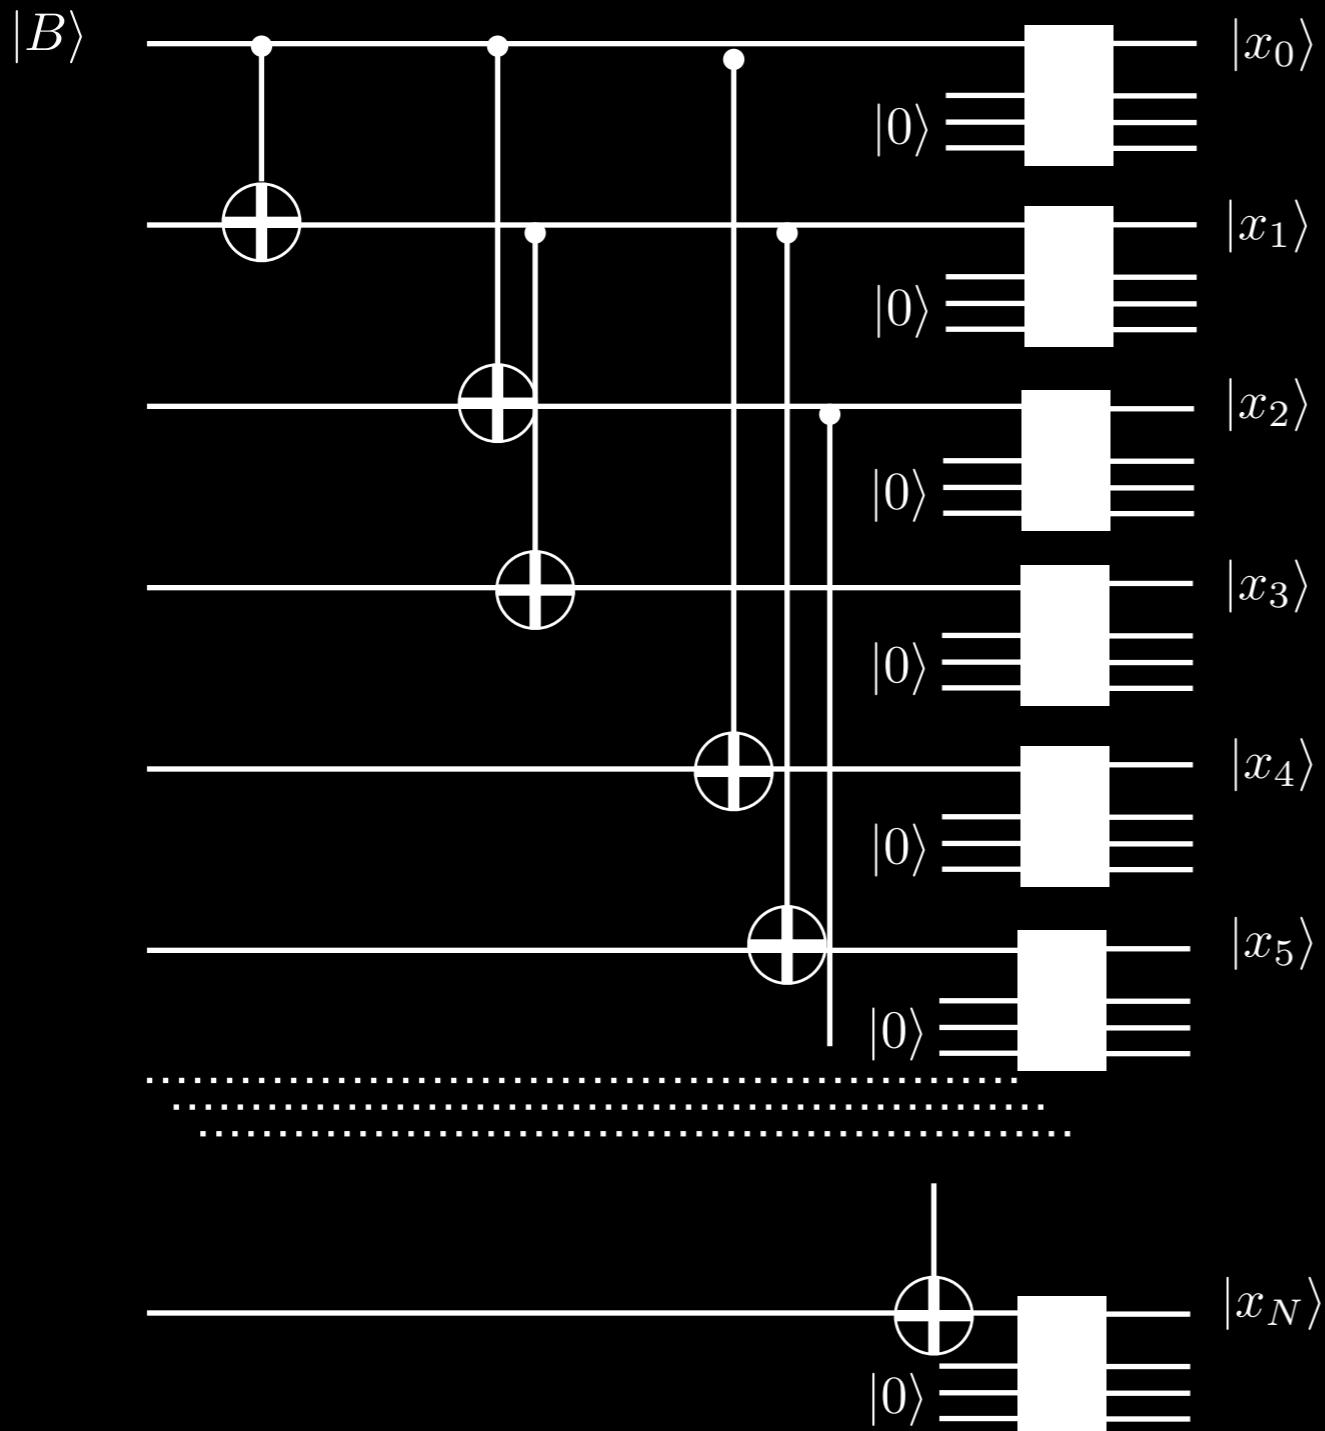
1. Start with $N^{*\text{poly}(d)}$ qubits
2. Copy the basis information

A quantum algorithm for parallel sieving



1. Start with $N^{\star}\text{poly}(d)$ qubits
2. Copy the basis information
3. Sample a vector and normalise it.

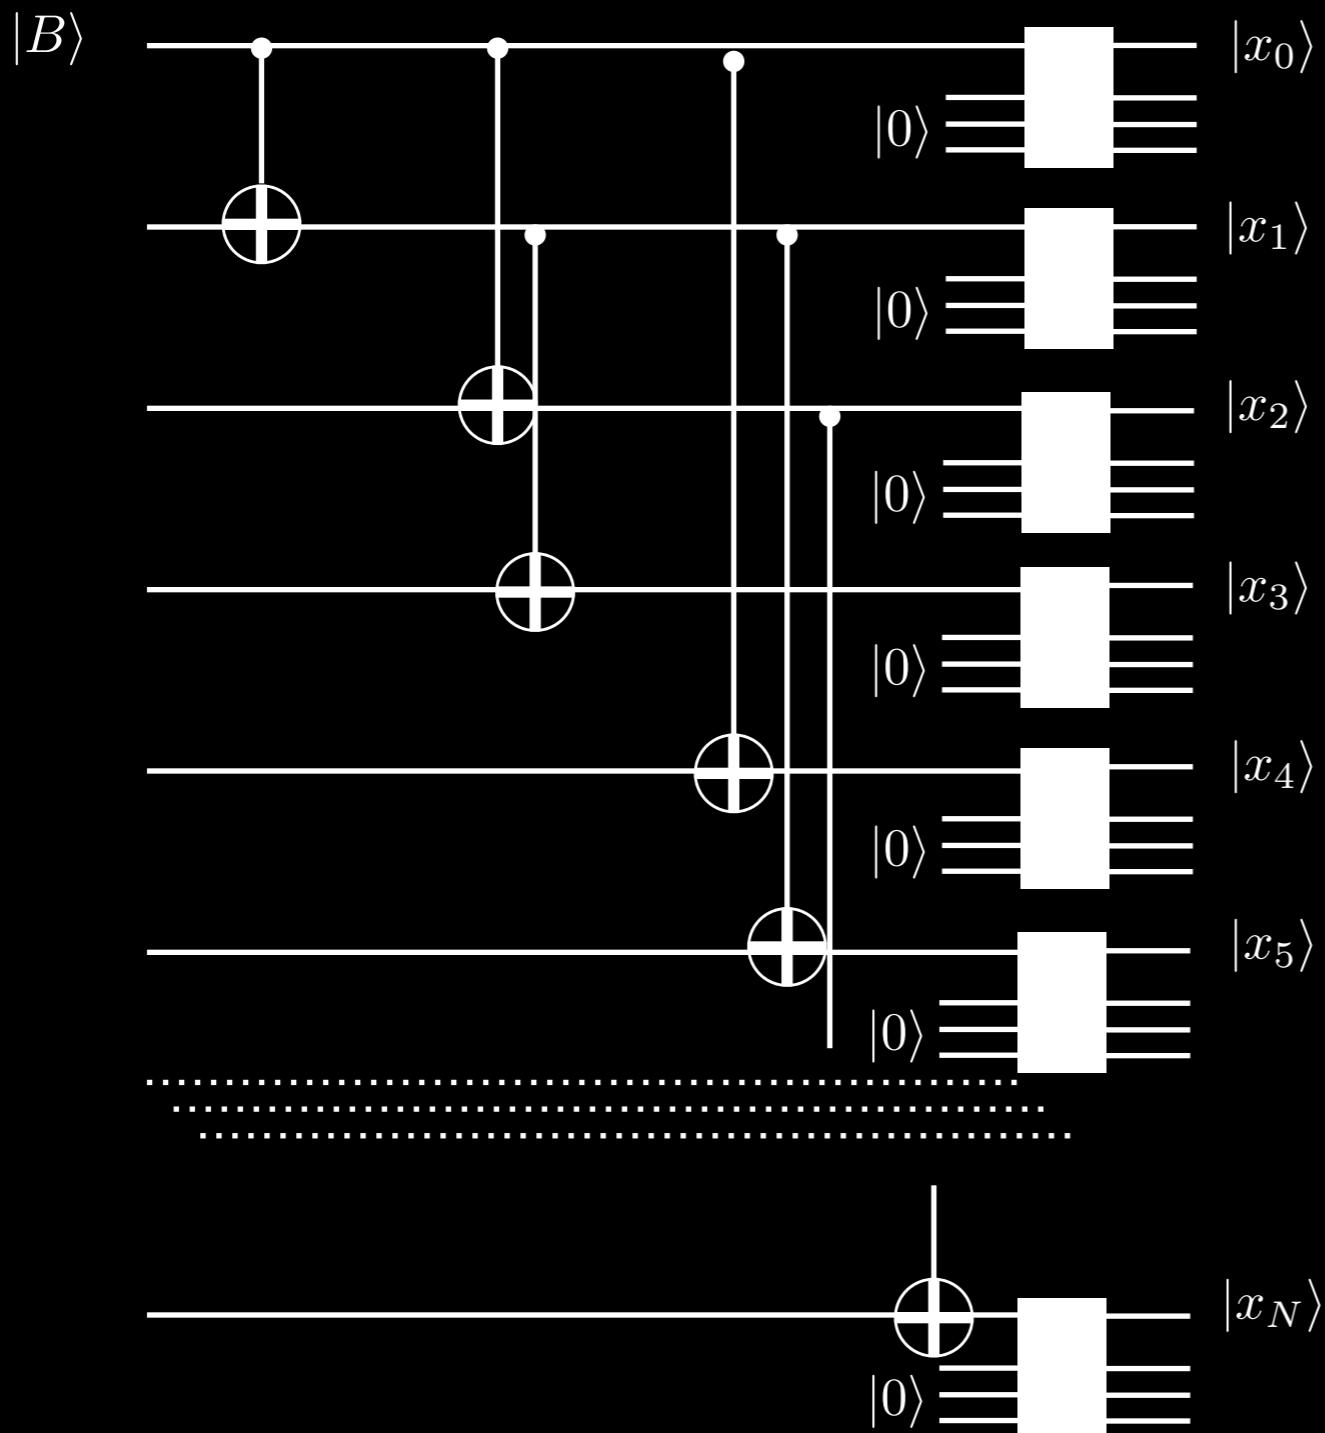
A quantum algorithm for parallel sieving



1. Start with $N^{\star}\text{poly}(d)$ qubits
2. Copy the basis information
3. Sample a vector and normalise it.

$\tilde{O}(N)$ width
 $O(\text{polylog}(N))$ depth

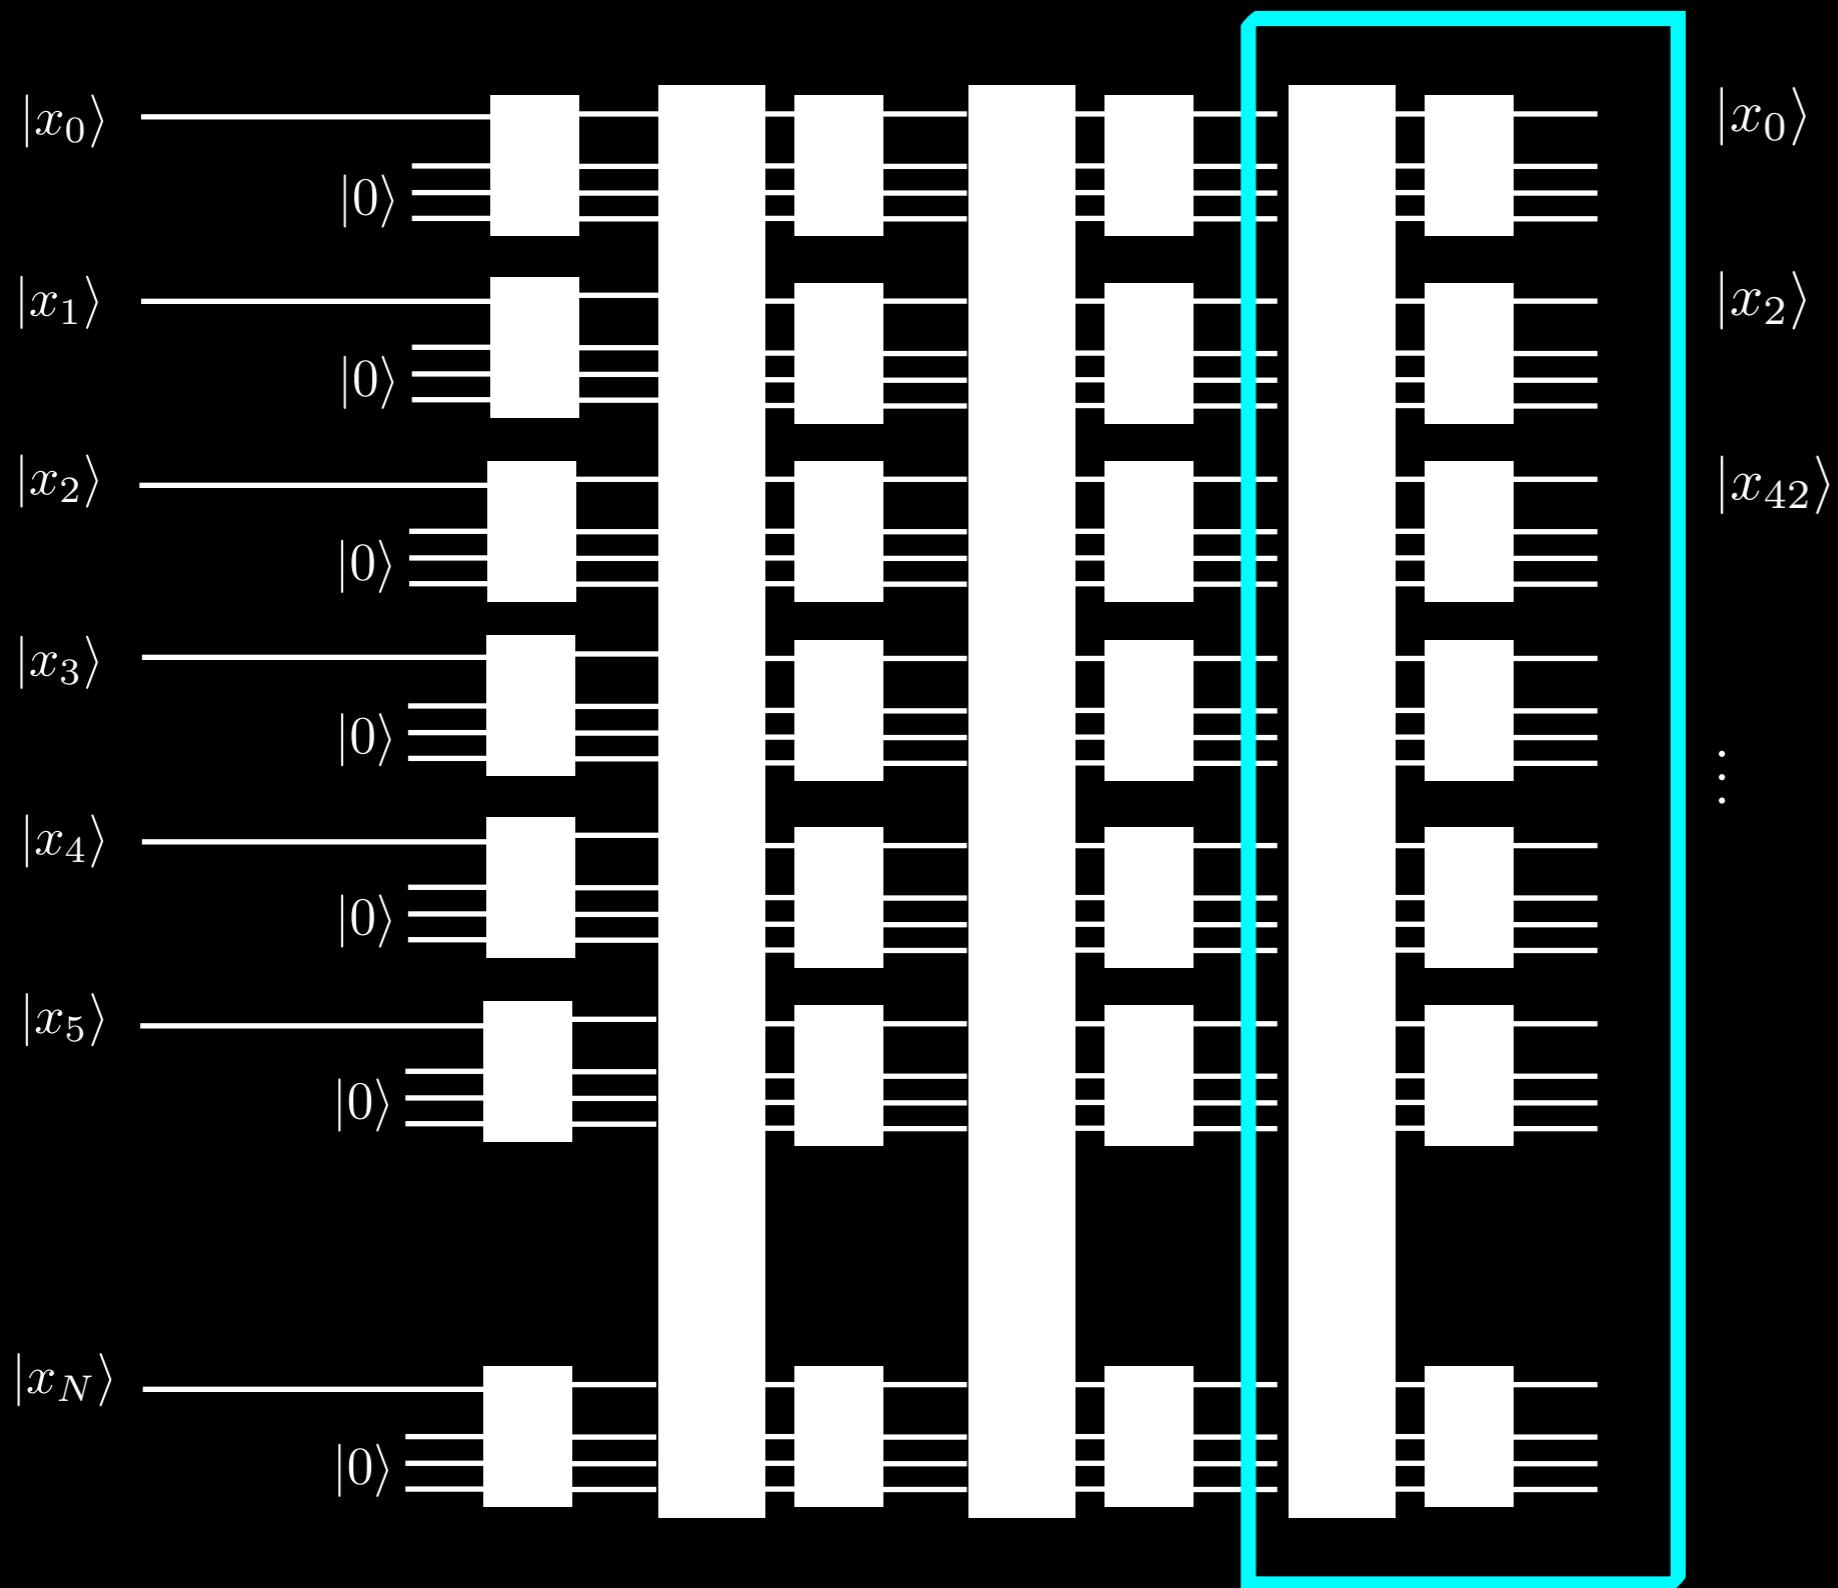
A quantum algorithm for parallel sieving



1. Start with $N^{\star}\text{poly}(d)$ qubits
2. Copy the basis information
3. Sample a vector and normalise it.

$\tilde{O}(N)$ width
 $O(\text{polylog}(N))$ depth

A quantum algorithm for parallel sieving



A quantum algorithm for parallel sieving

U_f recognises a solution

$$U_f : |x\rangle |y\rangle \rightarrow |x\rangle |y \oplus f(x)\rangle$$

set $|y\rangle = \frac{|0\rangle - |1\rangle}{\sqrt{2}}$, and the reduced unitary induces the following transformation

$$|x\rangle \rightarrow (-1)^{f(x)} |x\rangle$$

V loads the data

$$V : |j\rangle |0\rangle |y\rangle |x_1 x_2 \dots x_N\rangle \rightarrow |j\rangle |x_j\rangle |y\rangle |x_1 x_2 \dots x_N\rangle$$

Quantum Search

$$\begin{aligned} & |j_1\rangle |0\rangle |y_1\rangle |x_1 \dots x_N\rangle \\ \xrightarrow{V} & |j_1\rangle |x_{j_1}\rangle |y_1\rangle |x_1 \dots x_N\rangle \\ \xrightarrow{\otimes U_{f_i}} & |j_1\rangle |x_{j_1}\rangle |y_1 \oplus f_1(x_{j_1})\rangle |x_1 \dots x_N\rangle \\ \xrightarrow{V} & |j_1\rangle |0\rangle |y_1 \oplus f_1(x_{j_1})\rangle |x_1 \dots x_N\rangle \end{aligned}$$

Setting $|y_1\rangle = |- \rangle$, and $\mathcal{W} = VUV$ the reduced unitary becomes

$$\mathcal{W} : |j_1\rangle \rightarrow (-1)^{f_1(x_{j_1})} |j_1\rangle .$$

The unitary \mathcal{W} can be used as an oracle for Grover's algorithm. Further, using \mathcal{W} as an oracle, requires a circuit width of $\Omega(Nd)$, and depth $\Omega(N(\log N + D_f))$ to find one solution with high probability.

A quantum algorithm for parallel sieving

Define $f(x, y) = 1$ if $| \langle x, y \rangle | \geq 1/2$; 0 otherwise.

$$U_f : |x, y\rangle |b\rangle \rightarrow |x, y\rangle |b \oplus f(x, y)\rangle$$

Now for each 'wire' i , define $f_i(y) = 1$ if $| \langle x_i, y \rangle | \geq 1/2$; 0 otherwise

$$U_{f_i} : |y\rangle |b\rangle \rightarrow |y\rangle |b \oplus f_i(x_i, y)\rangle$$

$$V : |j_1 \dots j_N\rangle |y_1 \dots y_N\rangle |x_1 \dots x_N\rangle \rightarrow |j_1 \dots j_N\rangle |y_1 + x_{j_1} \dots y_N + x_{j_N}\rangle |x_1 x_2 \dots x_N\rangle$$

A quantum algorithm for parallel sieving

Parallel Quantum Search

$$\begin{aligned} & |j_1 \dots j_N\rangle |0 \dots 0\rangle |y_1 \dots y_N\rangle |x_1 \dots x_N\rangle \\ \xrightarrow{V} & |j_1 \dots j_N\rangle |x_{j_1} \dots x_{j_N}\rangle |y_1 \dots y_N\rangle |x_1 \dots x_N\rangle \\ \xrightarrow{\otimes U_{f_i}^{\otimes N}} & |j_1 \dots j_N\rangle |x_{j_1} \dots x_{j_N}\rangle |y_1 \oplus f_1(x_{j_1}) \dots y_N \oplus f_N(x_{j_N})\rangle |x_1 \dots x_N\rangle \\ \xrightarrow{V} & |j_1 \dots j_N\rangle |0 \dots 0\rangle |y_1 \oplus f_1(x_{j_1}) \dots y_N \oplus f_N(x_{j_N})\rangle |x_1 \dots x_N\rangle \end{aligned}$$

Setting $|y_i\rangle = |- \rangle$, for all i , $\mathcal{W} = VU_{f_i}^{\otimes N}V$ the reduced unitary becomes

$$\mathcal{W} : |j_1 \dots j_N\rangle \rightarrow (-1)^{f_1(x_{j_1})} |j_1\rangle \dots (-1)^{f_N(x_{j_N})} |j_N\rangle.$$

The unitary \mathcal{W} can be implemented using a uniform family of circuits of width $N \cdot (\log N + d) = N \cdot \text{polylog} N$ and depth $O(\log N \log(d \log N) + \text{poly}(d)) = \text{polylog}(N)$. Define R to be a rotation along $|0\rangle$ and run $R \circ \mathcal{W}, \Theta(\sqrt{N})$ times, on the state

$$\sum_{j_1, \dots, j_N} |j_1 \dots j_N\rangle |0 \dots 0\rangle |- \dots -\rangle |x_1 \dots x_N\rangle$$

A quantum algorithm for parallel sieving

Each wire, i , that had previously sample x_i , thus finds a x'_i , such that x_i and x'_i are remarkable close, and computes the difference and sets that as the new vector $y_i = \min\{x_i \pm x'_i\}$.

x_i is stored in register $L1[i]$

x'_i is stored in register $L2[i]$

y_i is stored in register $L3[i]$

Swap registers L3 and L1. And repeat this process $\text{poly}(d)$ times.

Search through L3, and output a vector shorter than the target norm (Euclidean norm less than λ).

This algorithm can be implemented with a quantum circuit of depth $\tilde{O}(\sqrt{N})$ and width $\tilde{O}(N)$, and recall $N = 2^{0.2075d}$, thus it take the quantum circuit $2^{0.1037d+o(d)}$ time steps and $2^{0.2075d+o(d)}$ memory

A classical algorithm for distributed sieving

by David Harvey and Paul Pollack

University of Massachusetts Amherst

and University of Colorado Boulder

Joint work with Michael Rubinstein and Andrew Sutherland

and the LMFDB team: Cremona, Fisher, Kedlaya, et al.

and the Sieve team: Ganthier, Klyve, et al.

and the Sage team: Cremona, Fisher, Kedlaya, et al.

and the LMFDB team: Cremona, Fisher, Kedlaya, et al.

and the Sieve team: Ganthier, Klyve, et al.

and the Sage team: Cremona, Fisher, Kedlaya, et al.

and the LMFDB team: Cremona, Fisher, Kedlaya, et al.

and the Sieve team: Ganthier, Klyve, et al.

and the Sage team: Cremona, Fisher, Kedlaya, et al.

and the LMFDB team: Cremona, Fisher, Kedlaya, et al.

and the Sieve team: Ganthier, Klyve, et al.

and the Sage team: Cremona, Fisher, Kedlaya, et al.

and the LMFDB team: Cremona, Fisher, Kedlaya, et al.

and the Sieve team: Ganthier, Klyve, et al.

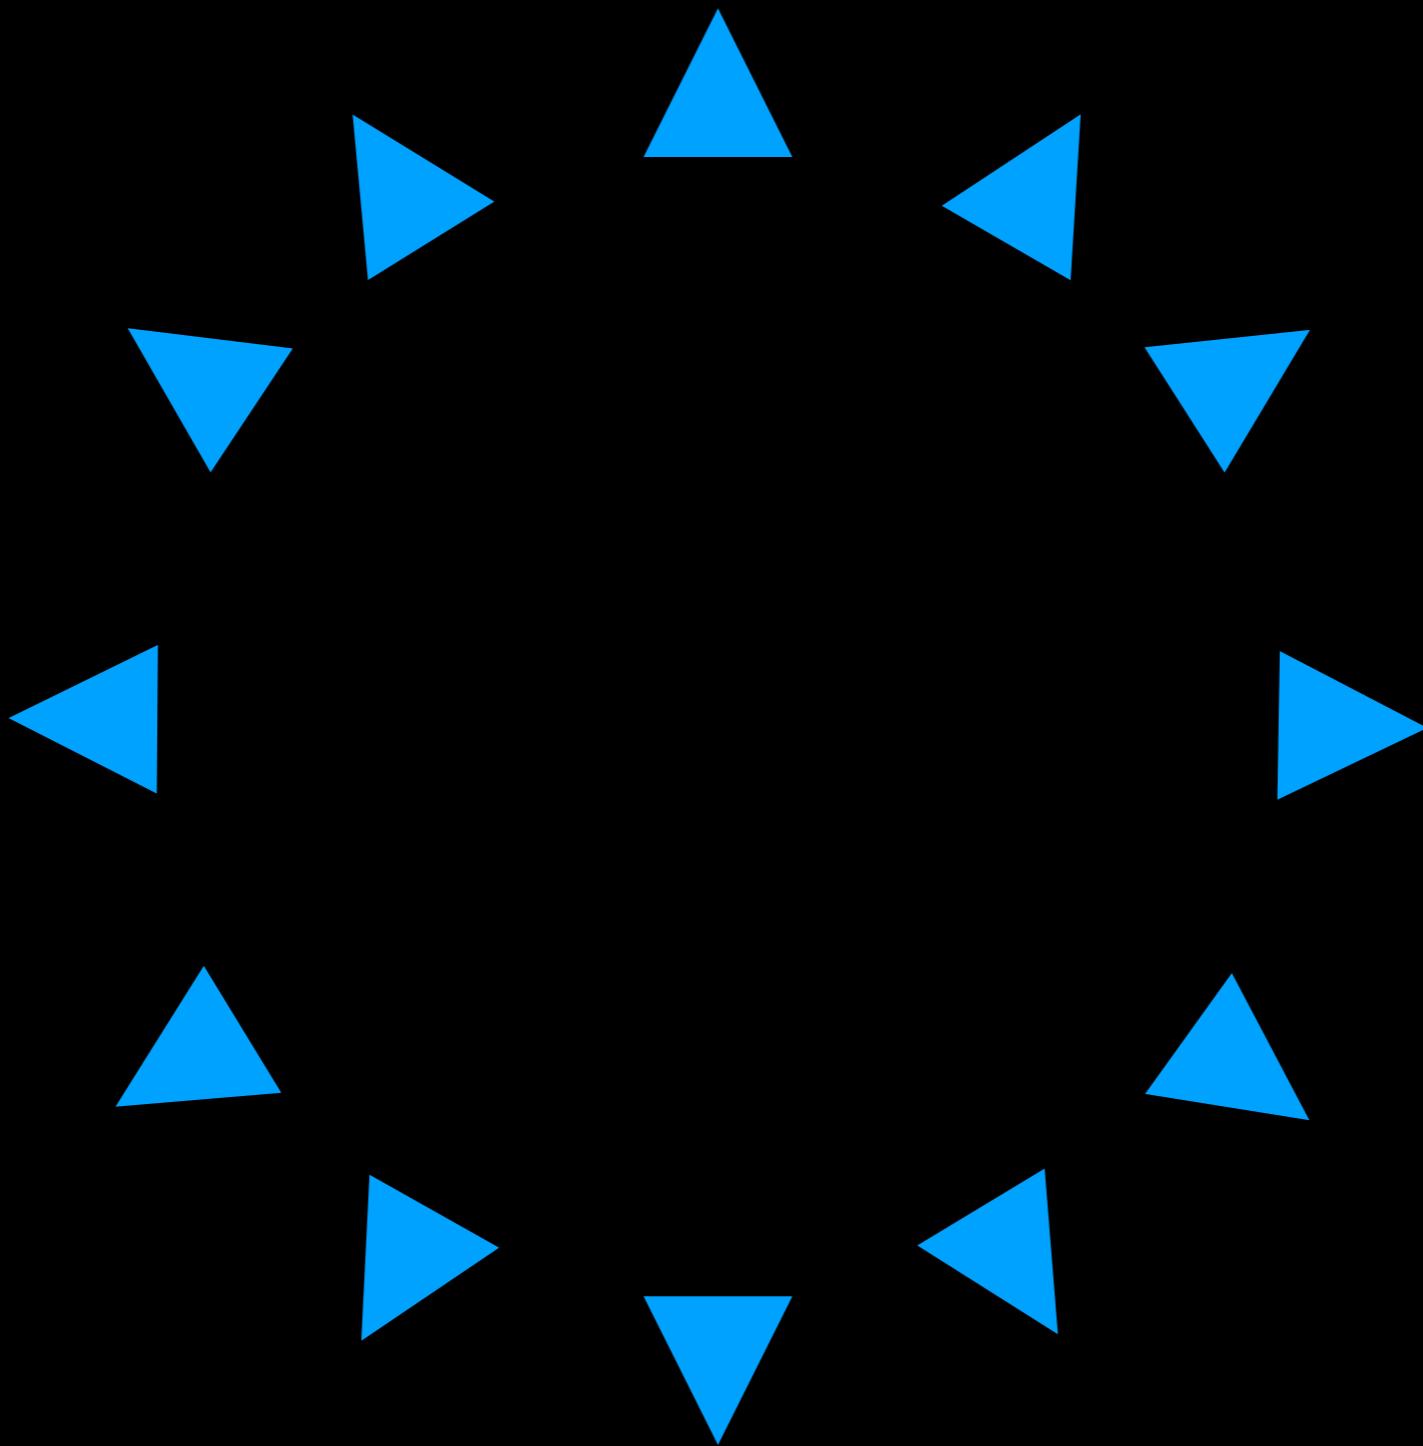
and the Sage team: Cremona, Fisher, Kedlaya, et al.

and the LMFDB team: Cremona, Fisher, Kedlaya, et al.

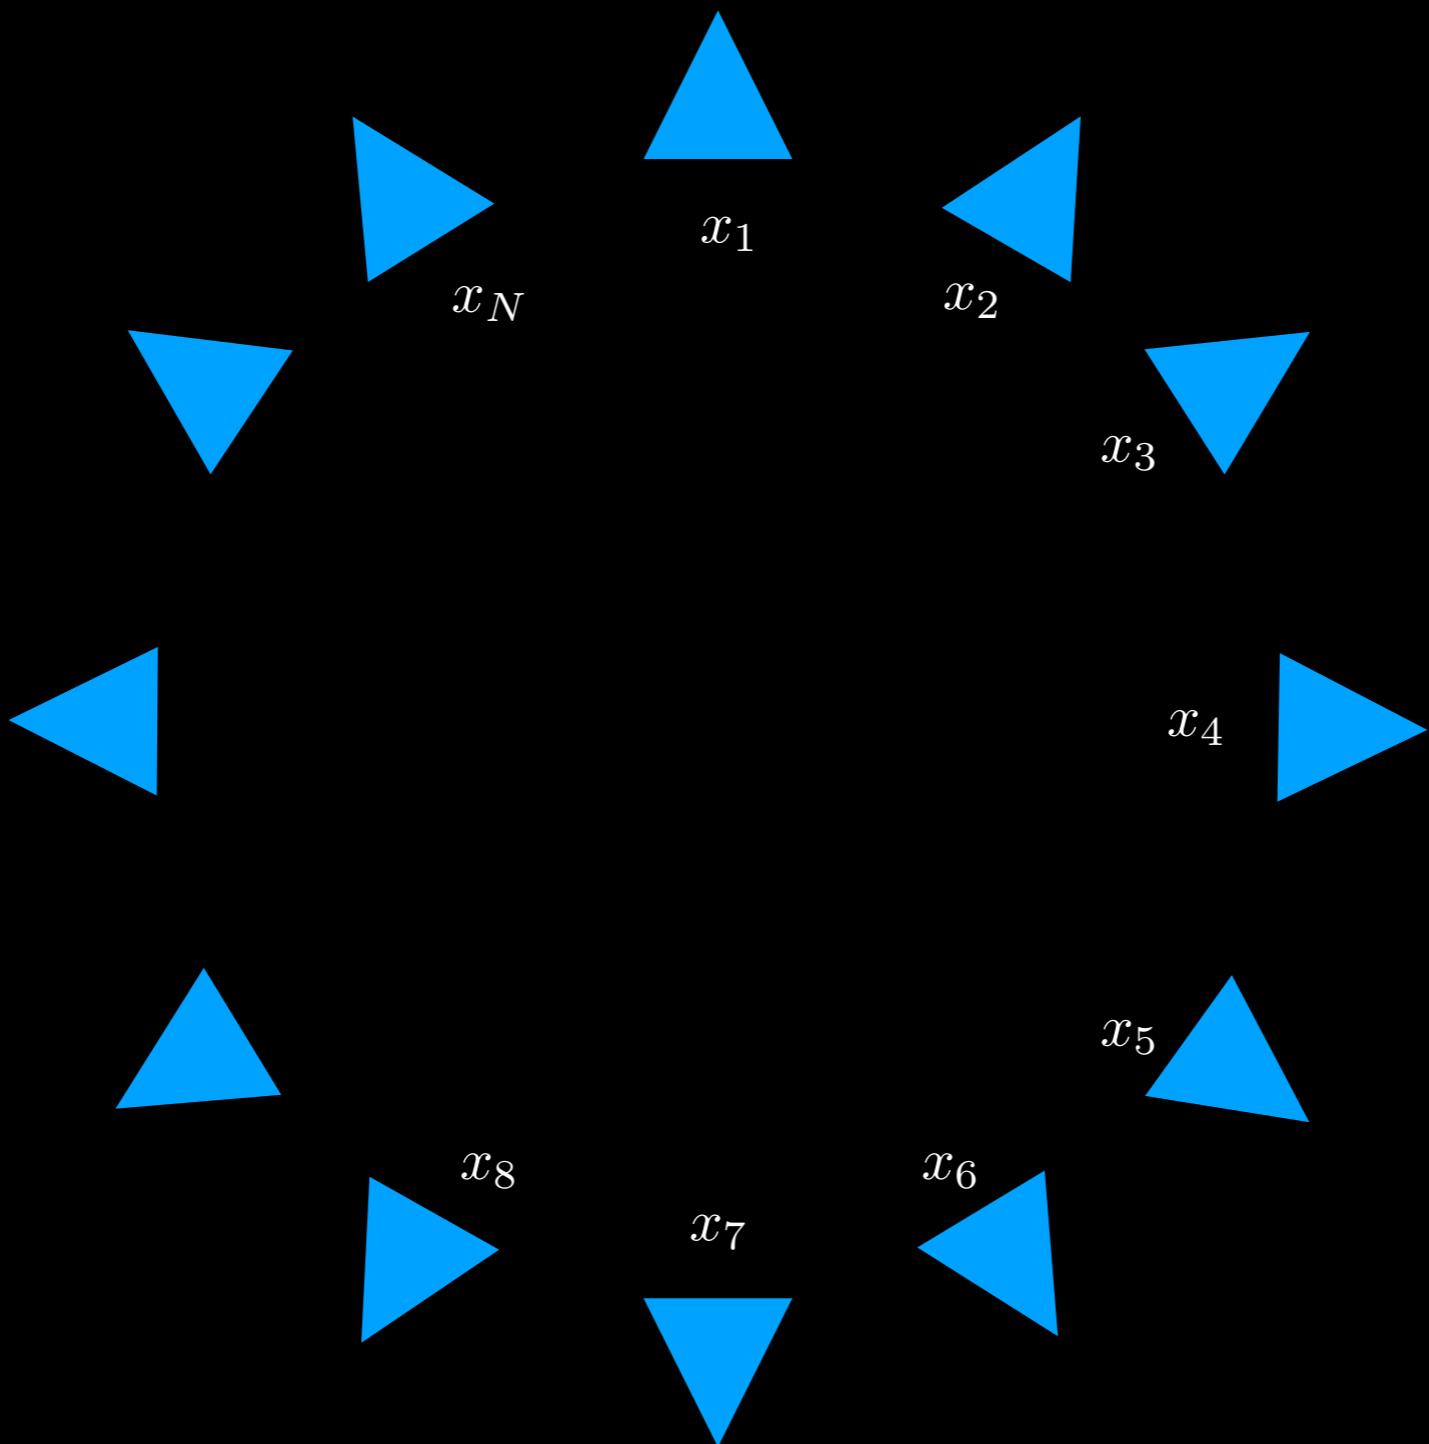
and the Sieve team: Ganthier, Klyve, et al.

and the Sage team: Cremona, Fisher, Kedlaya, et al.

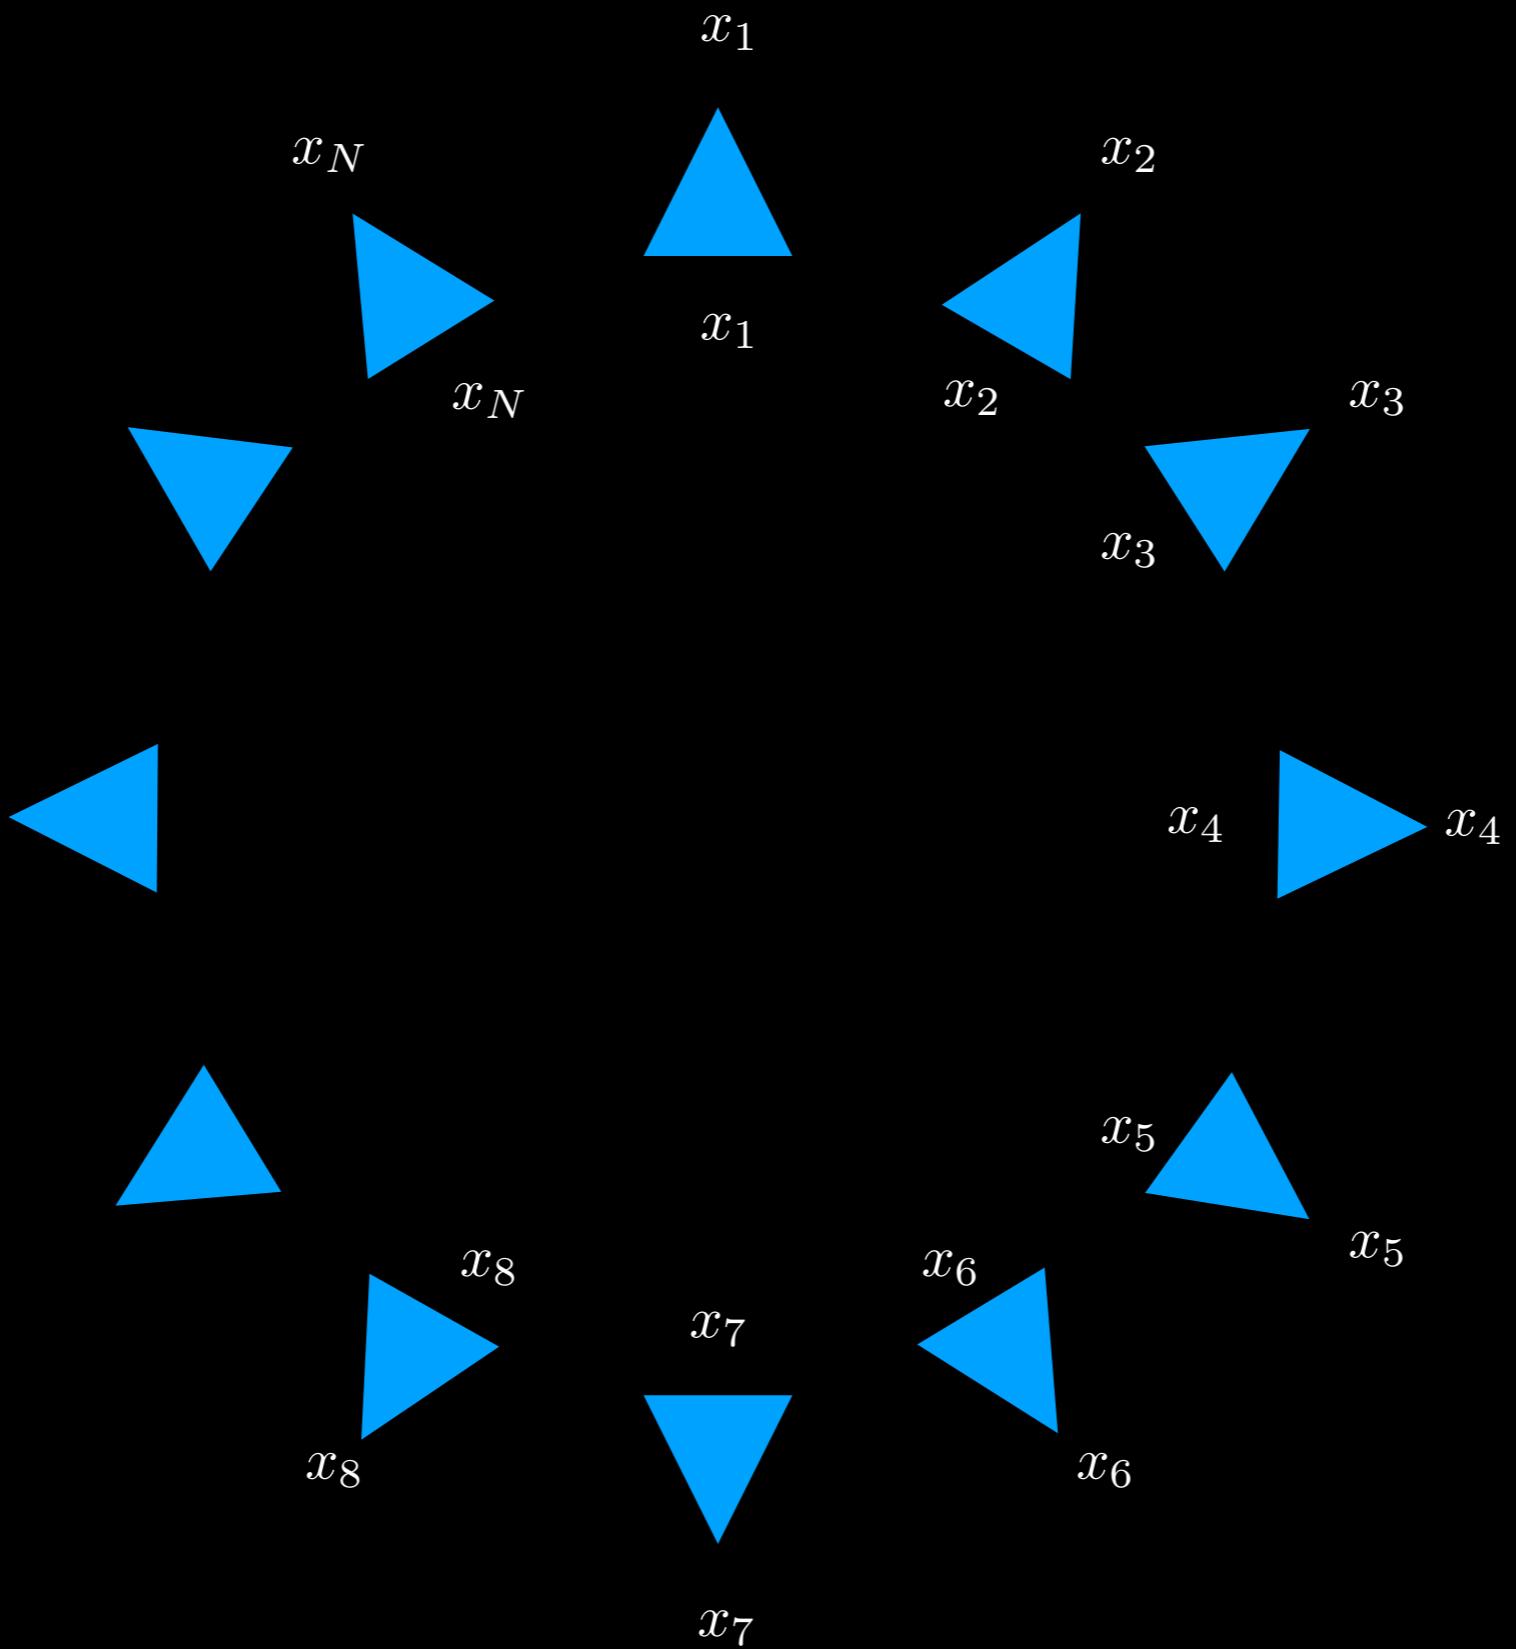
A classical algorithm for distributed sieving



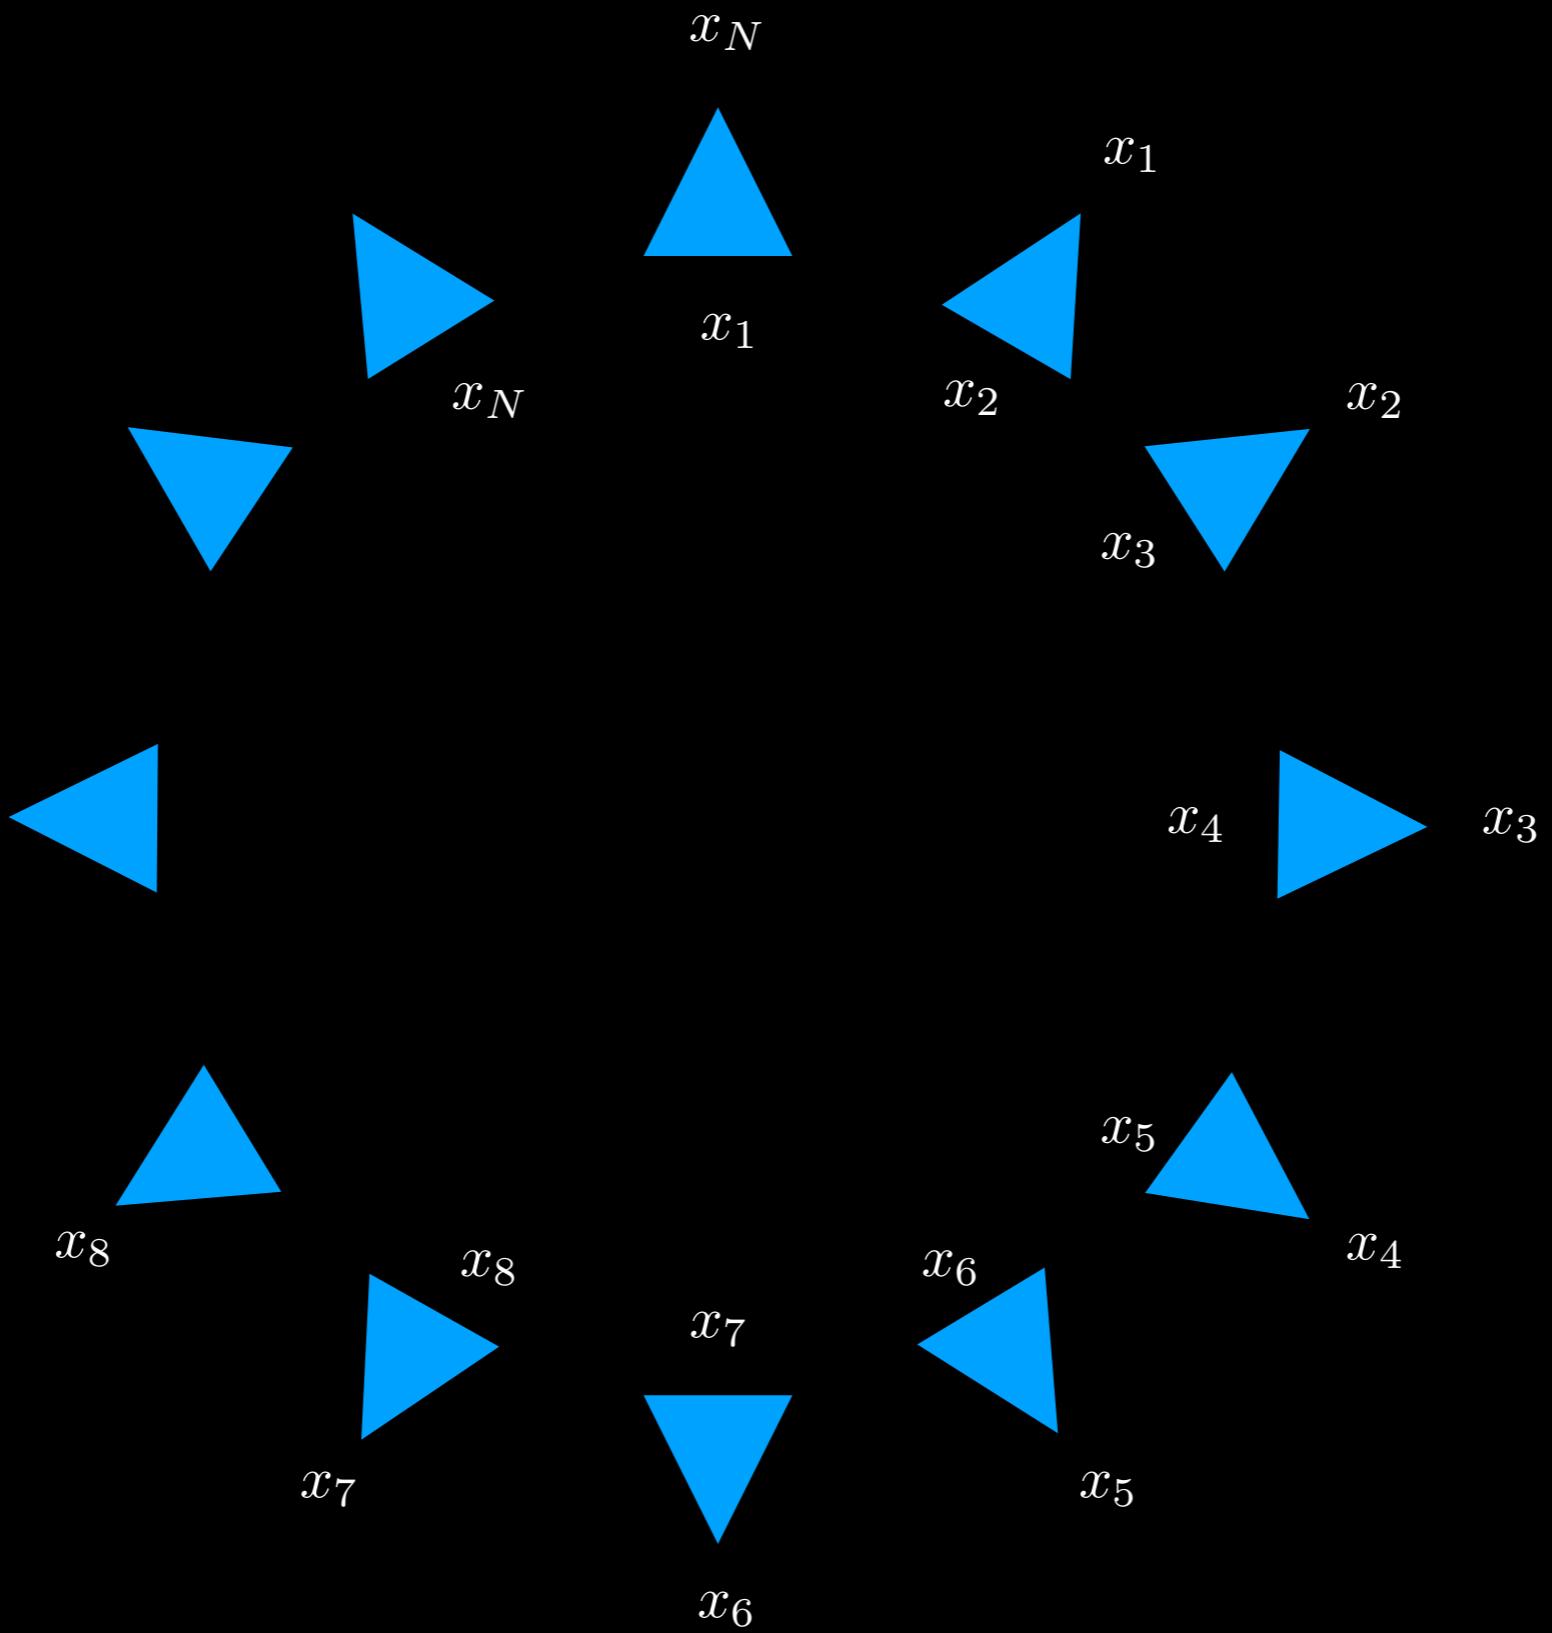
A classical algorithm for distributed sieving



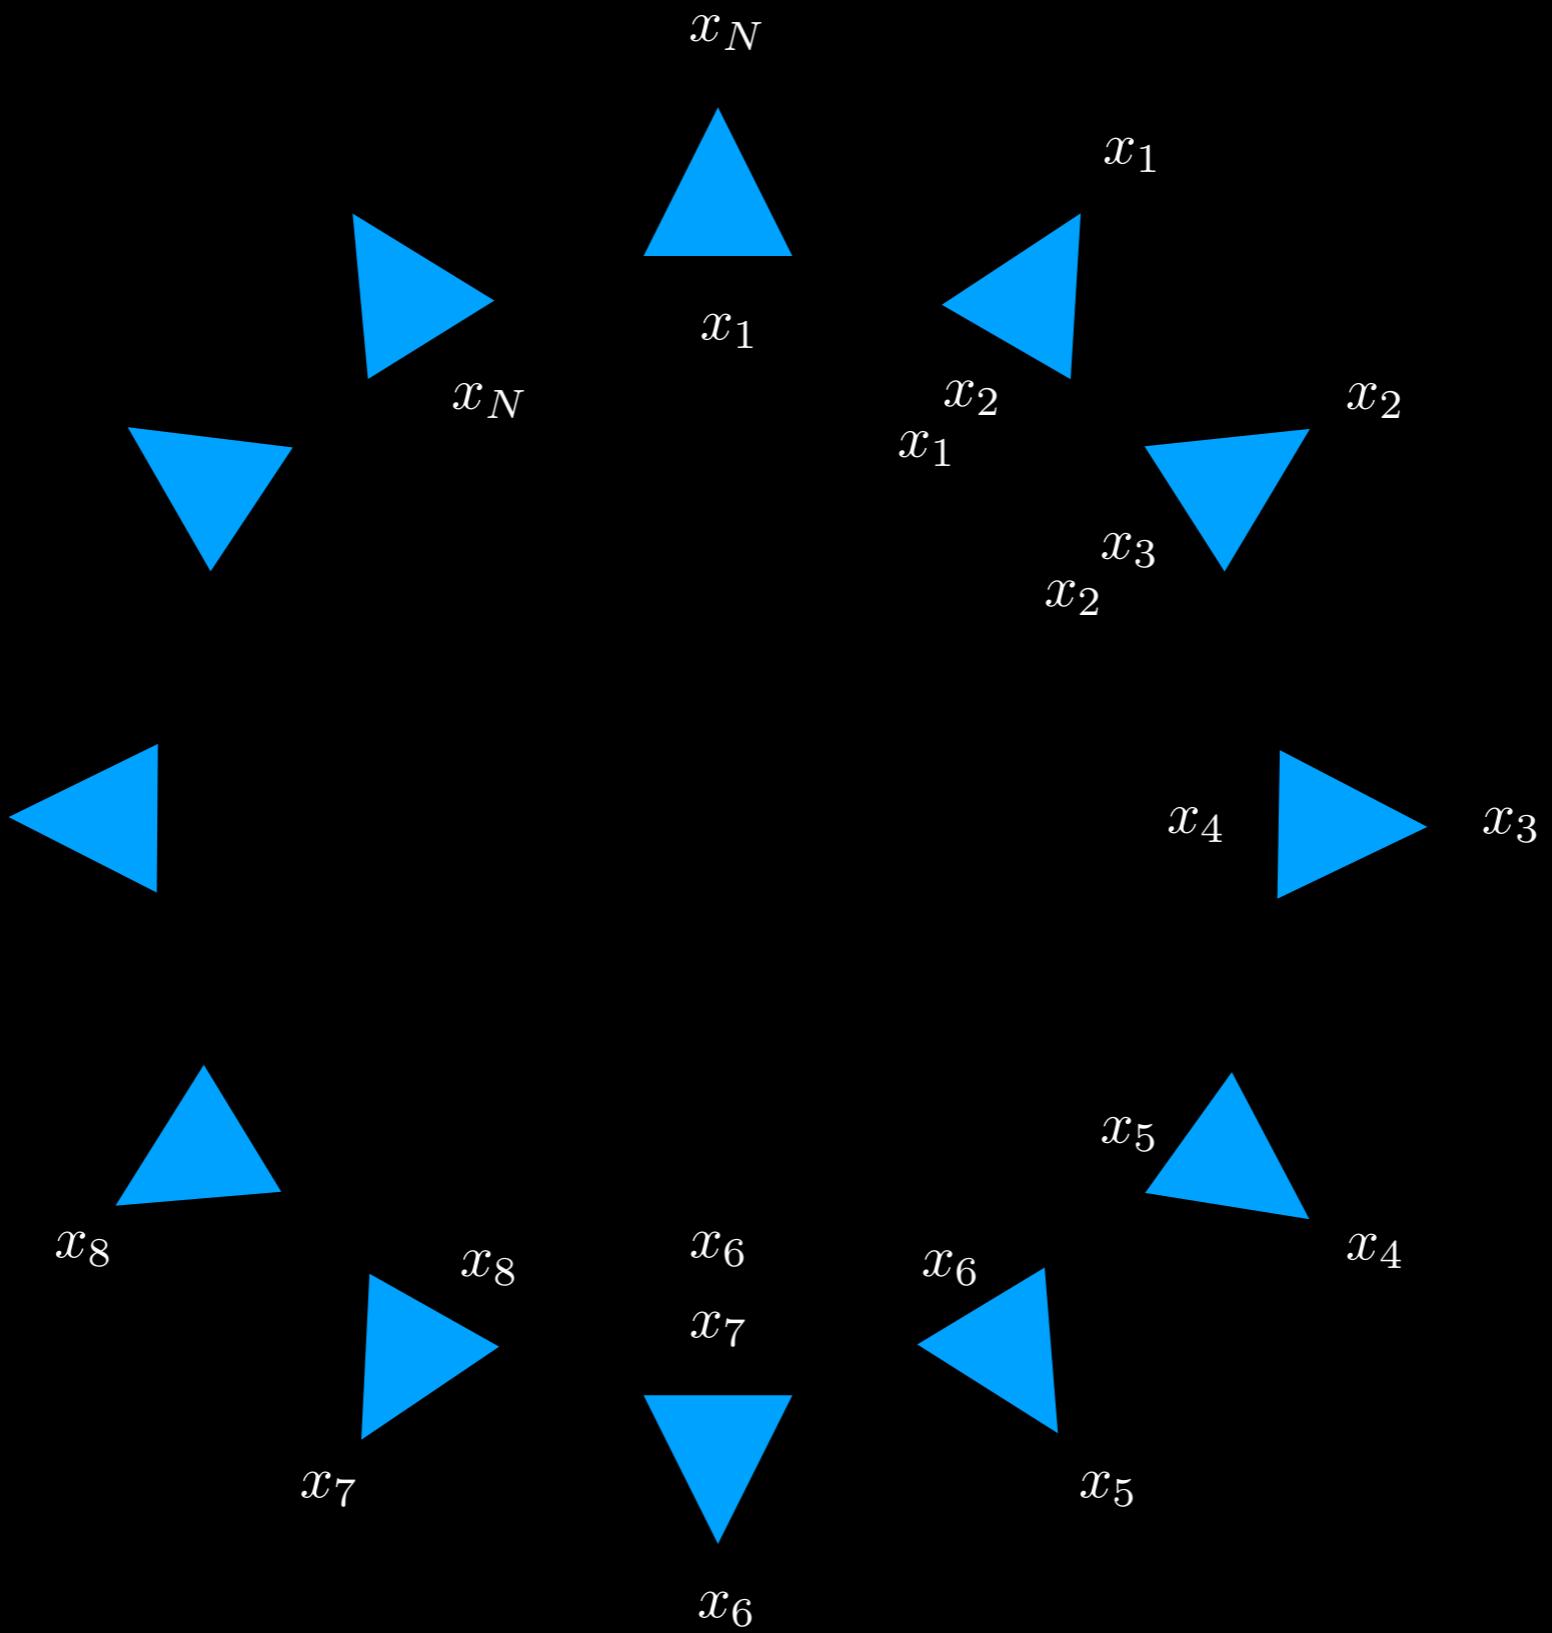
A classical algorithm for distributed sieving



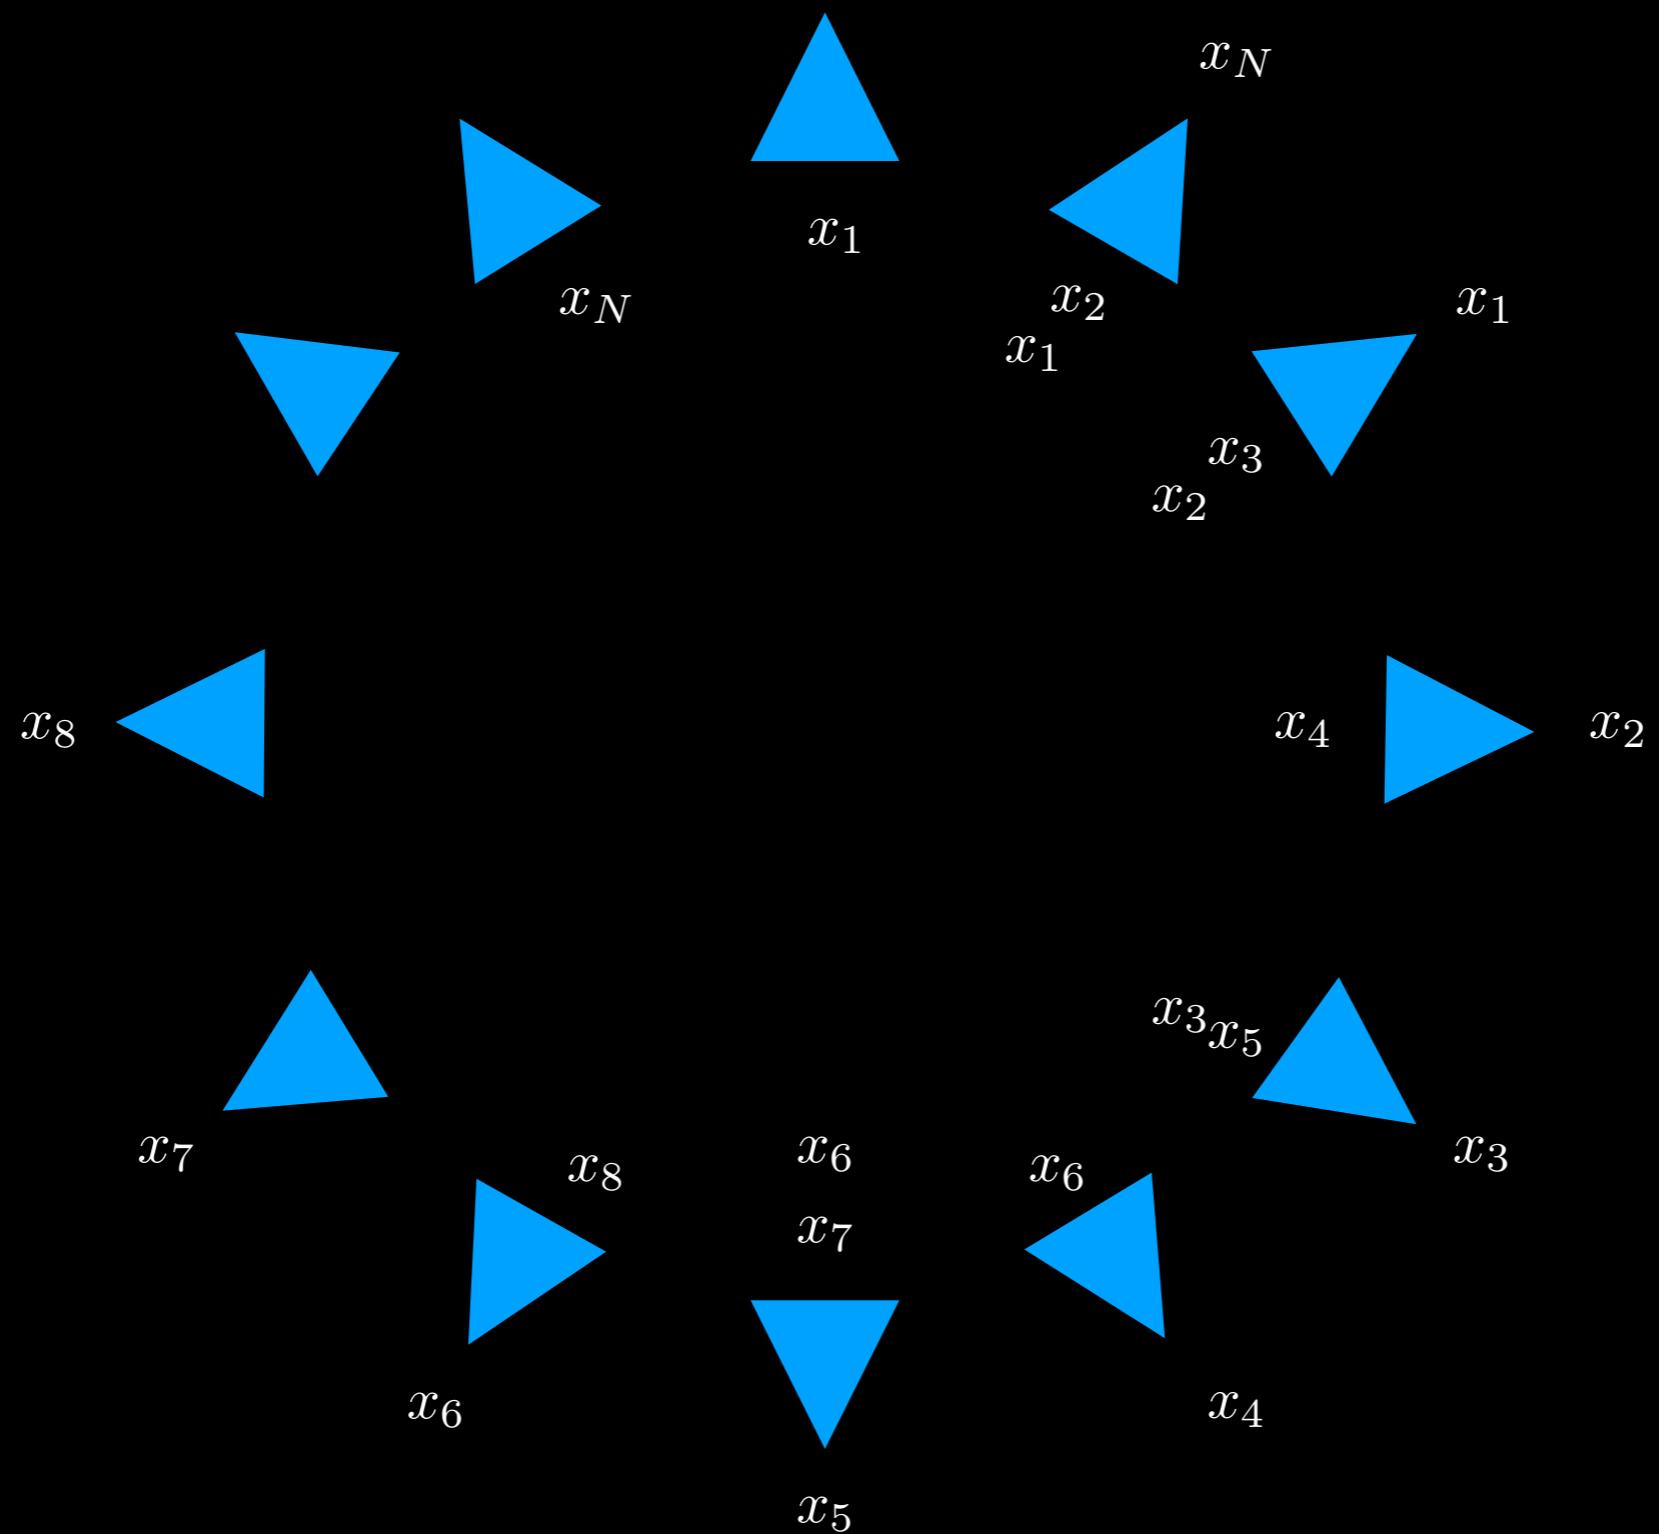
A classical algorithm for distributed sieving



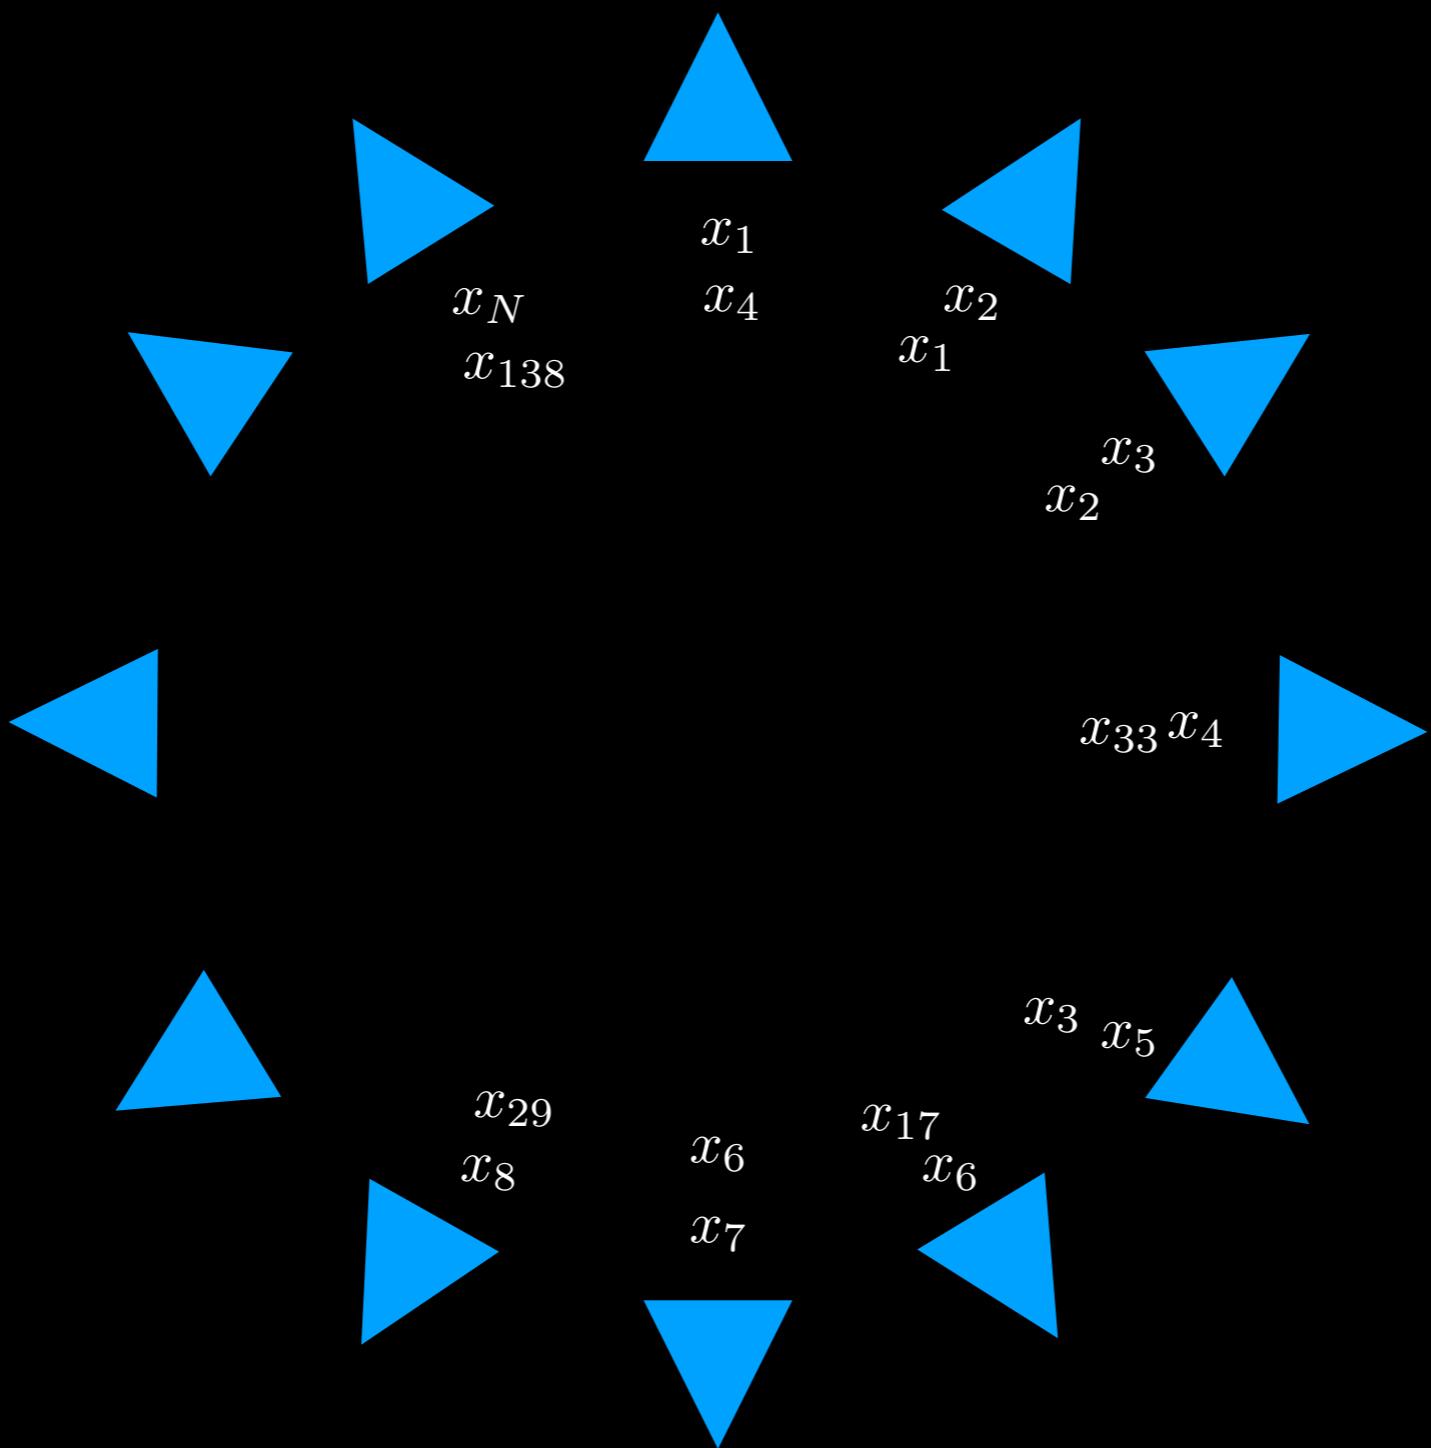
A classical algorithm for distributed sieving



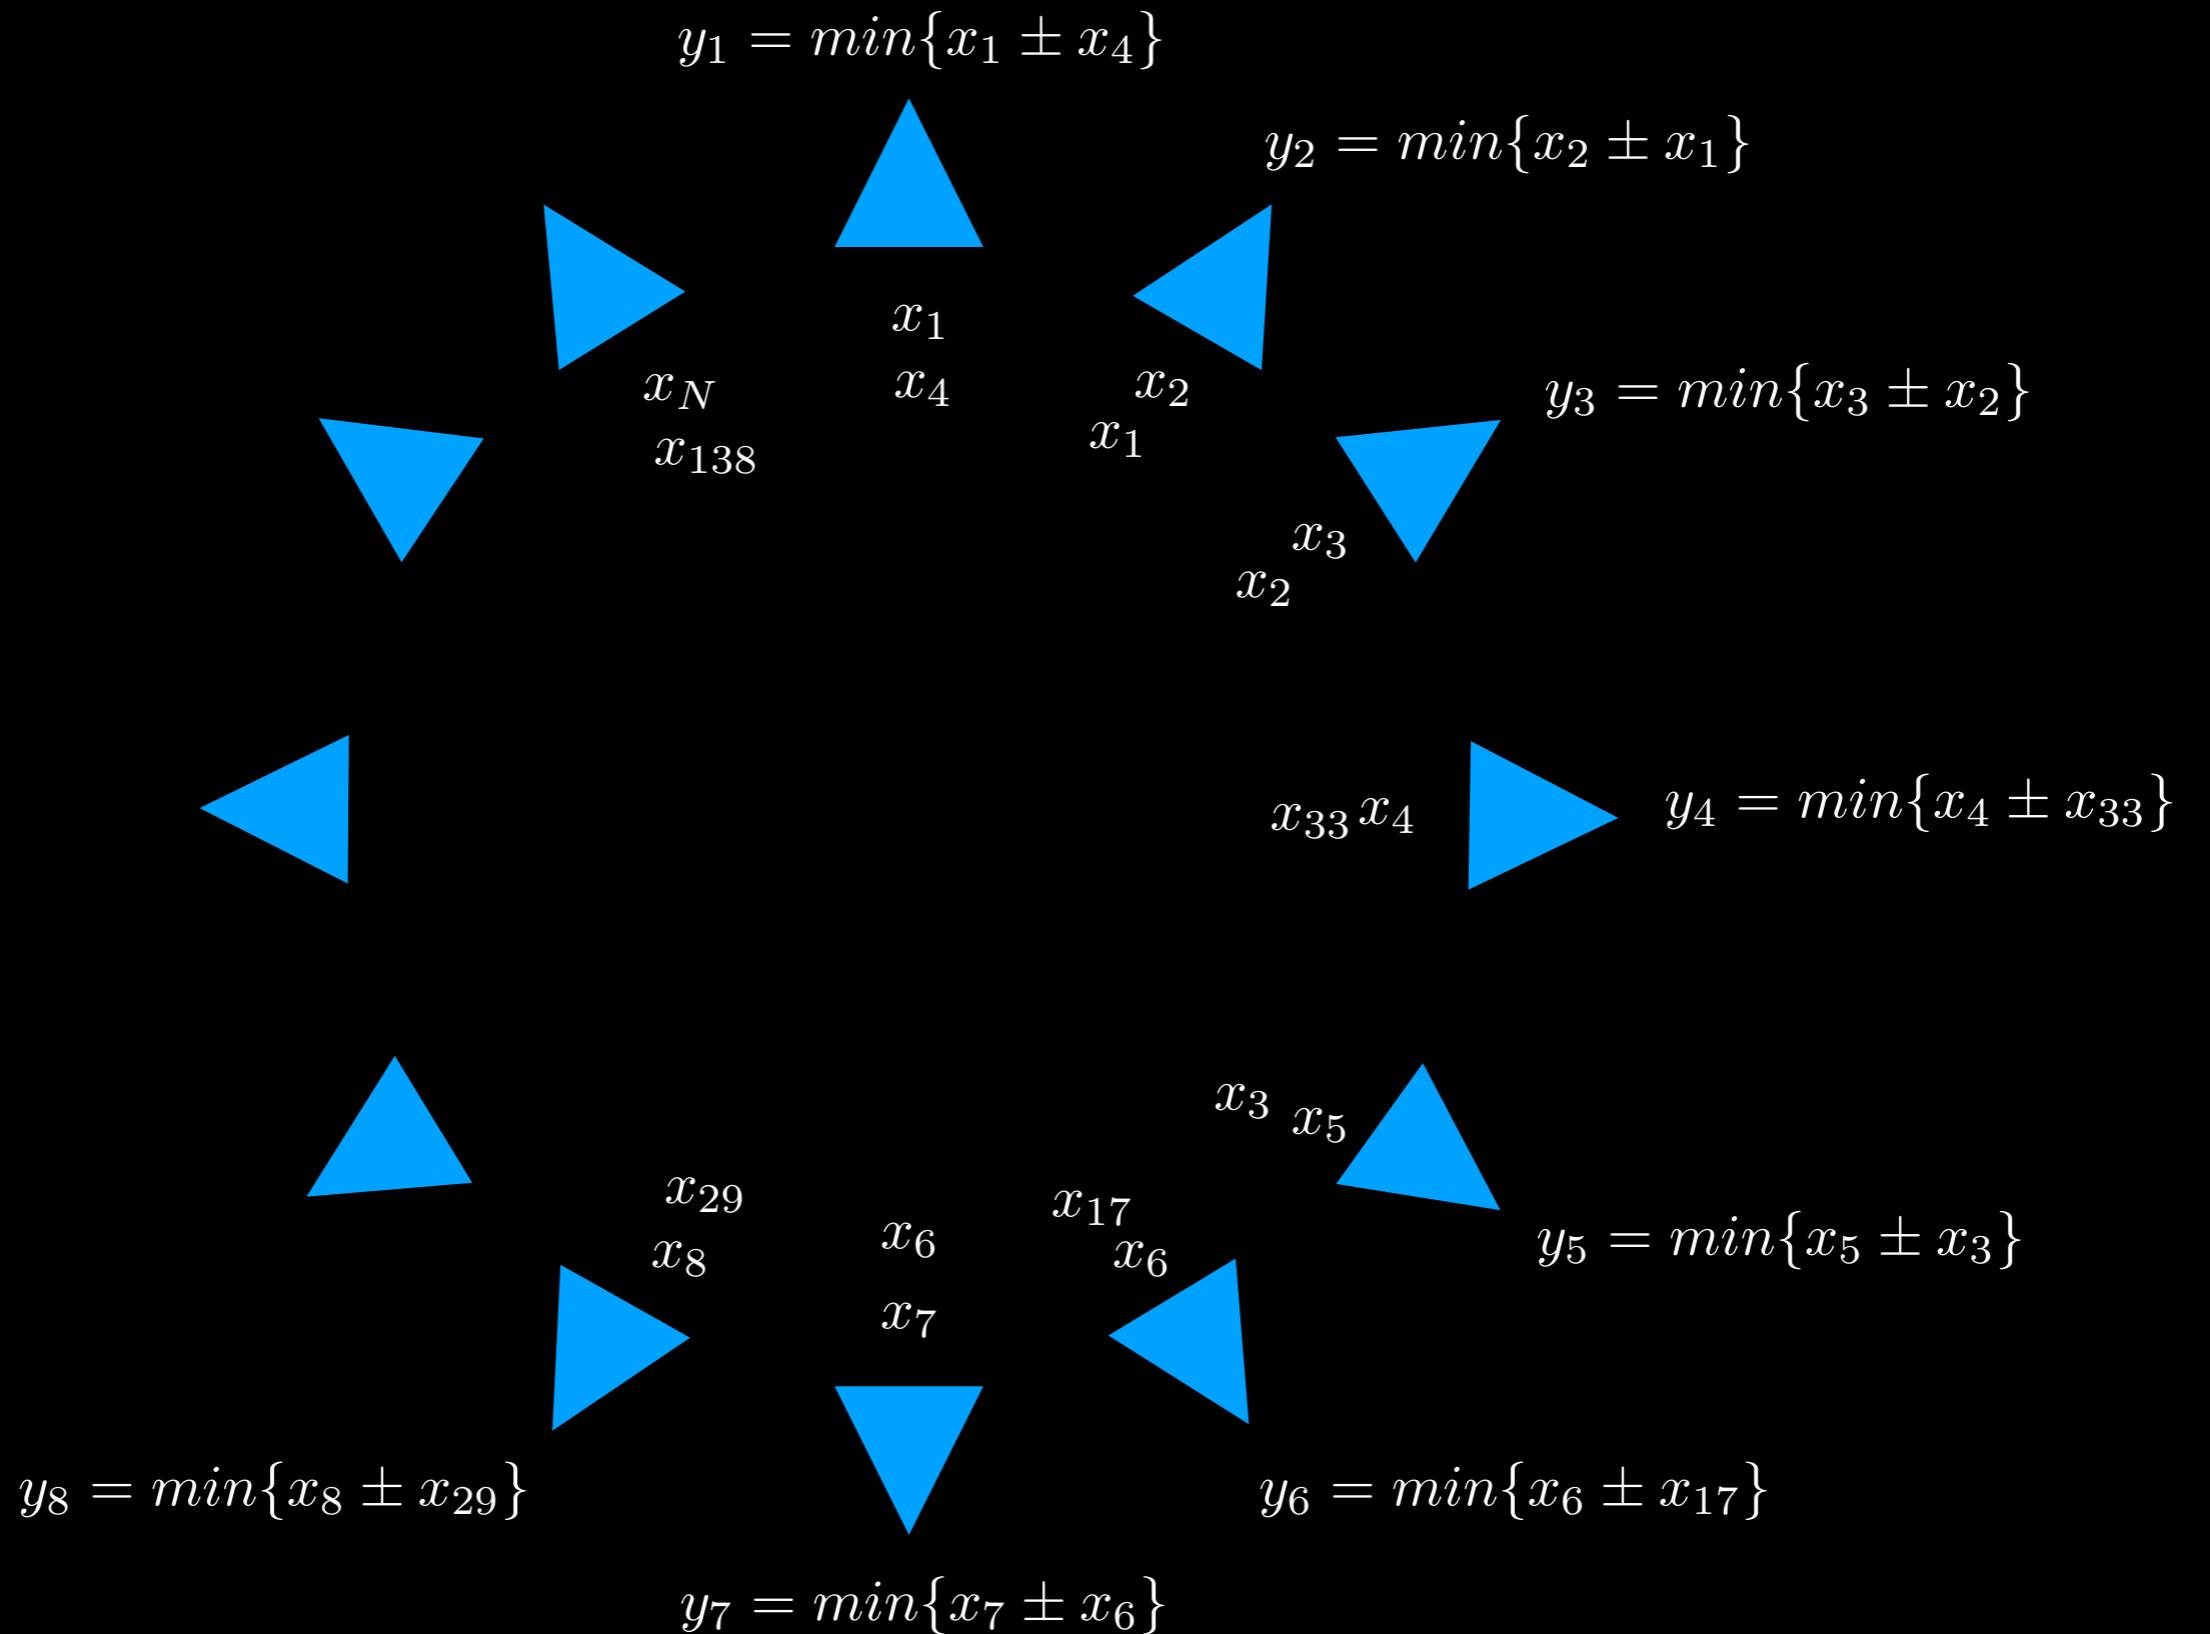
A classical algorithm for distributed sieving



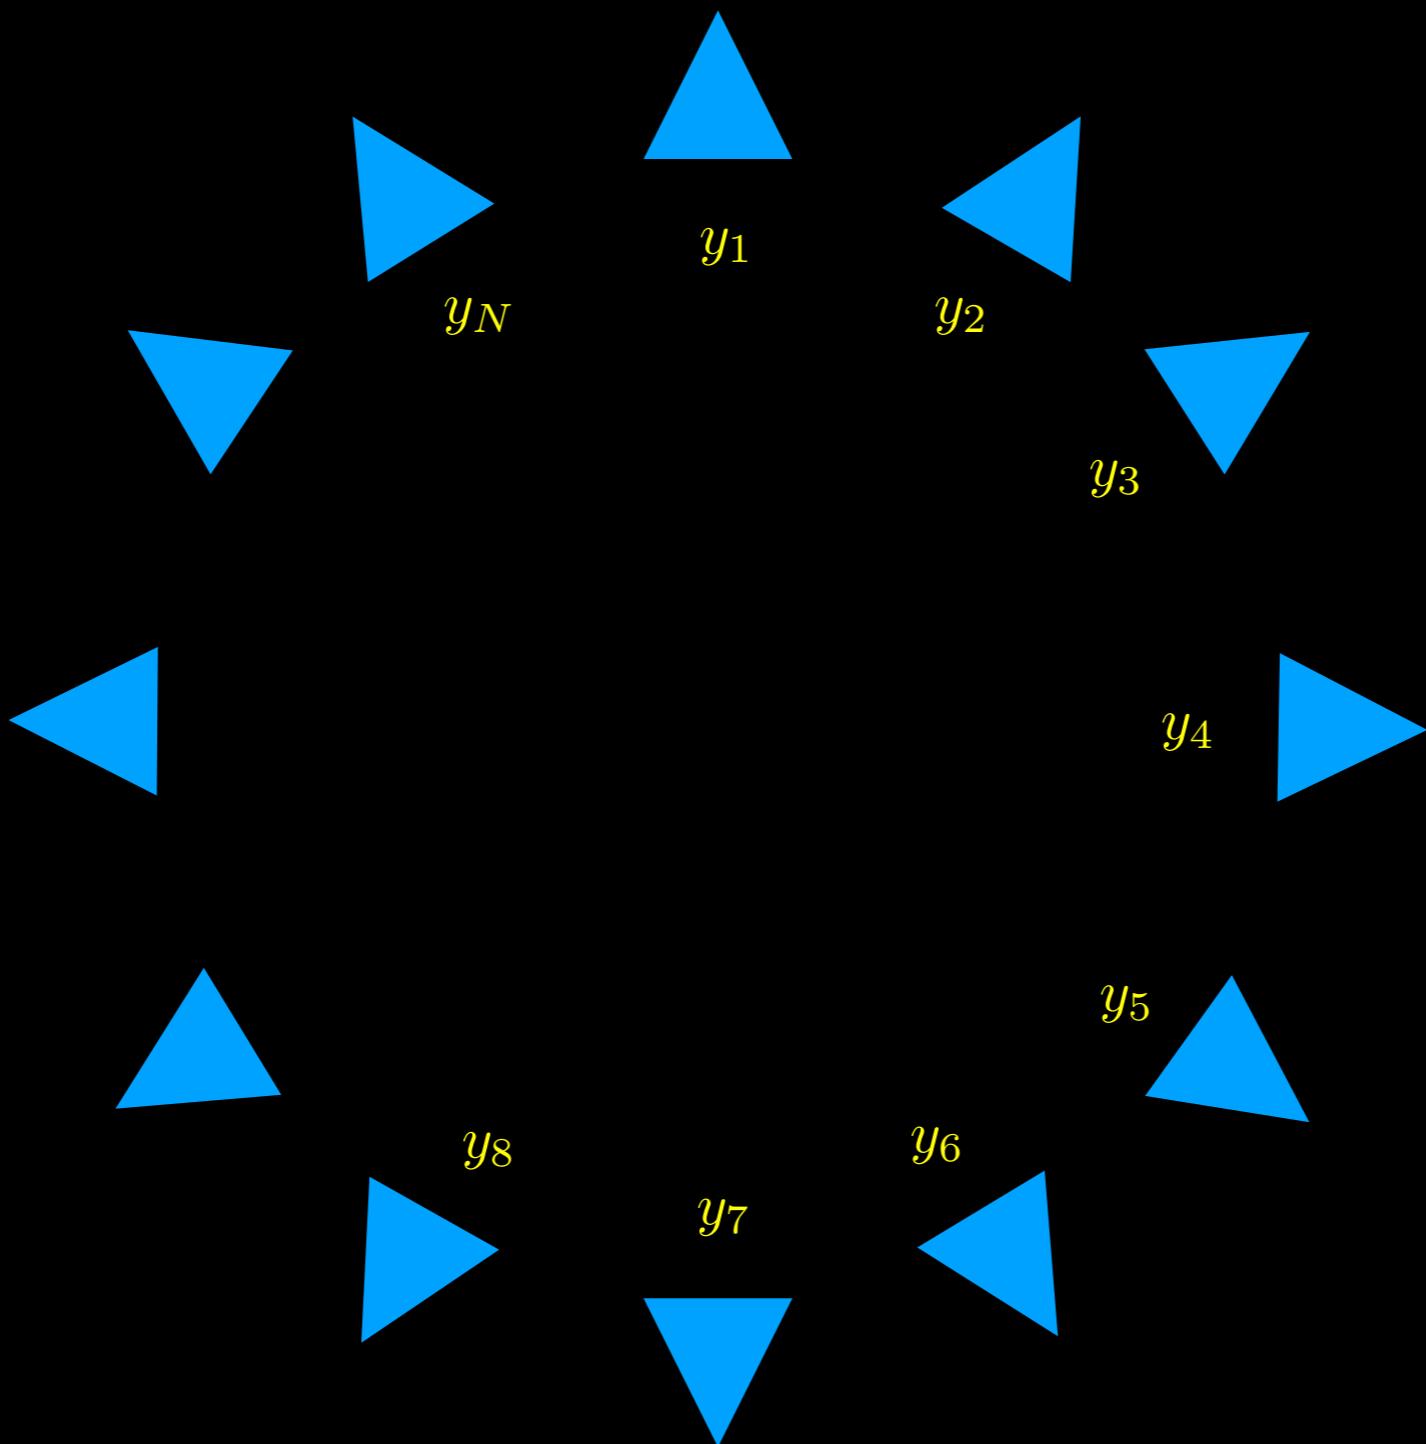
A classical algorithm for distributed sieving



A classical algorithm for distributed sieving



A classical algorithm for distributed sieving



This classical algorithm for distributed sieving can be implemented in $2^{0.2075d+o(d)}$ time steps and $2^{0.2075d+o(d)}$ memory.

Compare this to the quantum algorithm for parallel sieving that can be implemented in $2^{0.1037d+o(d)}$ time steps and $2^{0.2075d+o(d)}$ memory.

Quantum Algorithms for k-sieve

L1

L1

L2

2-sieve

For $x_i, x_j \in L_1$,
the pair (x_i, x_j) is ‘good’ iff $| \langle x_i, x_j \rangle | \geq \frac{1}{2}$

L1

L1

L2

L2

L3

2-sieve

For $x_i, x_j \in L_1$,
the pair (x_i, x_j) is ‘good’ iff $| \langle x_i, x_j \rangle | \geq \frac{1}{2}$

L1

L1

L2

L2

L3

L3

L4

L4

2-sieve

For $x_i, x_j \in L_1$,
the pair (x_i, x_j) is ‘good’ iff $| \langle x_i, x_j \rangle | \geq \frac{1}{2}$

L



k-sieve

For $k=3$, $x_i, x_j, x_k \in L_1$,
the tuple (x_i, x_j, x_k) is ‘good’ iff
 $| \langle x_i, x_j \rangle | \approx \frac{1}{3}$
 $| \langle x_j, x_k \rangle | \approx \frac{1}{3}$
 $| \langle x_k, x_i \rangle | \approx \frac{1}{3}$

L1

L1

L1

L1

L2

L2

L2

L2

k-sieve

For $k=3$, $x_i, x_j, x_k \in L_1$,
the tuple (x_i, x_j, x_k) is ‘good’ iff
 $| \langle x_i, x_j \rangle | \approx \frac{1}{3}$
 $| \langle x_j, x_k \rangle | \approx \frac{1}{3}$
 $| \langle x_k, x_i \rangle | \approx \frac{1}{3}$

L3

L1

L1

L1

L1

L2

L2

L2

L2

k-sieve

For $k=3$, $x_i, x_j, x_k \in L_1$,
the tuple (x_i, x_j, x_k) is ‘good’ iff
 $| \langle x_i, x_j \rangle | \approx \frac{1}{3}$
 $| \langle x_j, x_k \rangle | \approx \frac{1}{3}$
 $| \langle x_k, x_i \rangle | \approx \frac{1}{3}$

L3

L3

L3

L3

L

Sieving and k-clique listing

x_1

x_3

x_{138}

x_2

x_4

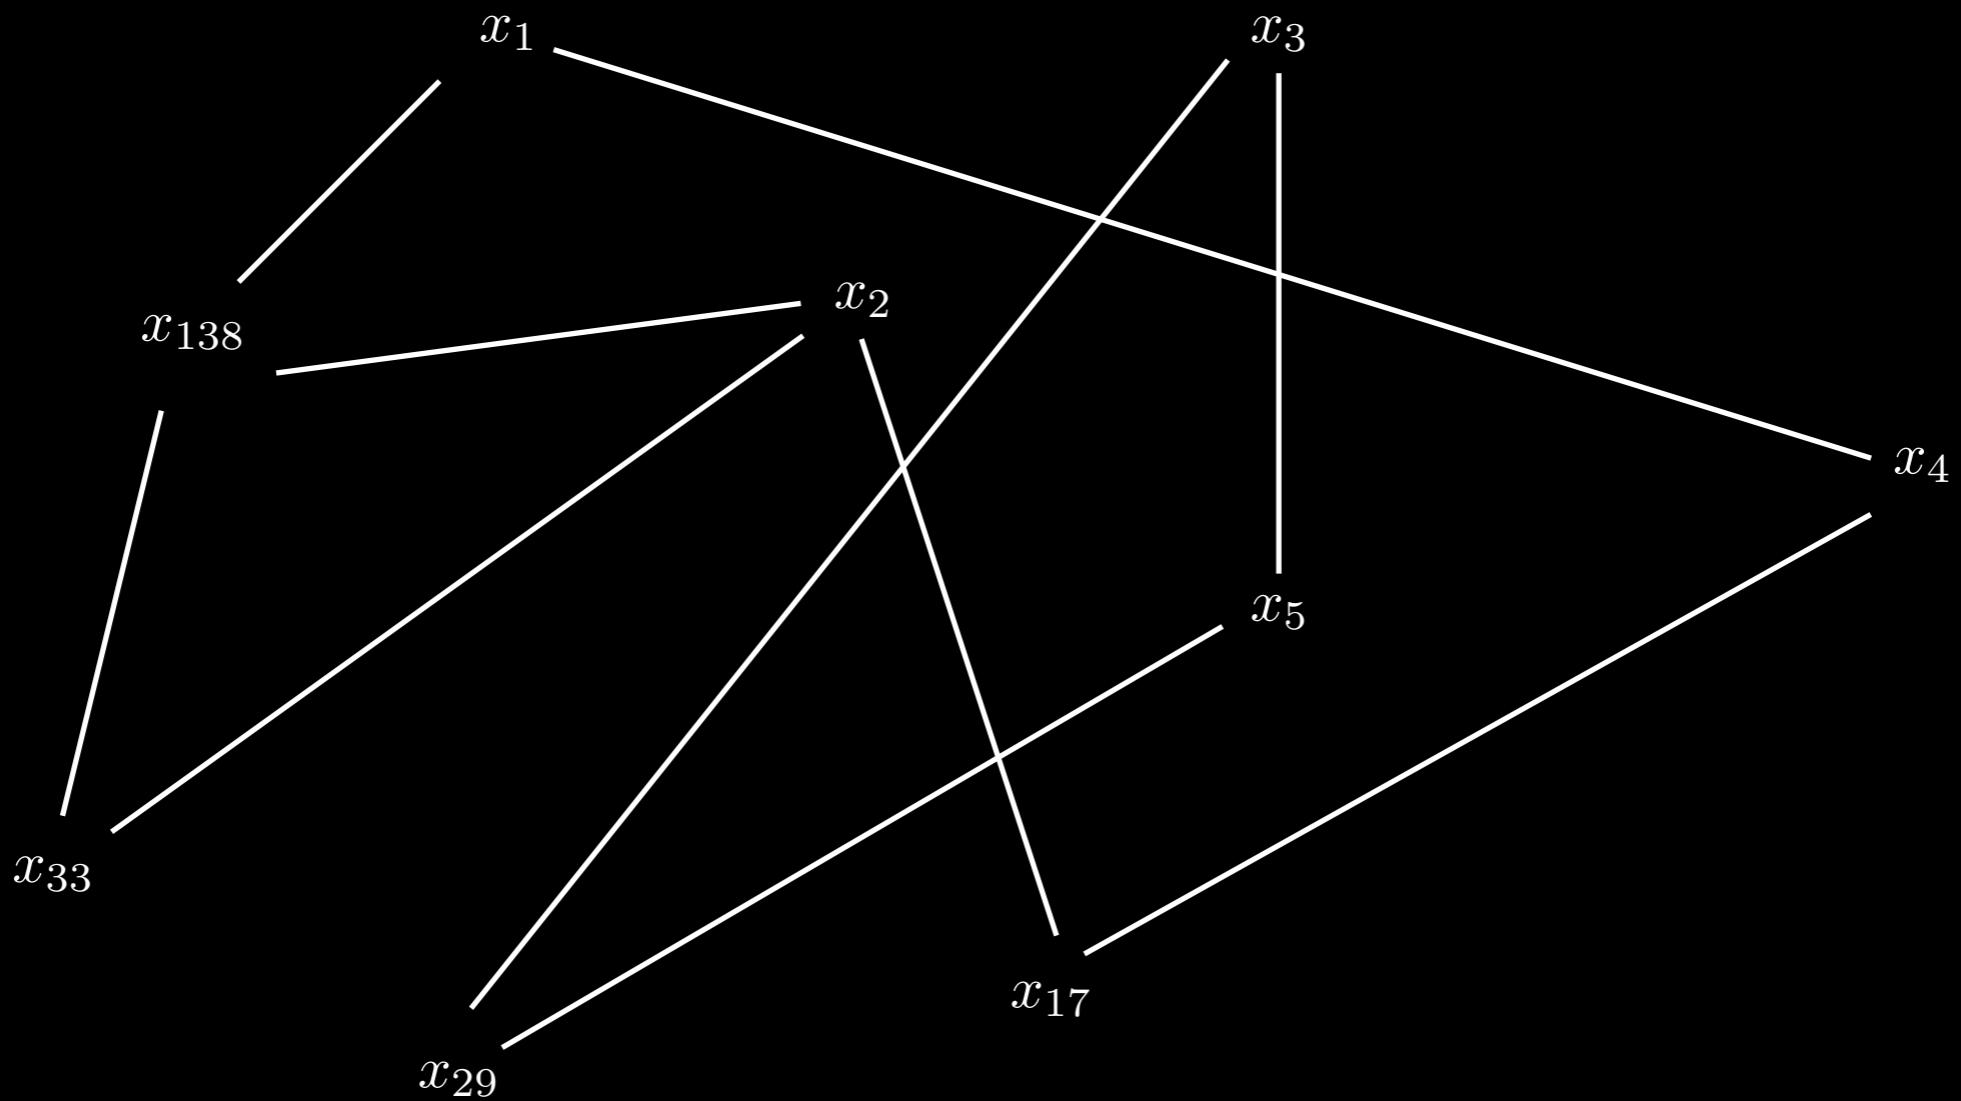
x_{33}

x_5

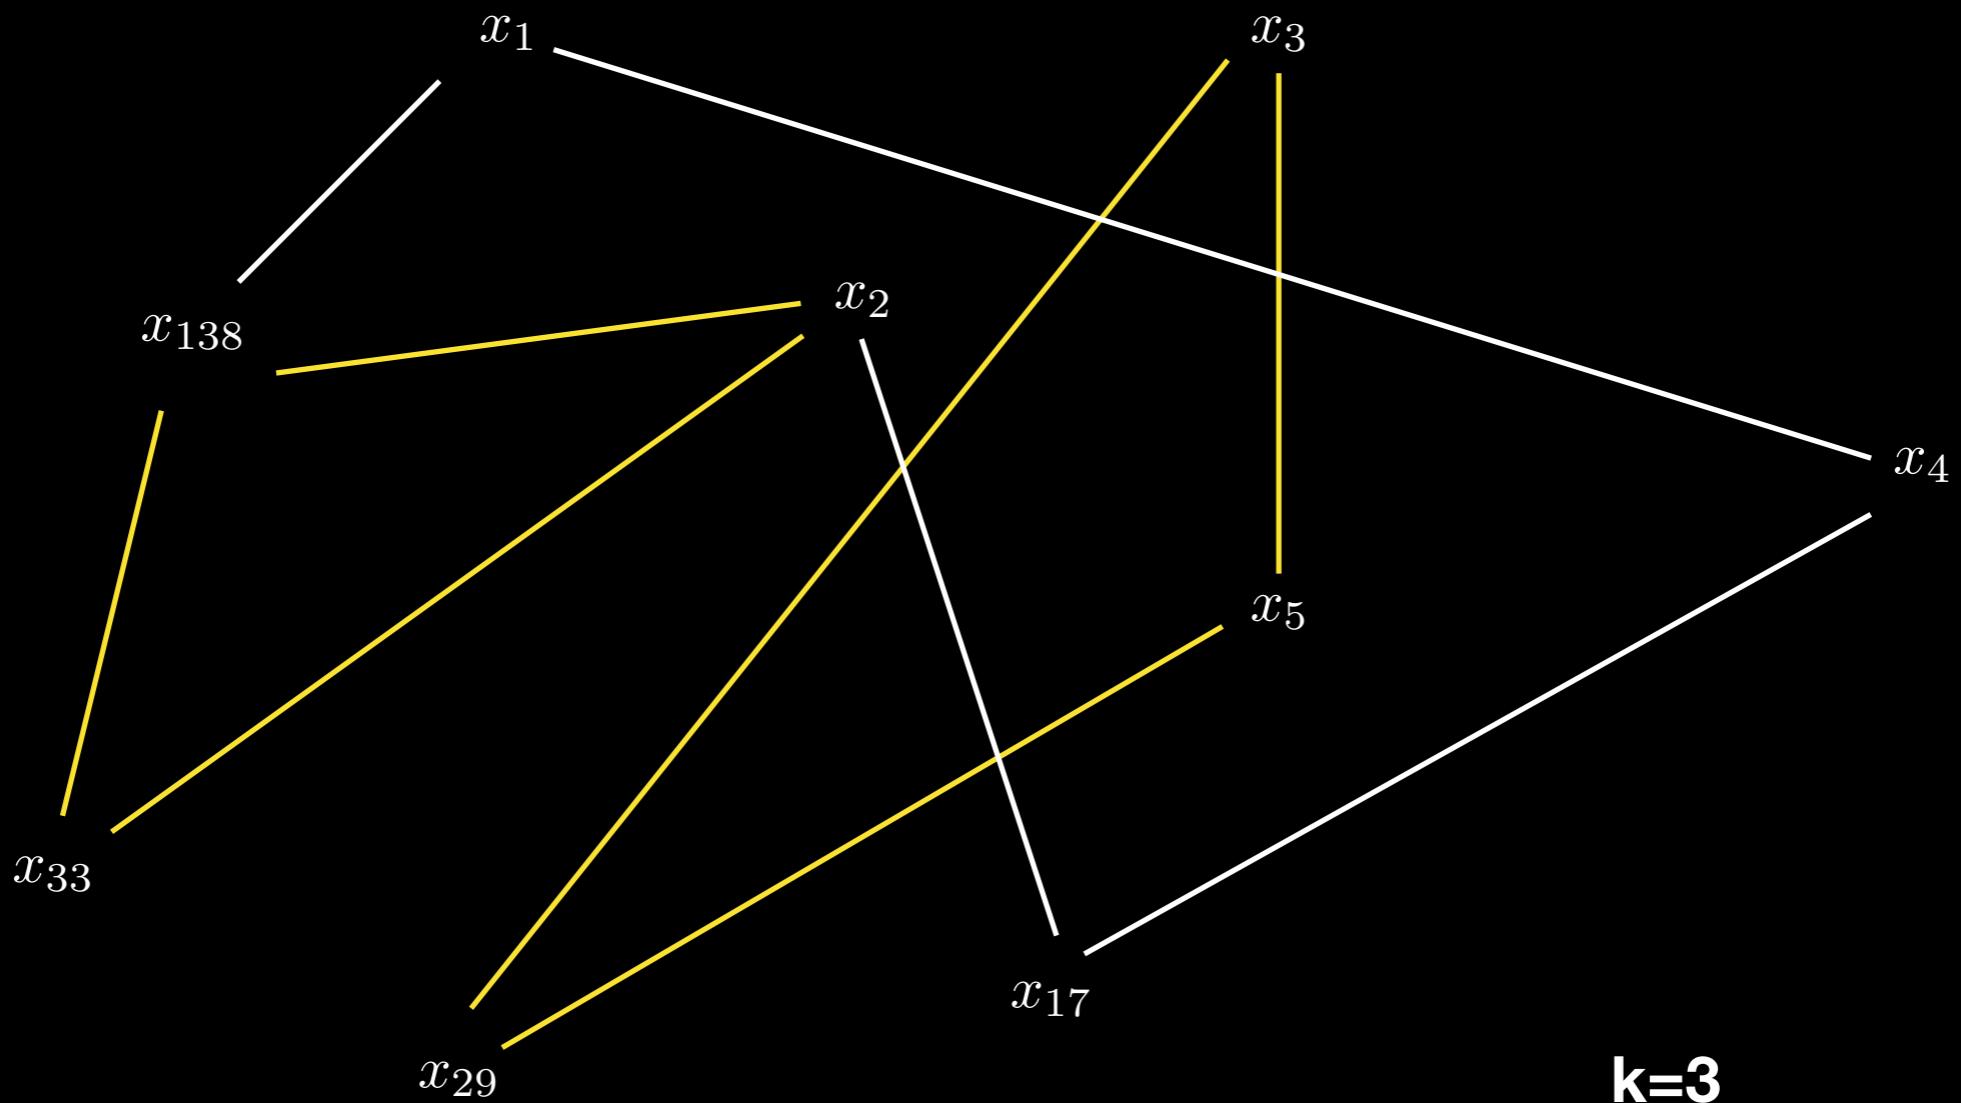
x_{17}

x_{29}

Sieving and k-clique listing



Sieving and k-clique listing



(x_{29}, x_5, x_3)

(x_{33}, x_2, x_{138})

Sieving and 3-clique listing

Triangle finder:

Buhrman et al. Quantum Algorithms for Element Distinctness,
SIAM J. of Computing, 34(6), 1324-1330, 2005

1. Use quantum search to find an edge $(x_i, x_j) \in E$
2. Use quantum search to find a node x_k , such that (x_i, x_j, x_k) is a triangle.
3. Apply Amplitude Amplification on steps 1-2.

Sieving and 3-clique listing

Triangle finder:

Buhrman et al. Quantum Algorithms for Element Distinctness,
SIAM J. of Computing, 34(6), 1324-1330, 2005

1. Use quantum search to find an edge $(x_i, x_j) \in E$
2. Use quantum search to find a node x_k , such that (x_i, x_j, x_k) is a triangle.
3. Apply Amplitude Amplification on steps 1-2.

Amplitude Amplification is a generic way to get speedups for any probabilistic algorithm. Given a Boolean function $f : \{0, 1\} \rightarrow \{0, 1\}$, suppose we have an algorithm that finds a solution $f(x) = 1$ with probability p when acting on the initial state $|0\rangle$.

Classically we would have to execute the algorithm roughly $\frac{1}{p}$ times before finding a solution.

Using Amplitude Amplification, we one need $\frac{1}{\sqrt{p}}$ repetitions.

Sieving and 3-clique listing

Triangle finder:

Buhrman et al. Quantum Algorithms for Element Distinctness,
SIAM J. of Computing, 34(6), 1324-1330, 2005

1. Use quantum search to find an edge $(x_i, x_j) \in E$
2. Use quantum search to find a node x_k , such that (x_i, x_j, x_k) is a triangle.
3. Apply Amplitude Amplification on steps 1-2.

In the graph, corresponding to 3-sieve, we set $N = \tilde{O}((3\sqrt{3}/4)^{d/2})$,
 $m = \tilde{O}(n^2(8/9)^{d/2})$, and it is expected there are $\Theta(N)$ triangles.

Step 1, needs $\Theta(\sqrt{\frac{N^2}{m}})$ time steps.

Step 2, needs $\Theta(\sqrt{N})$ time steps

Step 3, needs $\Theta(\sqrt{\frac{m}{N}})$ time steps.

Total cost: $\Theta(\sqrt{m})$ to find one triangle

Sieving and 3-clique listing

Triangle finder:

Buhrman et al. Quantum Algorithms for Element Distinctness,
SIAM J. of Computing, 34(6), 1324-1330, 2005

1. Use quantum search to find an edge $(x_i, x_j) \in E$
2. Use quantum search to find a node x_k , such that (x_i, x_j, x_k) is a triangle.
3. Apply Amplitude Amplification on steps 1-2.

In the graph, corresponding to 3-sieve, we set $N = \tilde{O}((3\sqrt{3}/4)^{d/2})$,
 $m = \tilde{O}(n^2(8/9)^{d/2})$, and it is expected there are $\Theta(N)$ triangles.

Step 1, needs $\Theta(\sqrt{\frac{N^2}{m}})$ time steps.

Step 2, needs $\Theta(\sqrt{N})$ time steps

Step 3, needs $\Theta(\sqrt{\frac{m}{N}})$ time steps.

Total cost: $\Theta(\sqrt{m})$ to find one triangle

Repeat this $N * \log N$ times, to list all triangles.

Sieving and 3-clique listing

Triangle finder:

Buhrman et al. Quantum Algorithms for Element Distinctness,
SIAM J. of Computing, 34(6), 1324-1330, 2005

1. Use quantum search to find an edge $(x_i, x_j) \in E$
2. Use quantum search to find a node x_k , such that (x_i, x_j, x_k) is a triangle.
3. Apply Amplitude Amplification on steps 1-2.

In the graph, corresponding to 3-sieve, we set $N = \tilde{O}((3\sqrt{3}/4)^{d/2})$,
 $m = \tilde{O}(n^2(8/9)^{d/2})$, and it is expected there are $\Theta(N)$ triangles.

Step 1, needs $\Theta(\sqrt{\frac{N^2}{m}})$ time steps.

Step 2, needs $\Theta(\sqrt{N})$ time steps

Step 3, needs $\Theta(\sqrt{\frac{m}{N}})$ time steps.

Total cost: $\Theta(\sqrt{m})$ to find one triangle

Repeat this $N * \log N$ times, to list all triangles.

Each sieving step needs $\tilde{O}(\sqrt{m} * N) = 2^{0.3349d+o(d)}$ time steps.

And we need $\tilde{O}(N) = 2^{0.1887d+o(d)}$ classical memory*.

And $\text{poly}(d)$ quantum memory.

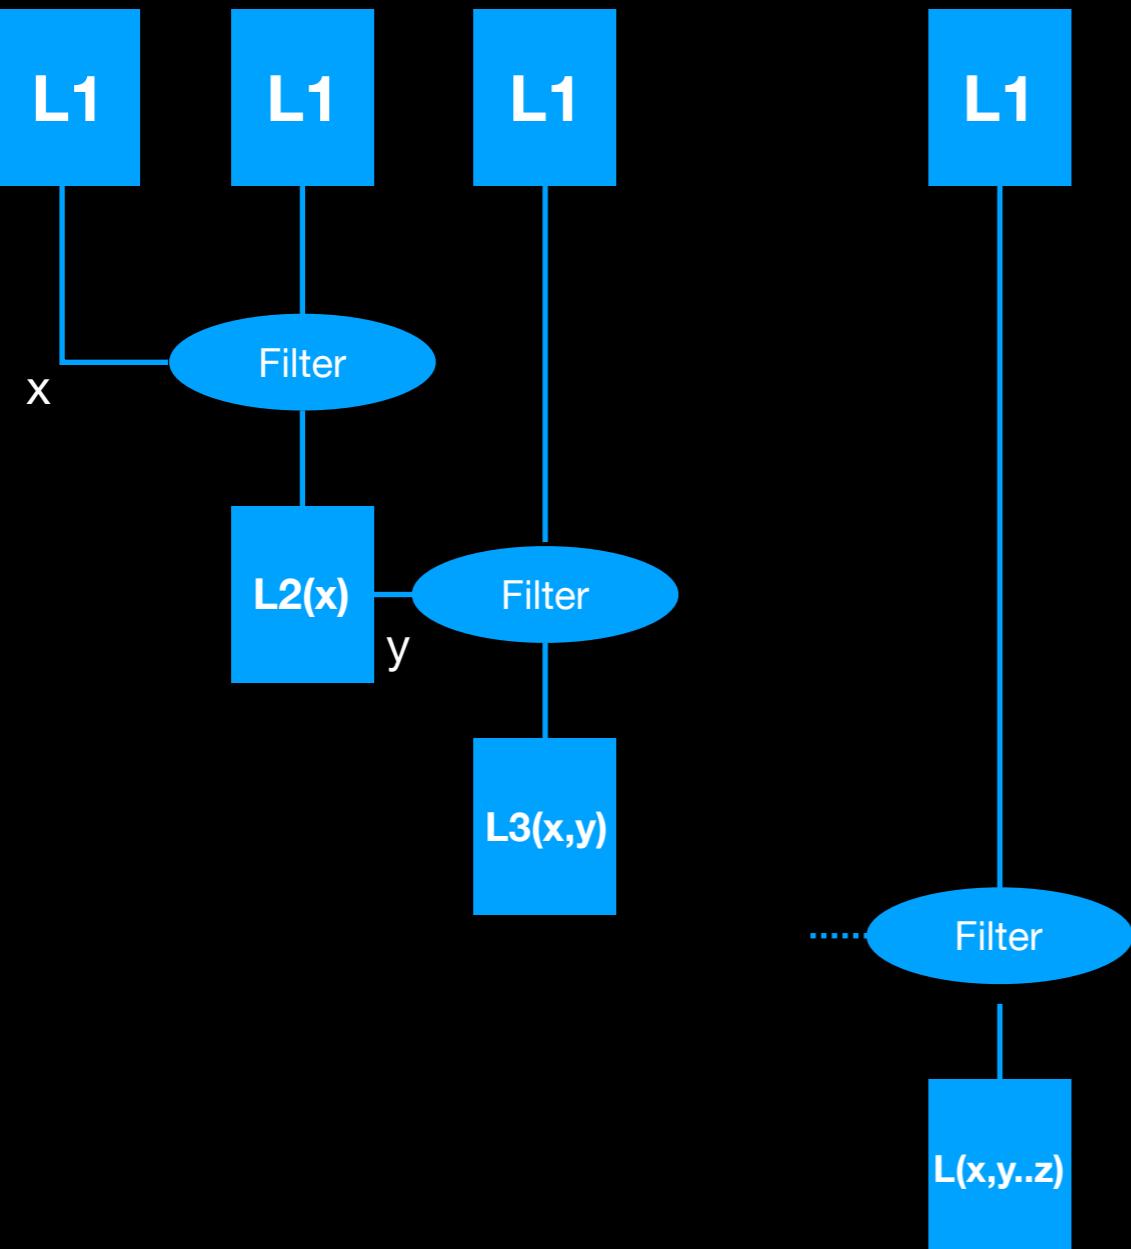
Sieving and k-clique listing

K-Clique finder:

1. Use quantum search to find an edge $(x_i, x_j) \in E$
2. Use quantum search to find a node x_k , such that (x_i, x_j, x_k) is a triangle.
3. Use quantum search to find a node x_l , such that (x_i, x_j, x_k, x_l) is a 4-clique.
- ⋮
- k-1. Use quantum search to find a node x_c , such that $(x_i, x_j, x_k, x_l \dots x_c)$ is a k-clique.
- k. Apply Amplitude Amplification on steps 1 – (k – 1).

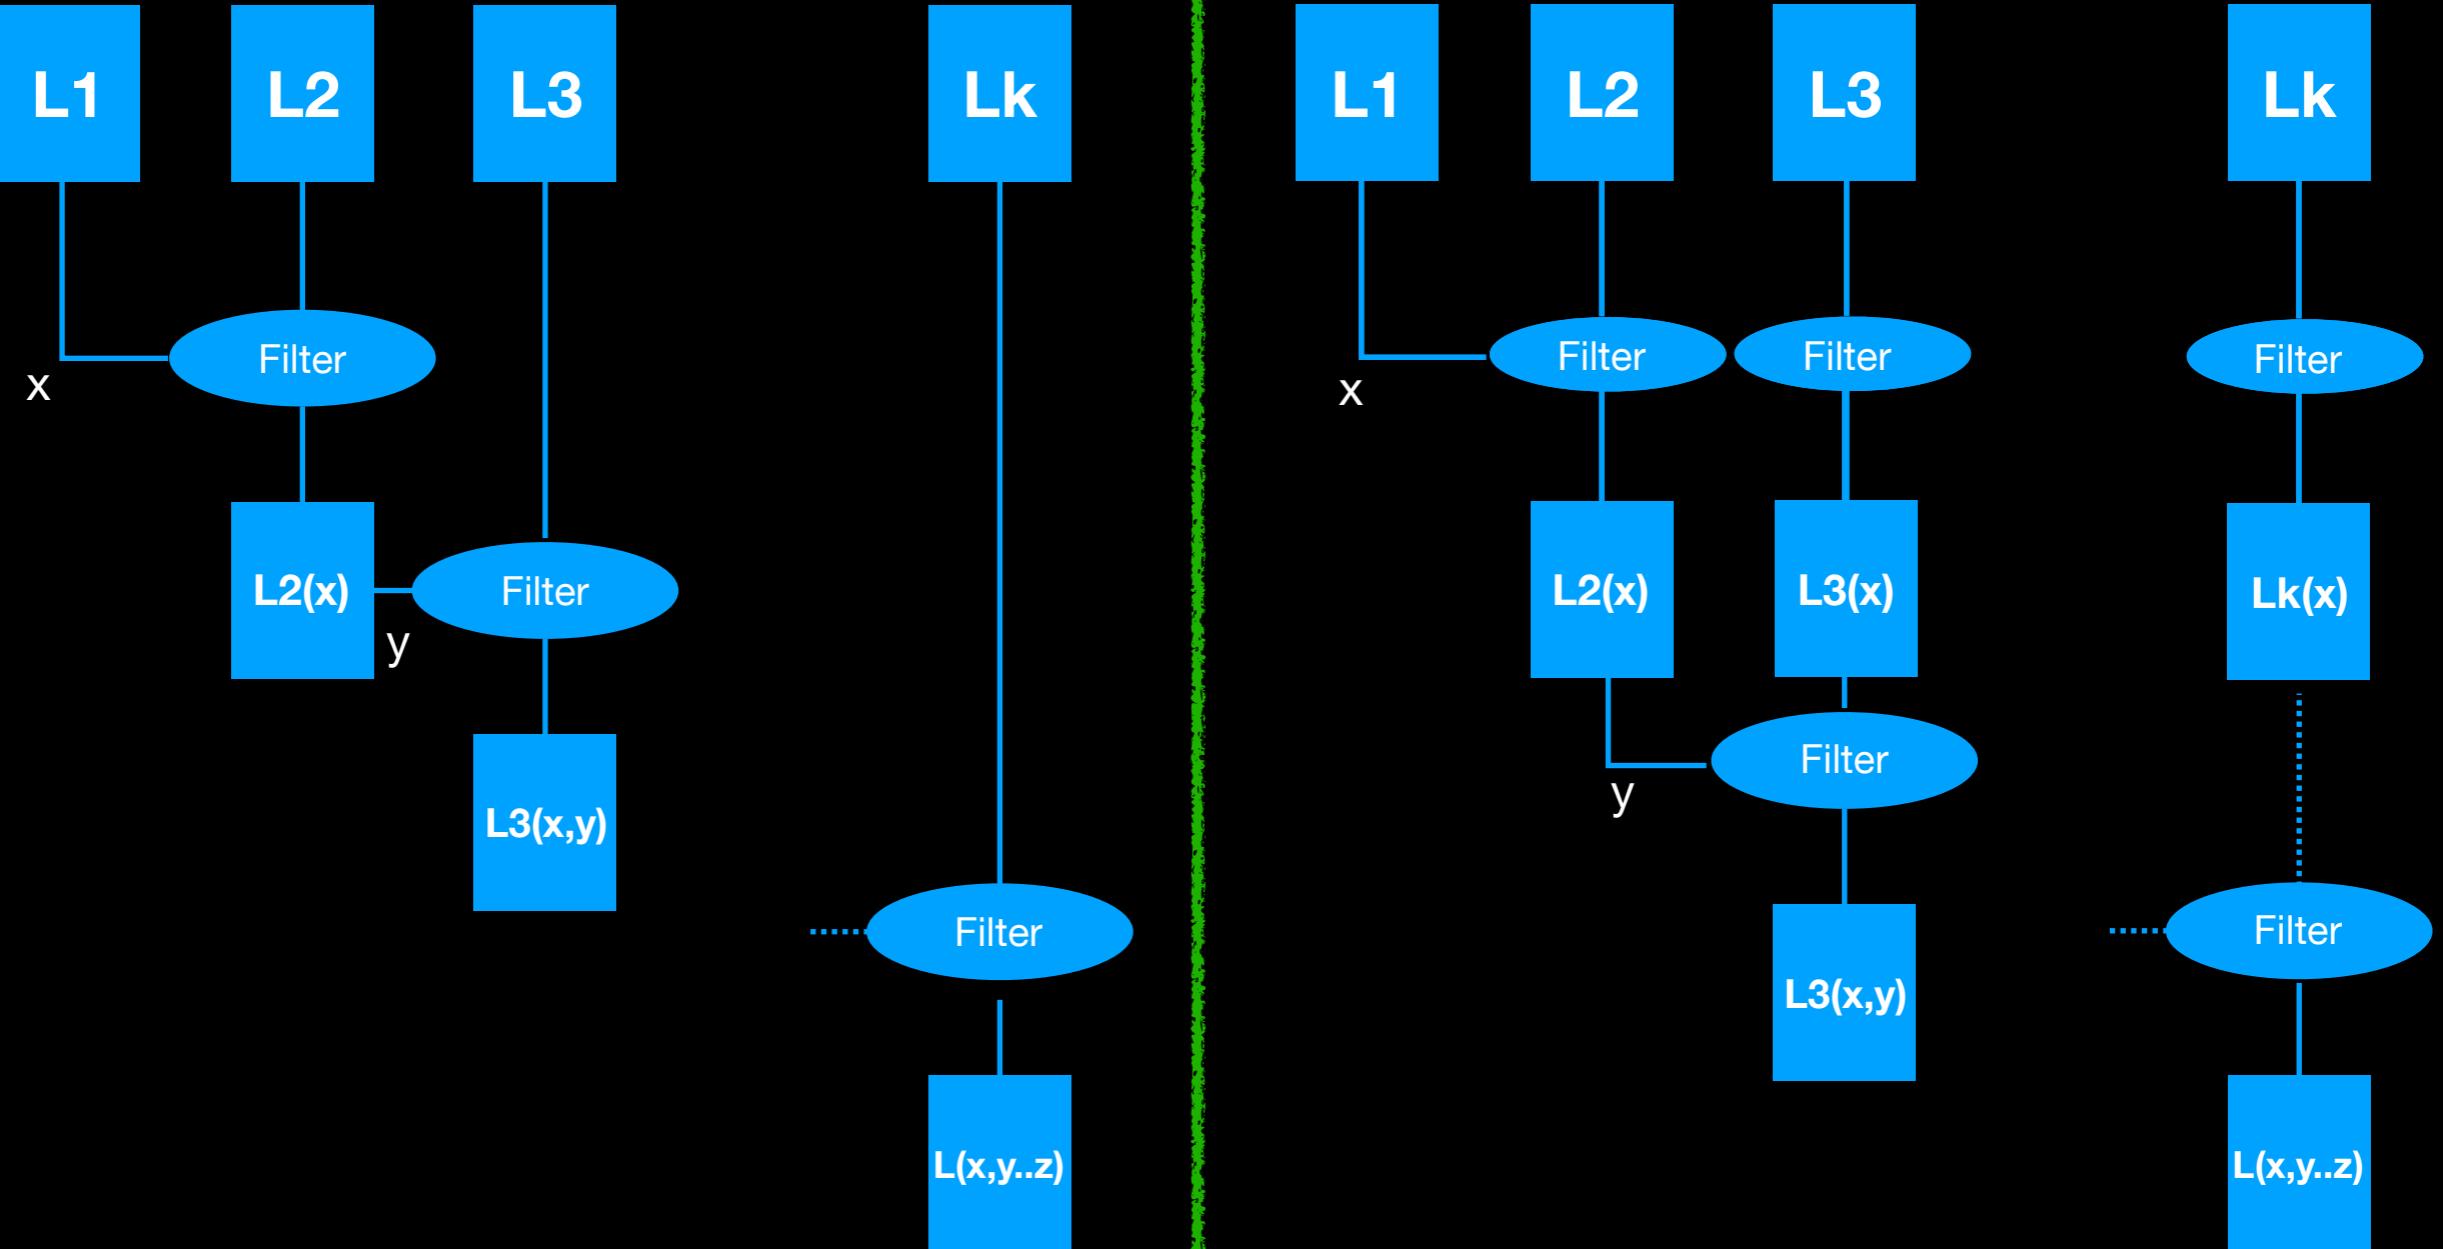
If there are N expected k-Cliques, Repeat the k-Clique finder algorithm $N * \log N$ times, to list them all.

Sieving and k-clique listing



Filter receives as input vectors from the two lists, and a real number C , the target inner product. This procedure only passes the vectors whose inner product is close to C .

Sieving and k-clique listing



Sieving and k-clique listing

Quantum
BLS16

K	2	3	4	5	...	28	29	30
TIME	0.3112	0.3306	0.3289	0.3219		0.2989	0.2989	0.2989
MEMORY	0.2075	0.1907	0.1796	0.1685	...	0.1395	0.1395	0.1395

optimising for time

Quantum
BLS16+HKL18

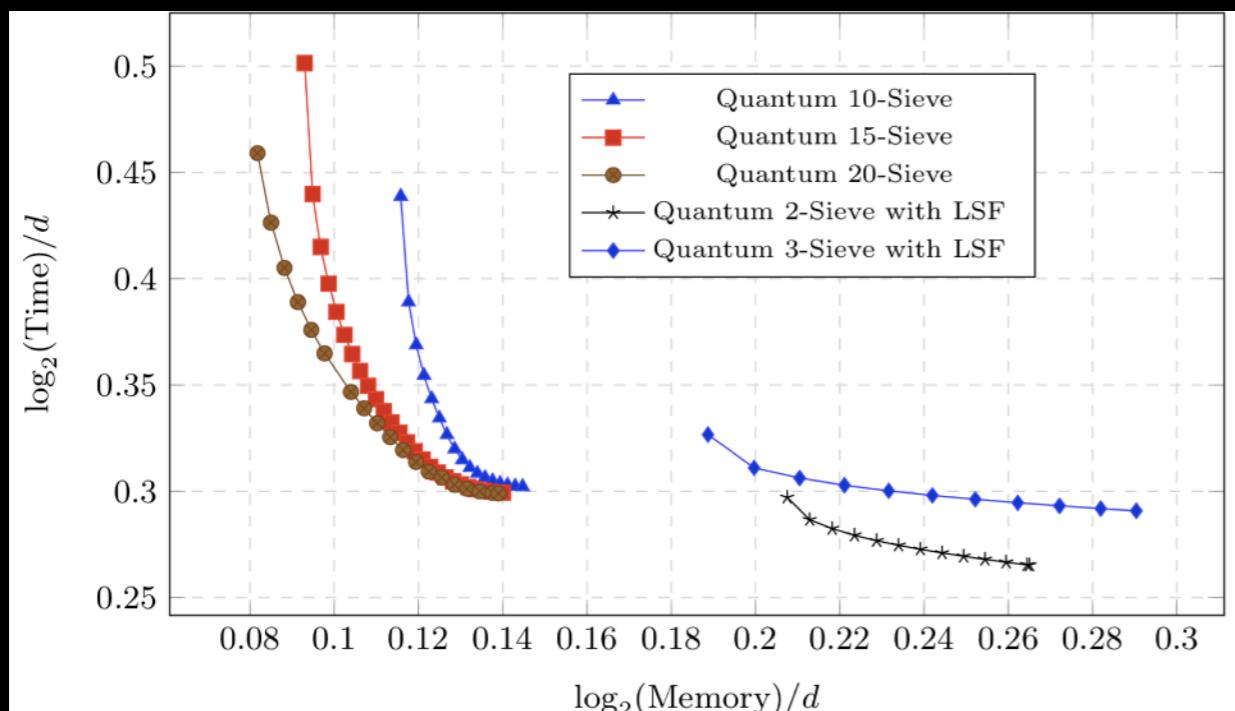
TIME	0.3112	0.3306	0.3197	0.3088		0.2989	0.2989	0.2989
MEMORY	0.2075	0.1907	0.1731	0.1638	...	0.1395	0.1395	0.1395

optimising for time

Quantum
BLS16+HKL18

TIME	0.3112	0.3349	0.3215	0.3305		0.6352	0.2423	0.6490
MEMORY	0.2075	0.1887	0.1724	0.1587	...	0.0637	0.0623	0.0609

optimising for memory



Each point on the curve represents (Memory, Time) optimised for time, for a fixed memory

In Summary

Can solve SVP, heuristically, in $2^{0.2989d+o(d)}$ time steps and $2^{1395d+o(d)}$ classical memory, in the **QRAM model**, requiring $\text{poly}(d)$ quantum memory.

Can solve SVP, heuristically, in $2^{1037d+o(d)}$ time steps in the **quantum circuit model** requiring $2^{0.2075d+o(d)}$ quantum memory.

In Summary

Can solve SVP, heuristically, in $2^{0.2989d+o(d)}$ time steps and $2^{1395d+o(d)}$ classical memory, in the **QRAM model**, requiring $\text{poly}(d)$ quantum memory.

Can solve SVP, heuristically, in $2^{1037d+o(d)}$ time steps in the **quantum circuit model** requiring $2^{0.2075d+o(d)}$ quantum memory.

Open questions

1. k-sieve for $k=f(n)$
2. Parallel k-sieve

In Summary

Can solve SVP, heuristically, in $2^{0.2989d+o(d)}$ time steps and $2^{1395d+o(d)}$ classical memory, in the **QRAM model**, requiring $\text{poly}(d)$ quantum memory.

Can solve SVP, heuristically, in $2^{1037d+o(d)}$ time steps in the **quantum circuit model** requiring $2^{0.2075d+o(d)}$ quantum memory.

Open questions

1. k-sieve for $k=f(n)$
2. Parallel k-sieve